



OpenGeoHUB
Connect • Create • Share • Repeat

<https://shorturl.at/uXNUp>

Optimizing Global Time Series ML Predictions From Modeling to Production Deployment



Leandro Parente



leandro.parente@opengeohub.org



<https://opengeohub.org>



+ OGH crew



Global Pasture Watch

A consortium of research partners developing
global grassland and livestock monitoring data
for impact



Land &
Carbon Lab



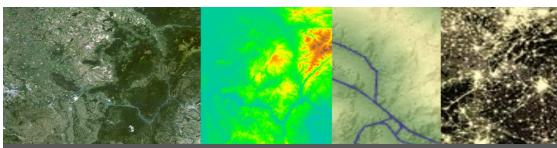
IIASA



WORLD
RESOURCES
INSTITUTE

+ User communities





Earth Observation / Gridded data

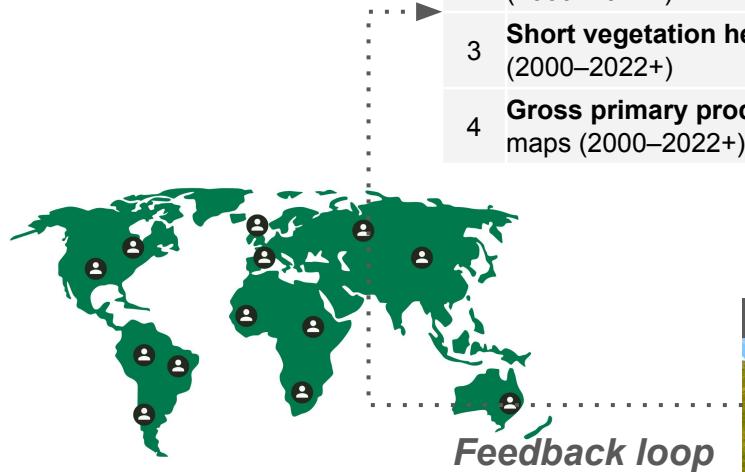


Reference samples

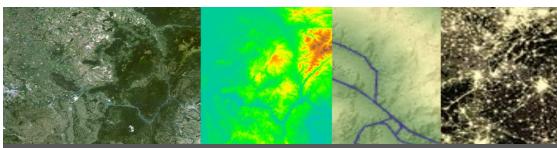


Machine learning

| # | Mapping product | Temporal resolution | Spatial resolution | Pixel value |
|---|--|---------------------|--------------------|---------------------------------|
| 1 | Grassland dynamics maps (2000–2022+) | Annual | 30-m | Two class of grass |
| 2 | Livestock density maps (2000–2022+) | Annual | 1-km | Number of animals per hectare |
| 3 | Short vegetation height maps (2000–2022+) | Annual | 30-m | Canopy height in meters |
| 4 | Gross primary productivity maps (2000–2022+) | Bi-monthly | 30-m | Kg of carbon per m ² |



Use Cases



Earth Observation / Gridded data

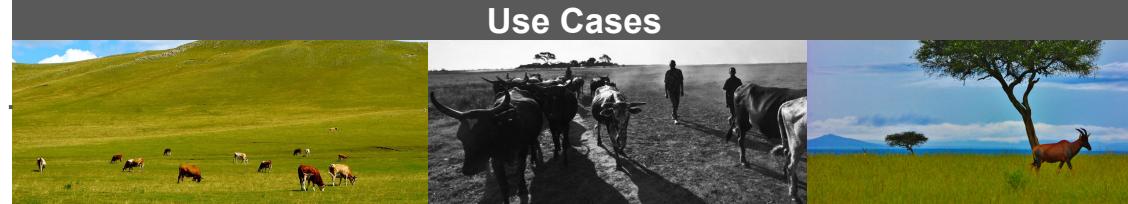
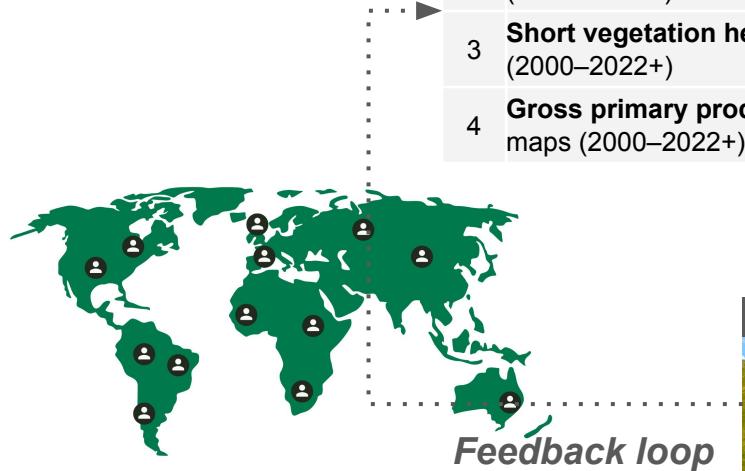


Reference samples

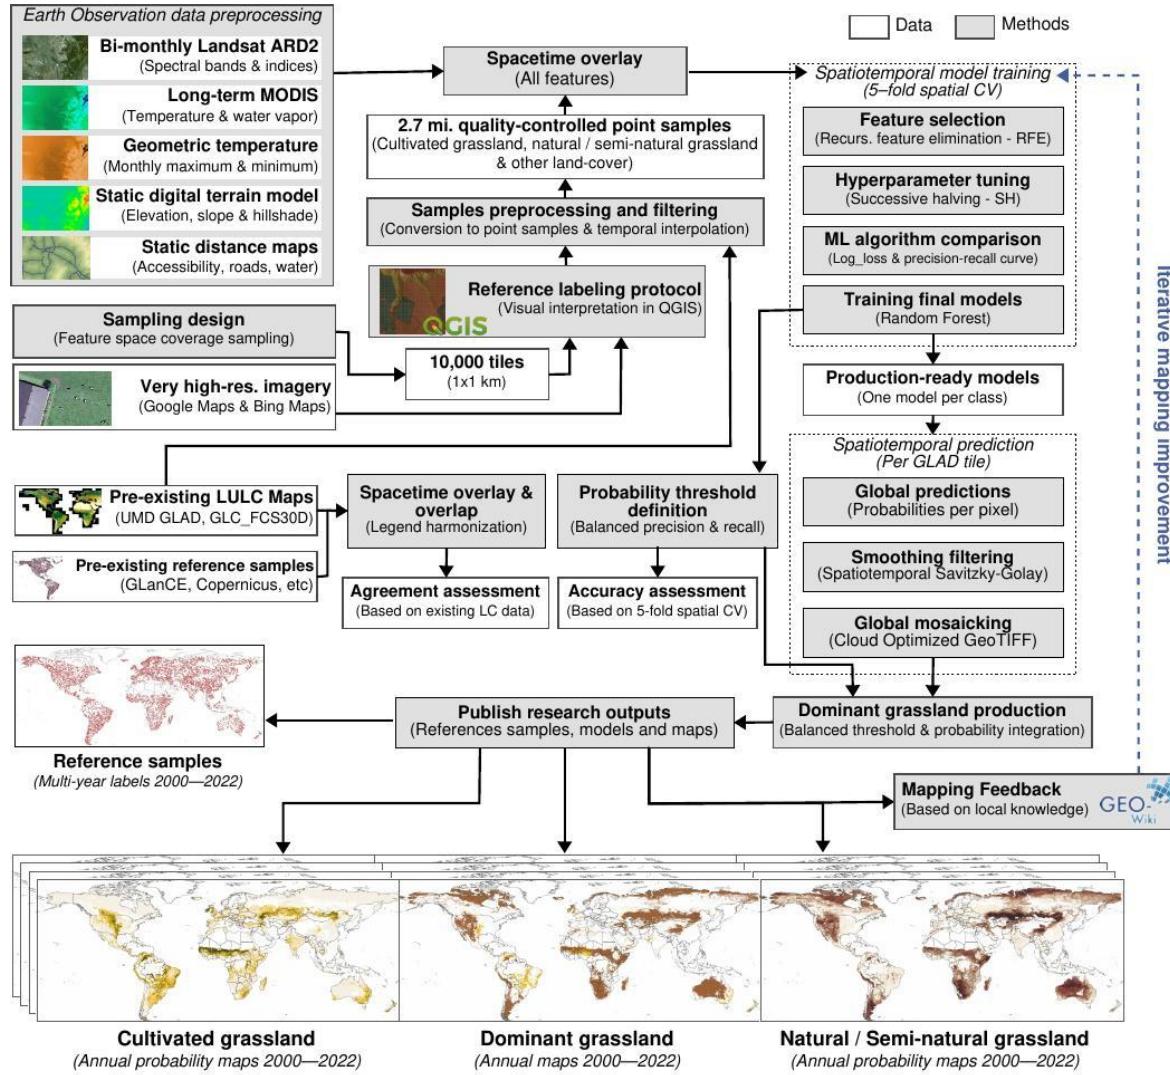


Machine learning

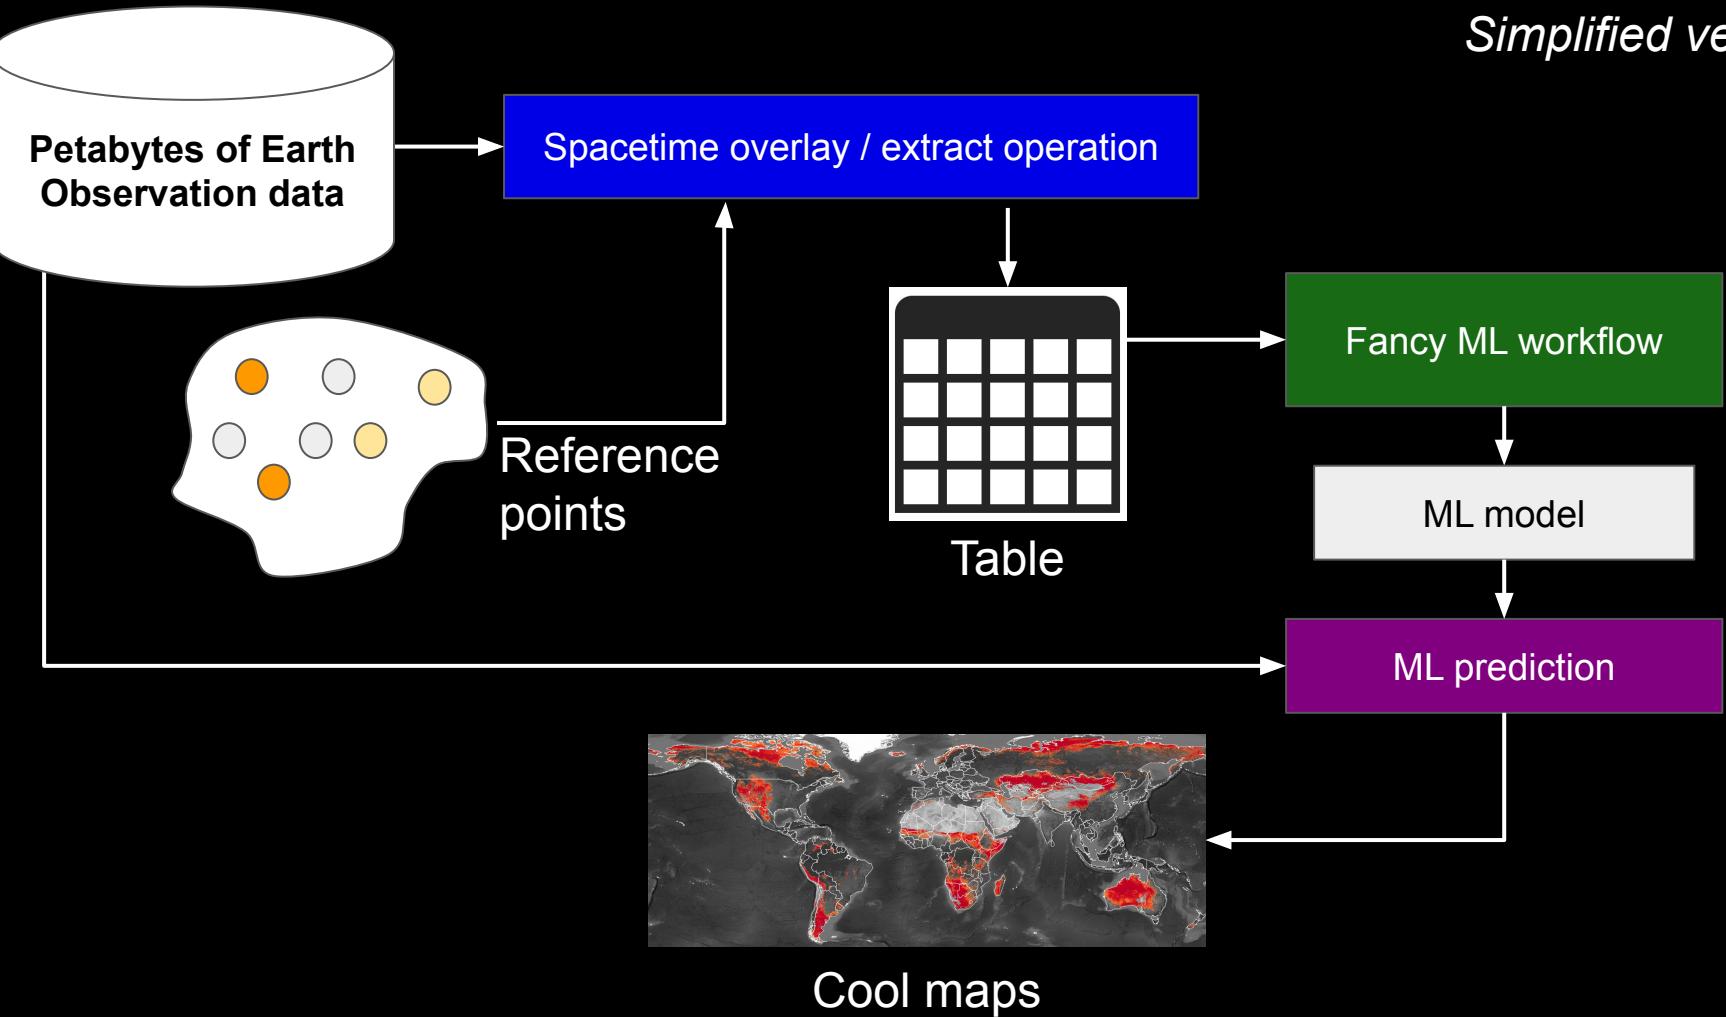
| # | Mapping product | Temporal resolution | Spatial resolution | Pixel value |
|---|---|---------------------|--------------------|---------------------------------|
| 1 | Grassland dynamics maps (2000–2022+) | Annual | 30-m | Two class of grass |
| 2 | Livestock density maps (2000–2022+) | Annual | 1-km | Number of animals per hectare |
| 3 | Short vegetation height maps (2000–2022+) | Annual | 30-m | Canopy height in meters |
| 4 | Gross primary productivity maps (2000–2022+) | Bi-monthly | 30-m | Kg of carbon per m ² |



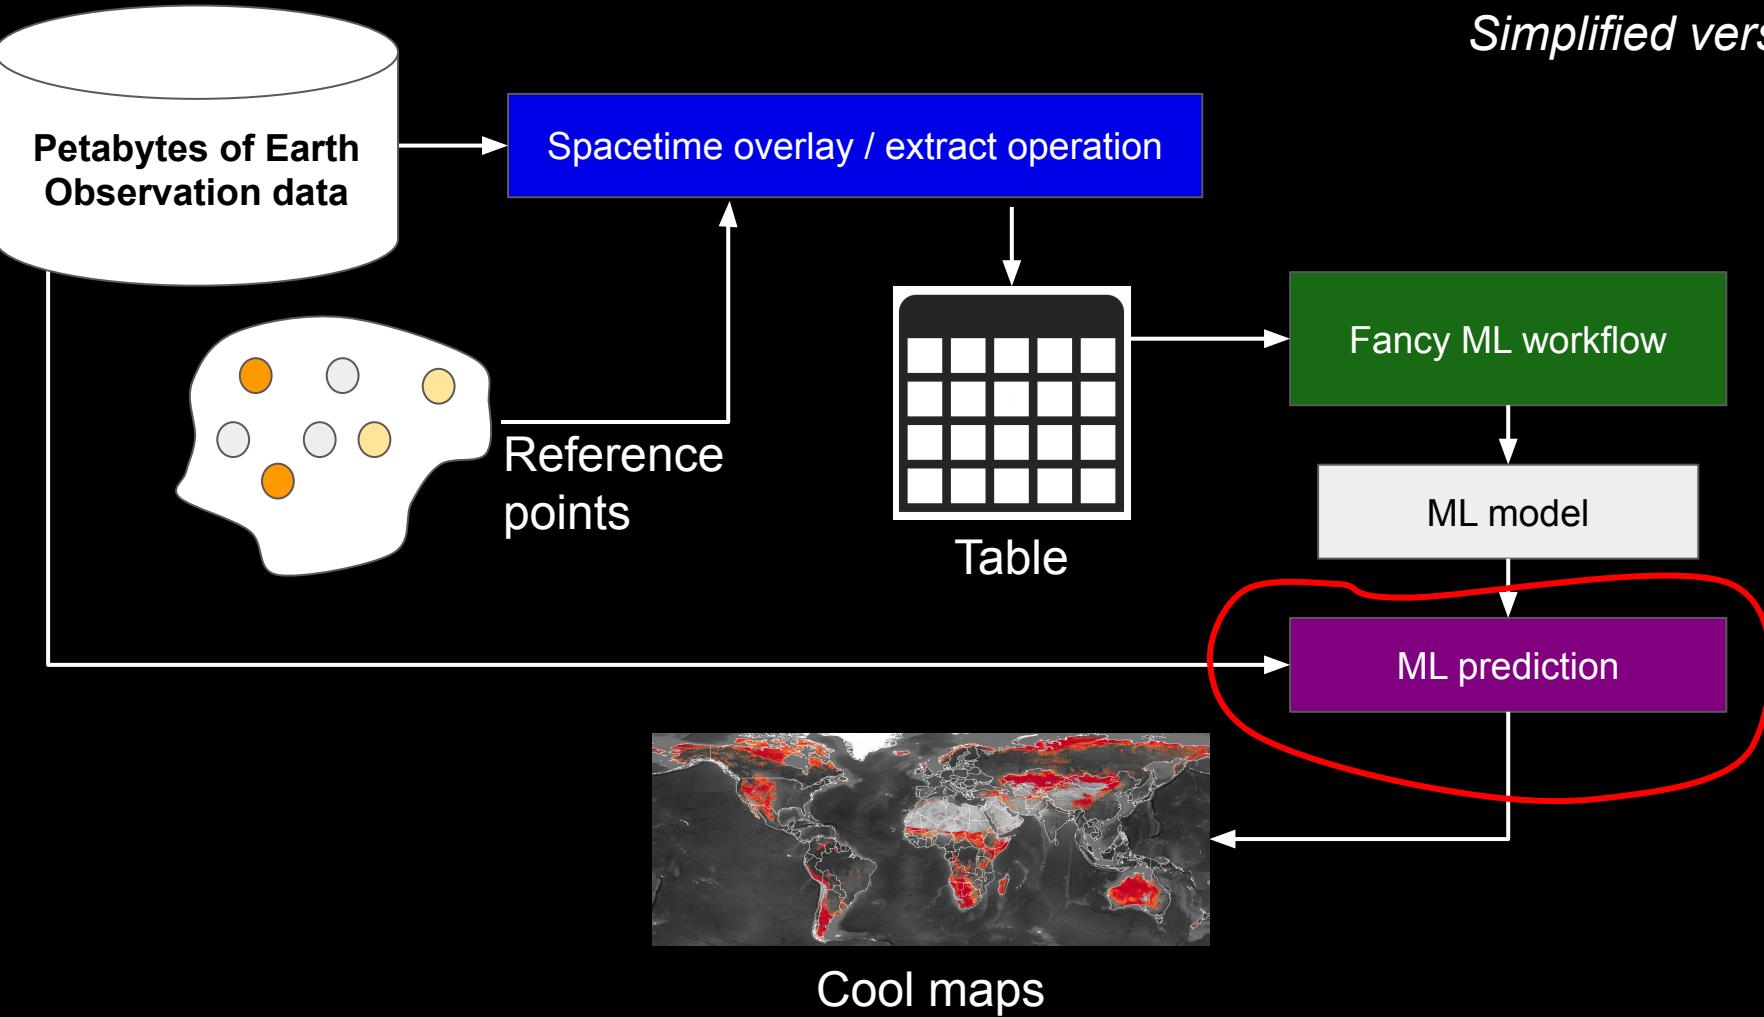
Use Cases



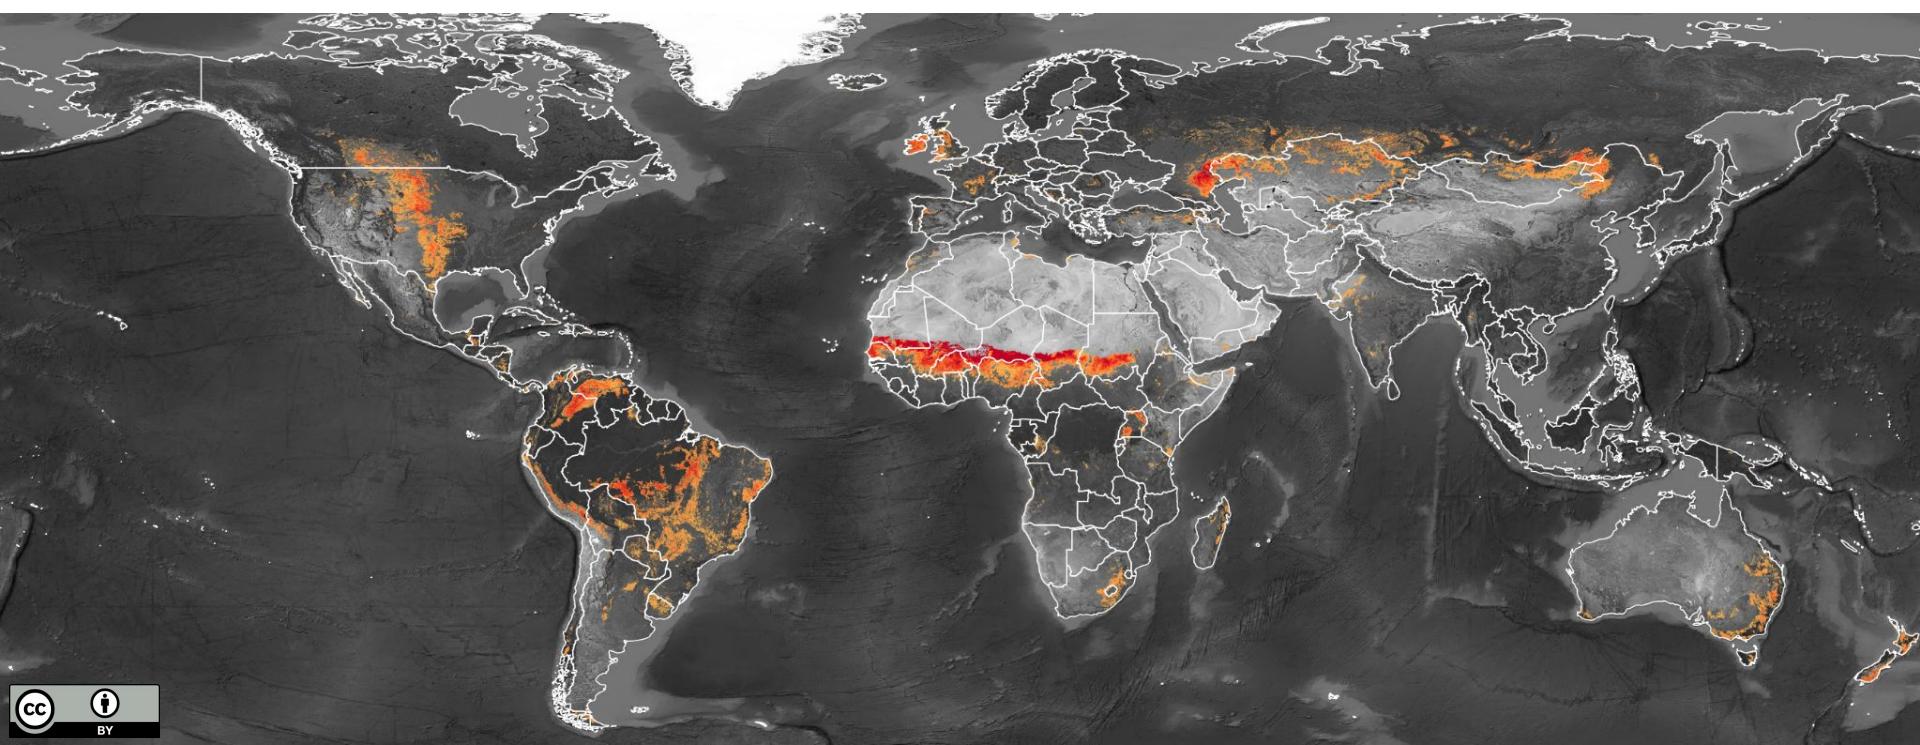
Simplified version



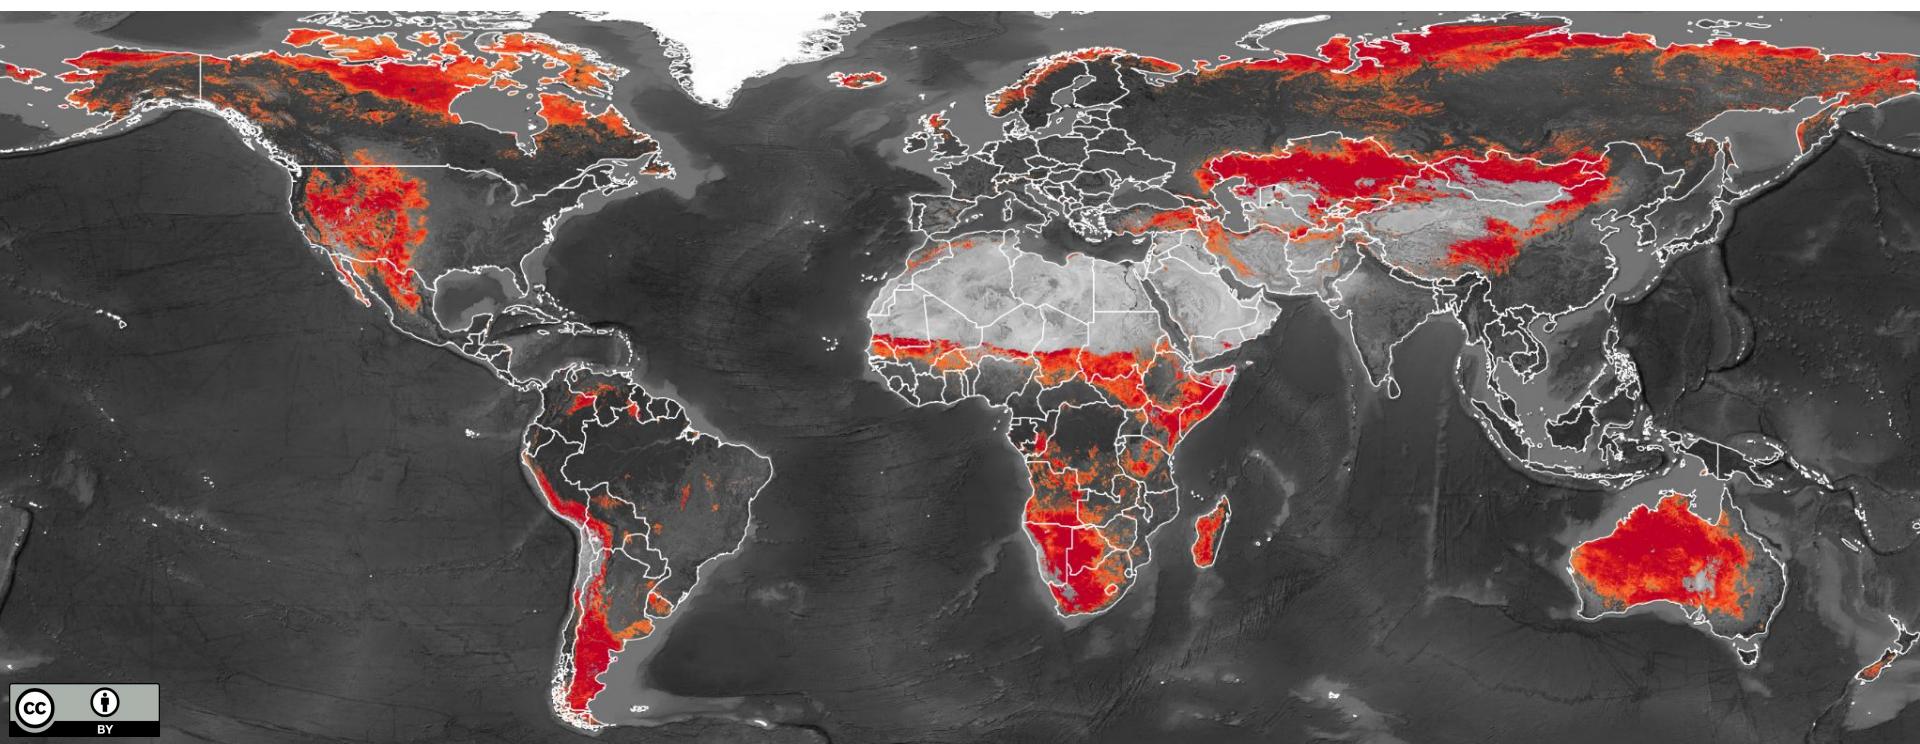
Simplified version



Cultivated grassland - 2022

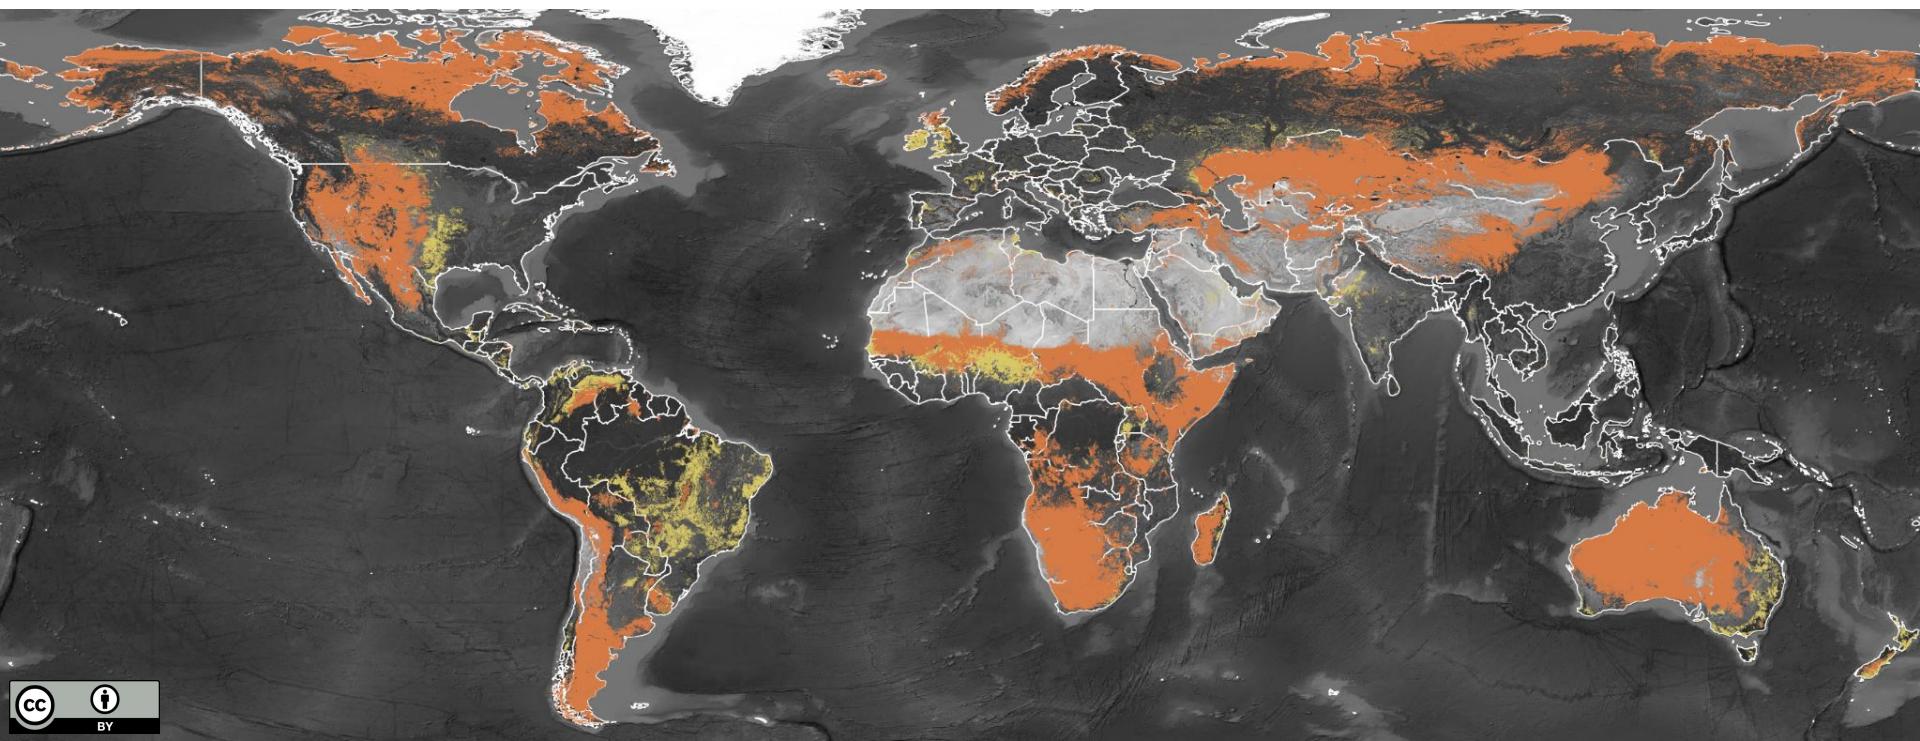


Natural / Semi-natural grassland - 2022

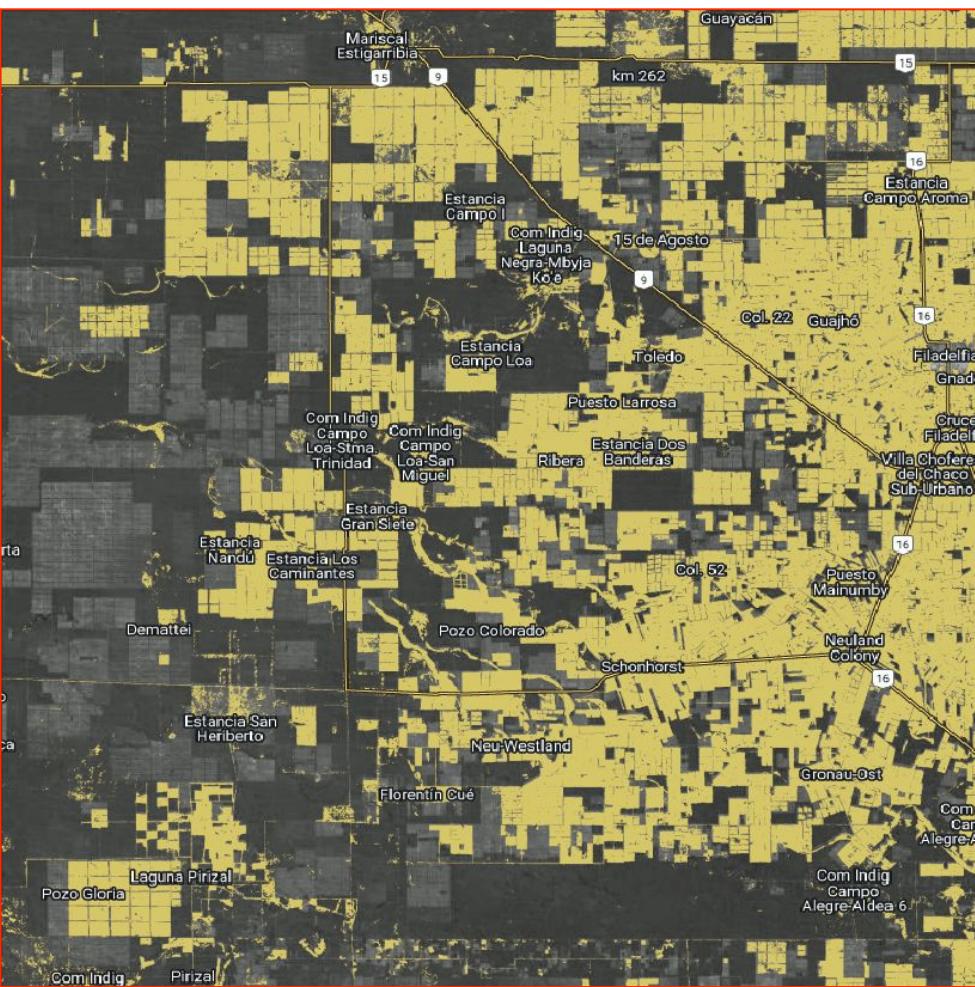


Dominant grassland - 2022

- Cultivated
- Natural / Semi-natural

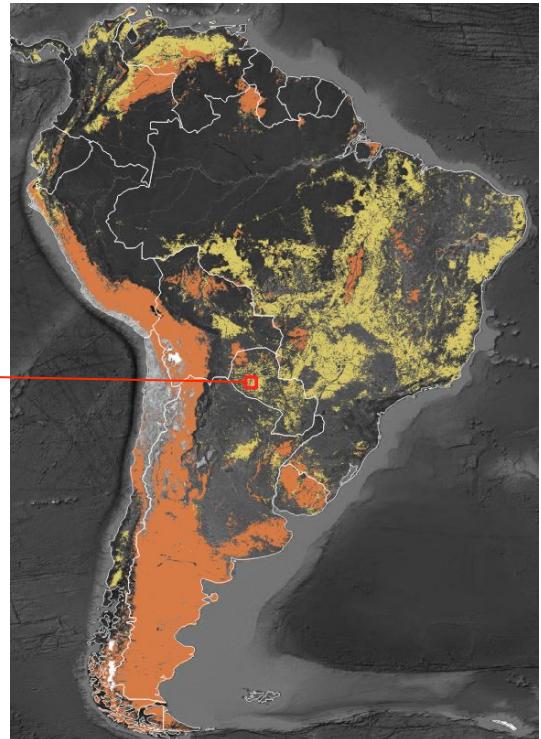


Dominant grassland (2000–2022)



Cultivated

Natural / Semi-natural



Filadelfia, Paraguay

Mapping global grassland dynamics 2000—2022 at 30m spatial resolution using spatiotemporal Machine Learning

Leandro Parente^{1*}, Lindsey Sloat², Vinicius Mesquita³, Davide Consoli¹, Radost Stanimirova², Tomislav Hengl¹, Carmelo Bonannella^{1,4}, Nathália Teles⁵, Ichsaní Wheeler¹, Steffen Ehrmann^{5,6,8}, Maria Hunter³, Laerte Ferreira³, Ana Paula Mattos³, Bernard Oliveira³, Carsten Meyer^{6,7,8}, Murat Şahin¹, Martijn Witjes^{1,4}, Steffen Fritz⁵, Ziga Malek⁵, and Fred Stolle²

¹OpenGeoHub Foundation, Doorwerth, The Netherlands

²Land & Carbon Lab, World Resources Institute, Washington DC, USA

³Remote Sensing and GIS Laboratory (LAPIG/UFG), Goiânia, Brazil

⁴Laboratory of Geo-Information Science and Remote Sensing, Wageningen University & Research, Wageningen, The Netherlands

⁵International Institute for Applied Systems Analysis (IIASA), Laxenburg, Austria

⁶German Centre for Integrative Biodiversity Research (iDiv) Halle-Jena-Leipzig, Leipzig, Germany

⁷Institute of Geosciences and Geography, Martin Luther University Halle-Wittenberg, Halle (Saale), Germany

⁸Institute of Biology, Leipzig University, Leipzig, Germany

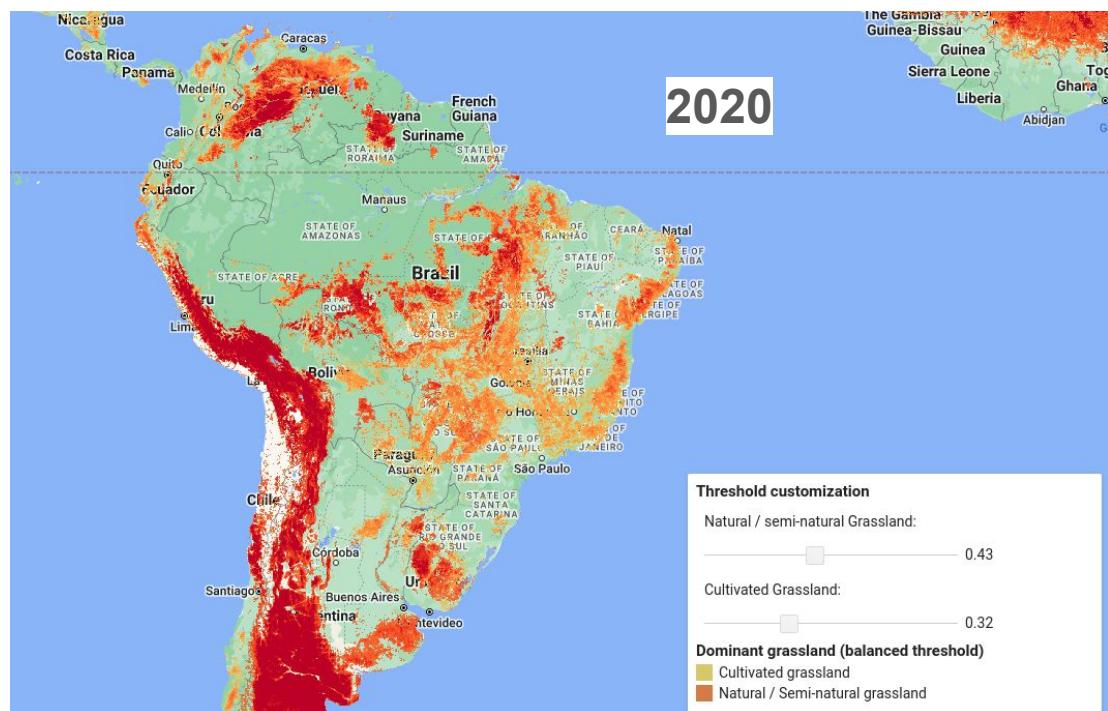
*corresponding author(s): Leandro Parente (leandro.parente@opengeohub.org)

Global maps will be available though:

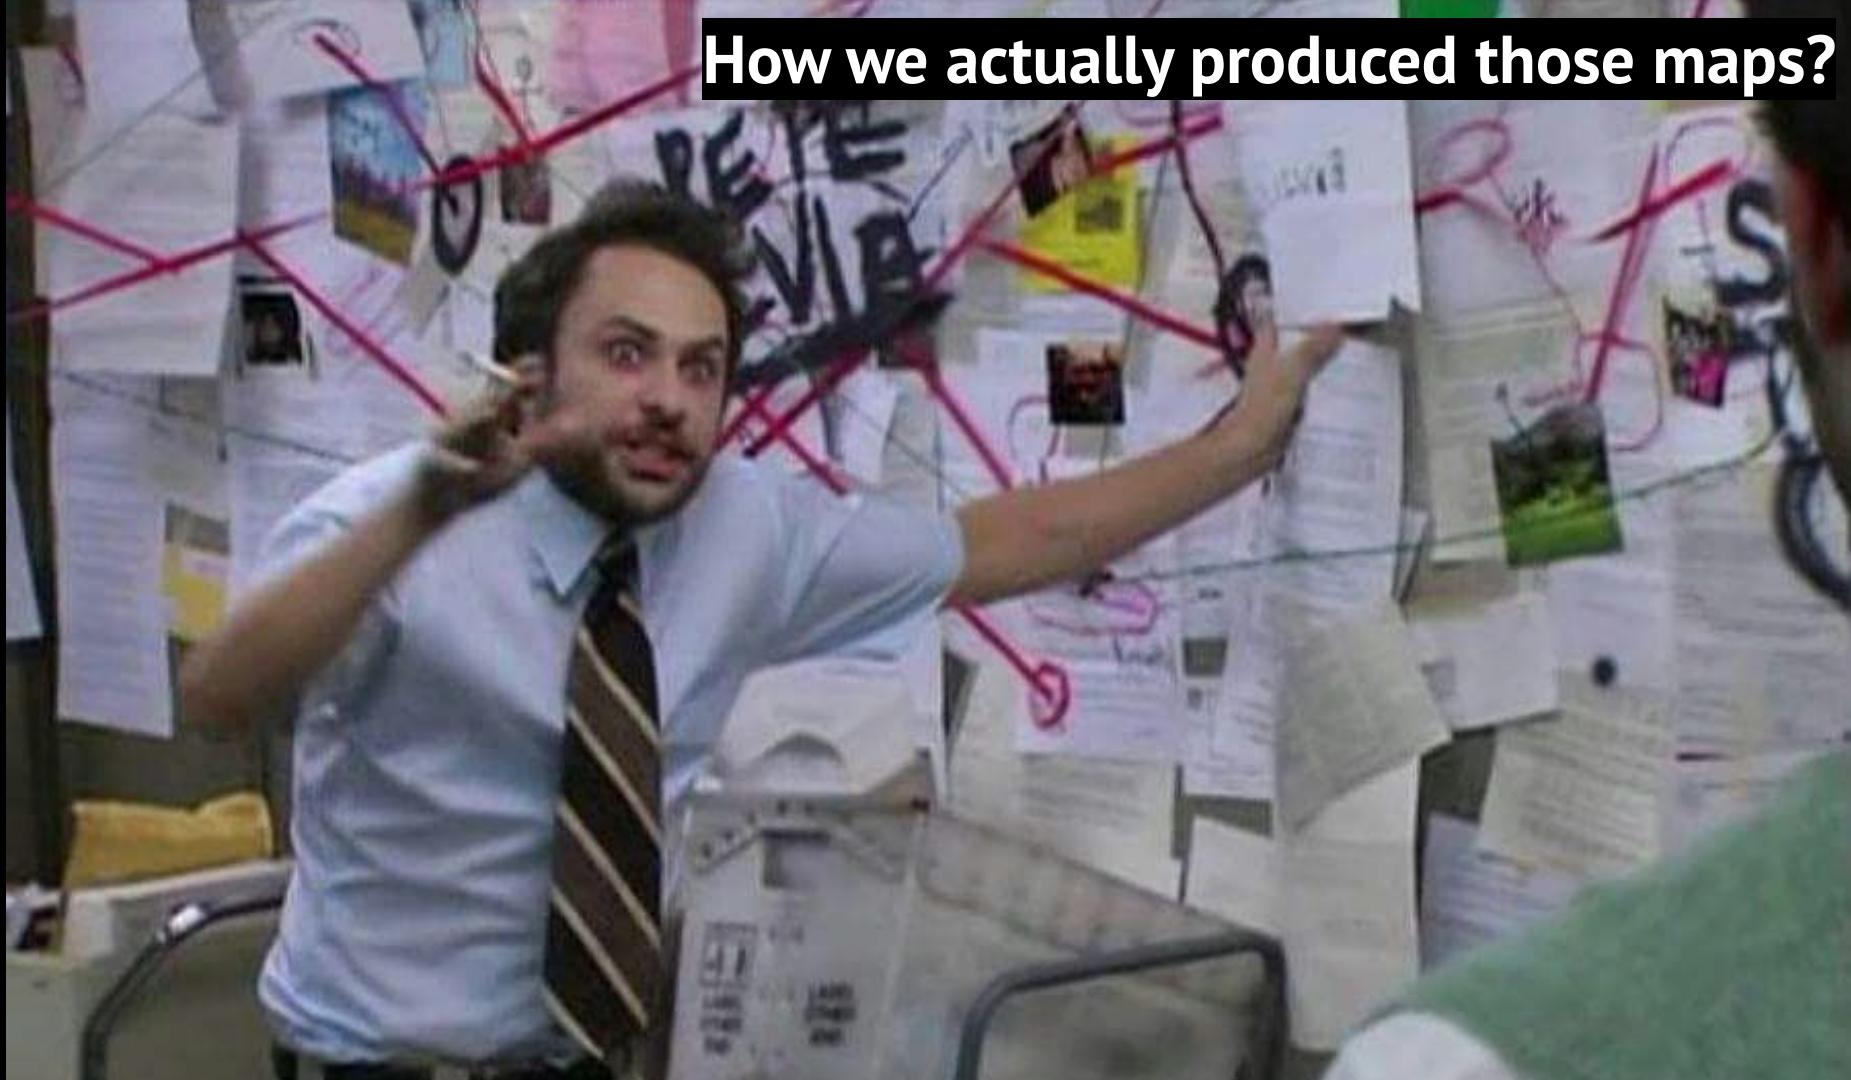
- Google Earth Engine
- Cloud-optimized GeoTIFF (COG)
- SpatioTemporal Asset Catalog (STAC)

<https://doi.org/10.21203/rs.3.rs-4514820/v1>

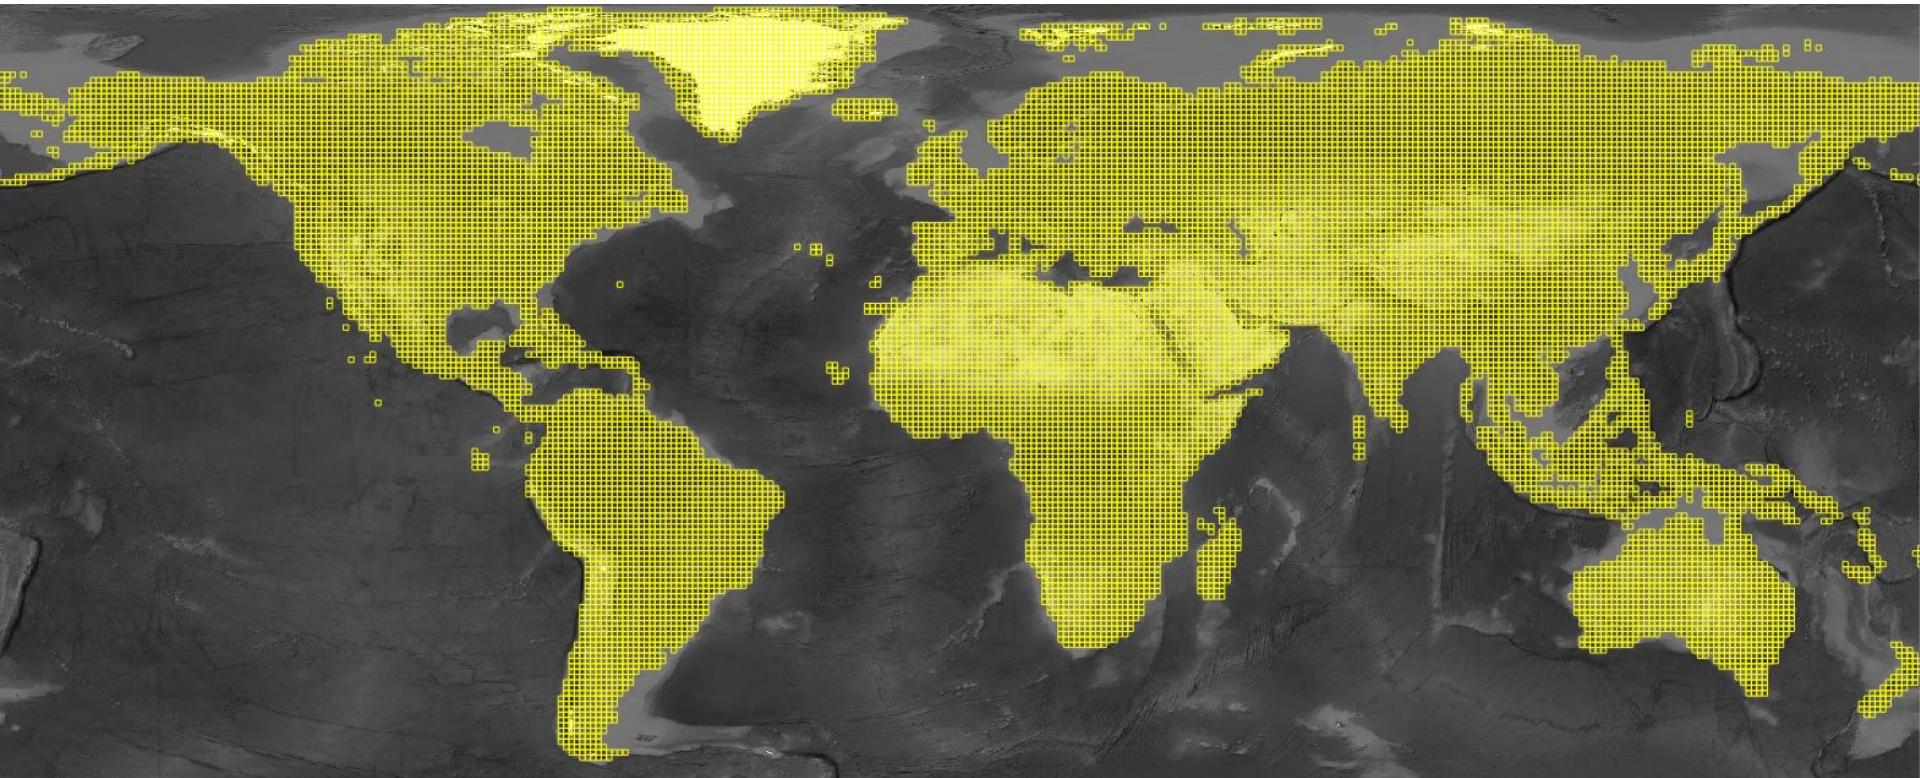
https://leandro_opengeohub.users.earthengine.app/view/gpw--global-grassland-2020



How we actually produced those maps?



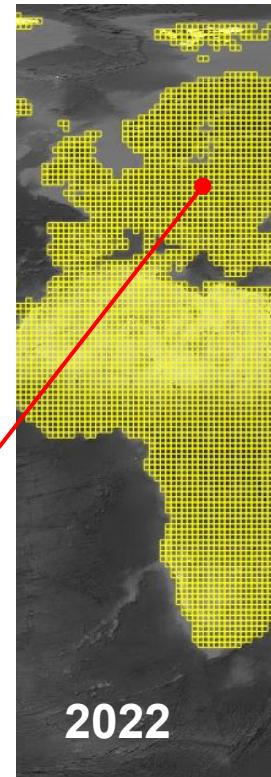
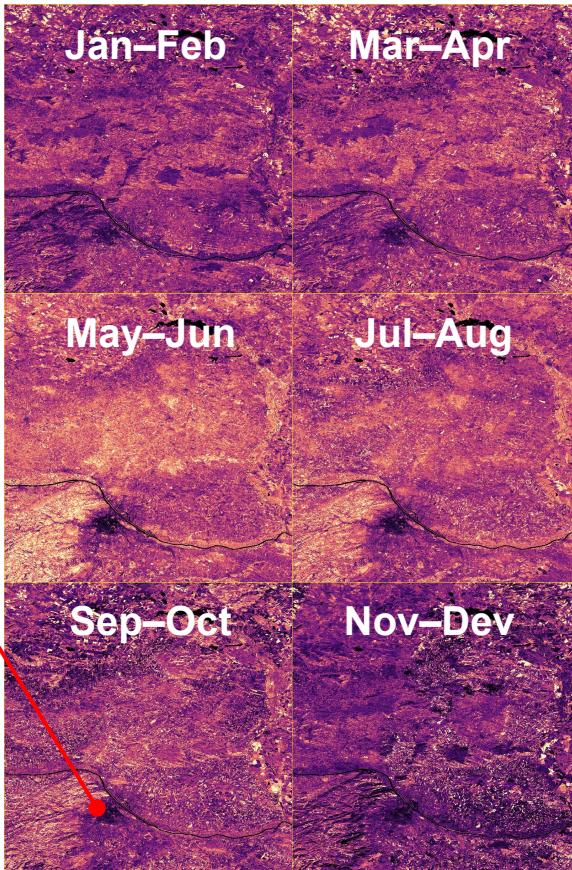
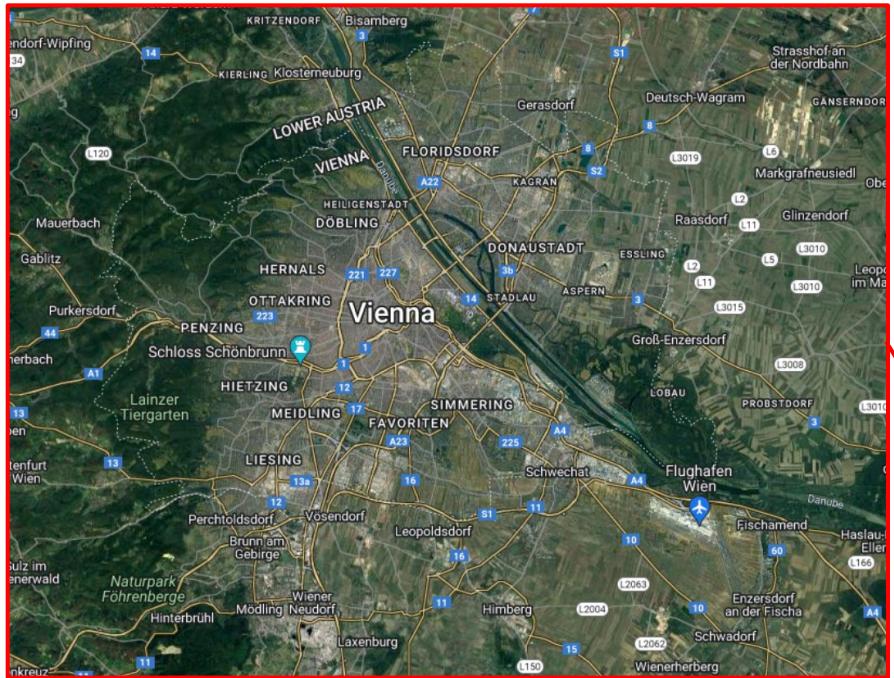
First, let's split the world in multiple tiles



18,667 Tiles (4004 x 4004 pixels)

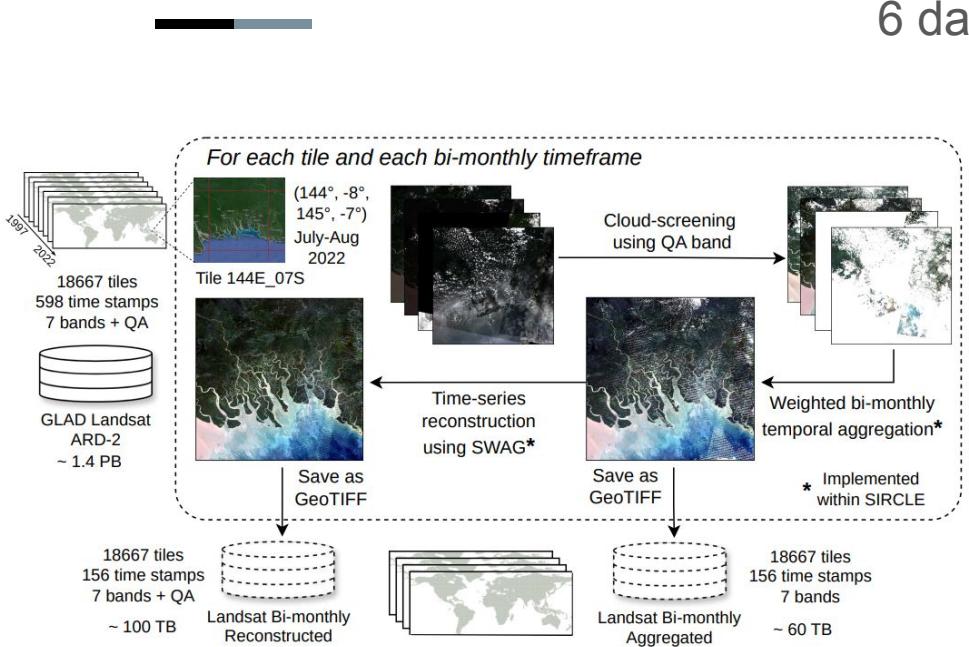
Example tile covering Viena (016E_48N)

6 dates *Near-infrared (NIR)*

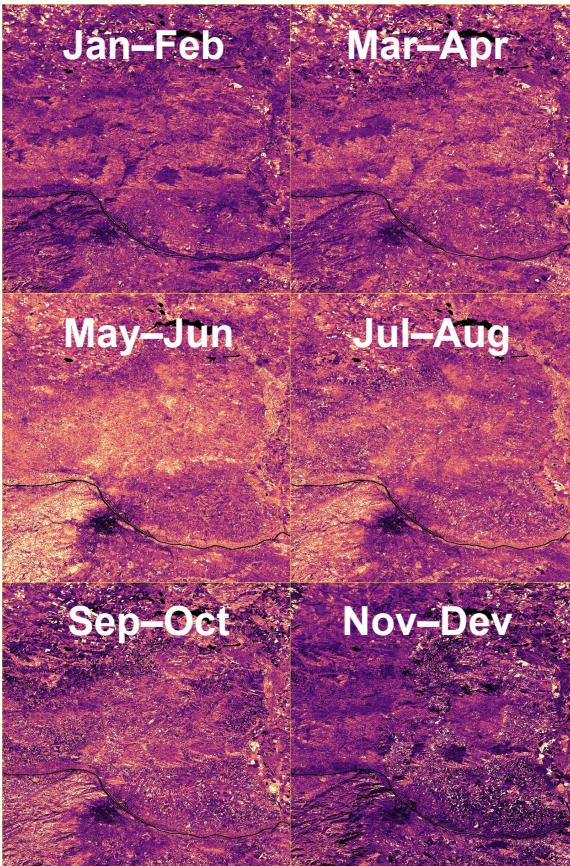


4004 x 4004
pixels

Example tile covering Vienna (016E_48N)



6 dates *Near-infrared (NIR)*

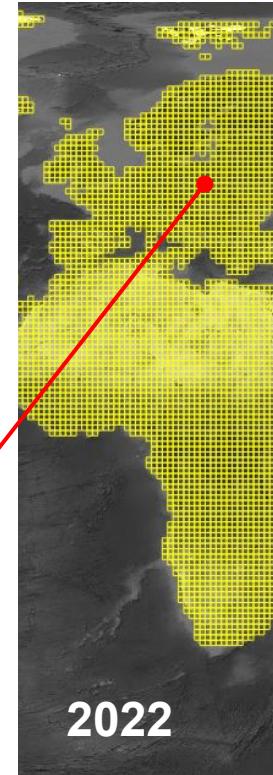
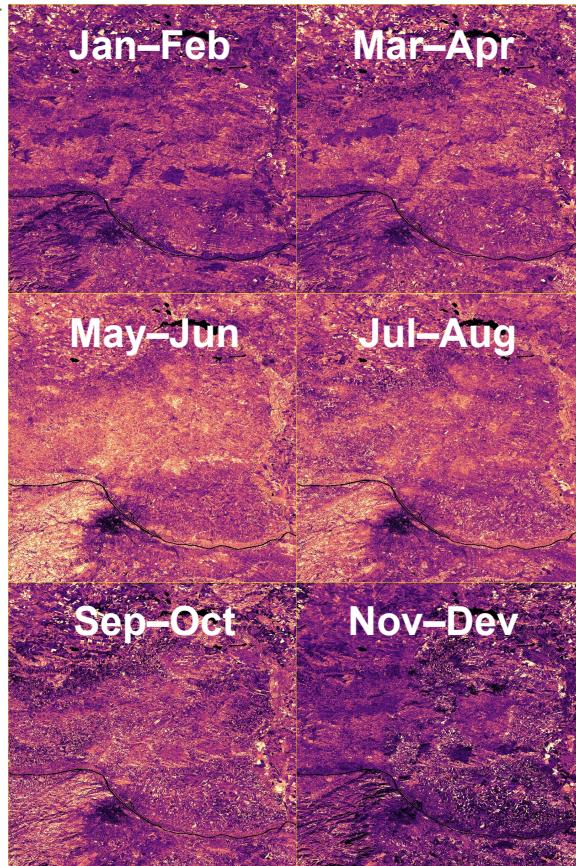
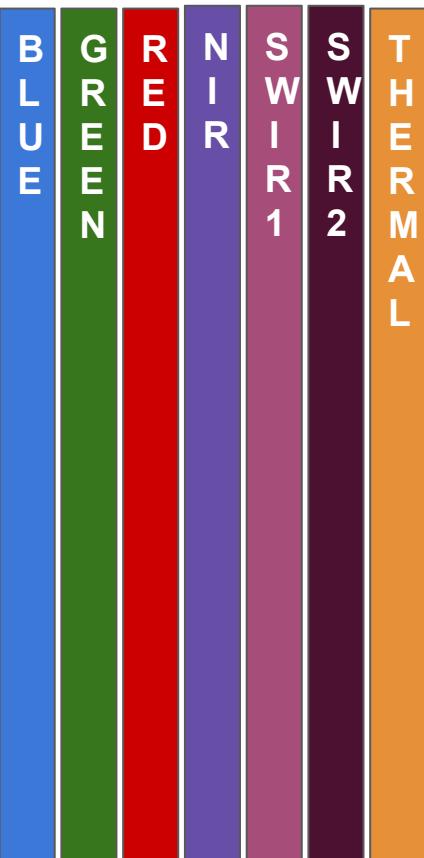


<https://doi.org/10.21203/rs.3.rs-4465582/v1>

4004 x 4004
pixels

Temporal layers - Landsat ARD

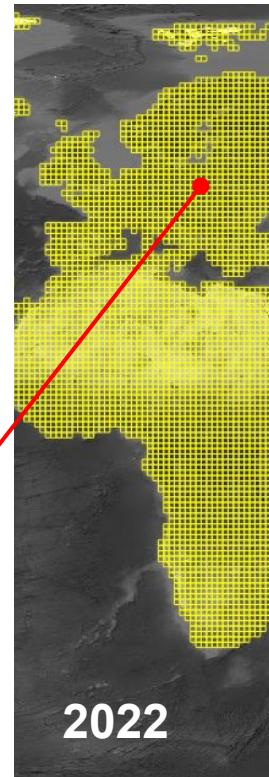
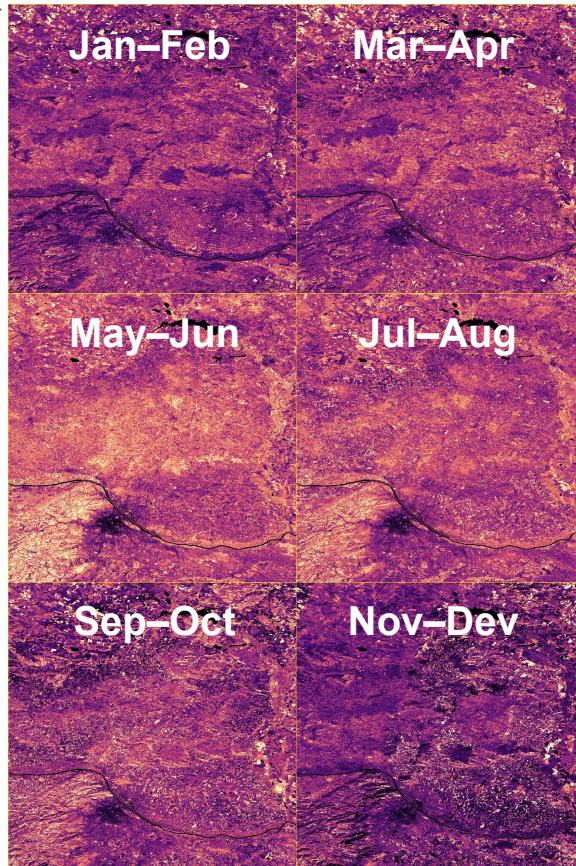
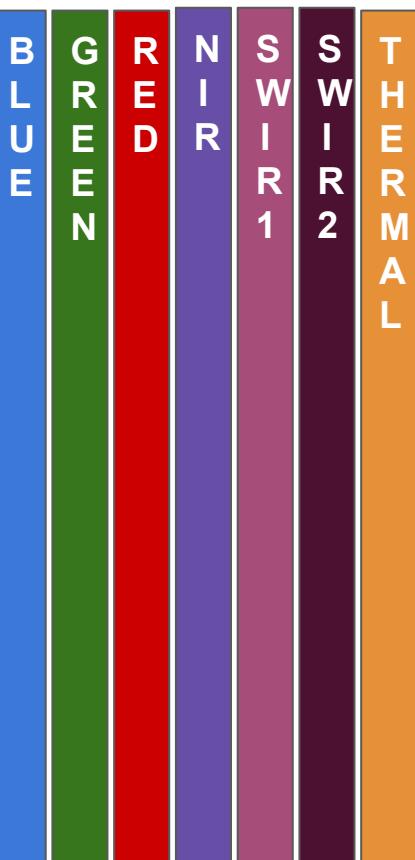
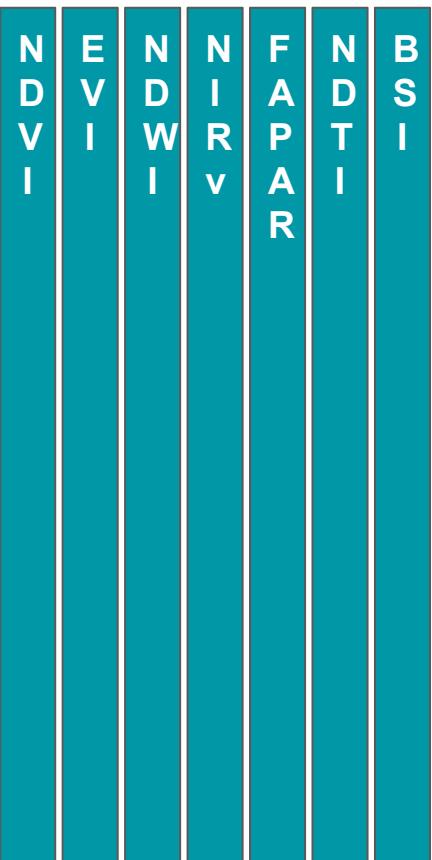
7 bands x 6 dates



4004 x 4004
pixels

Temporal layers - Landsat ARD

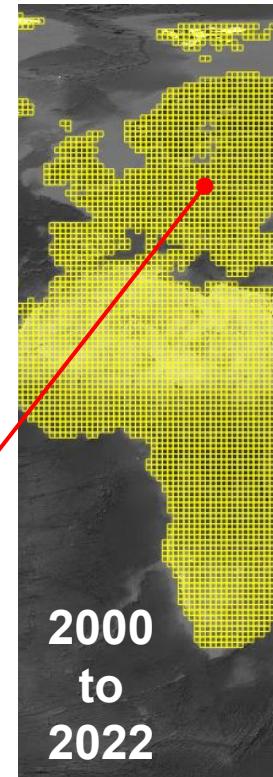
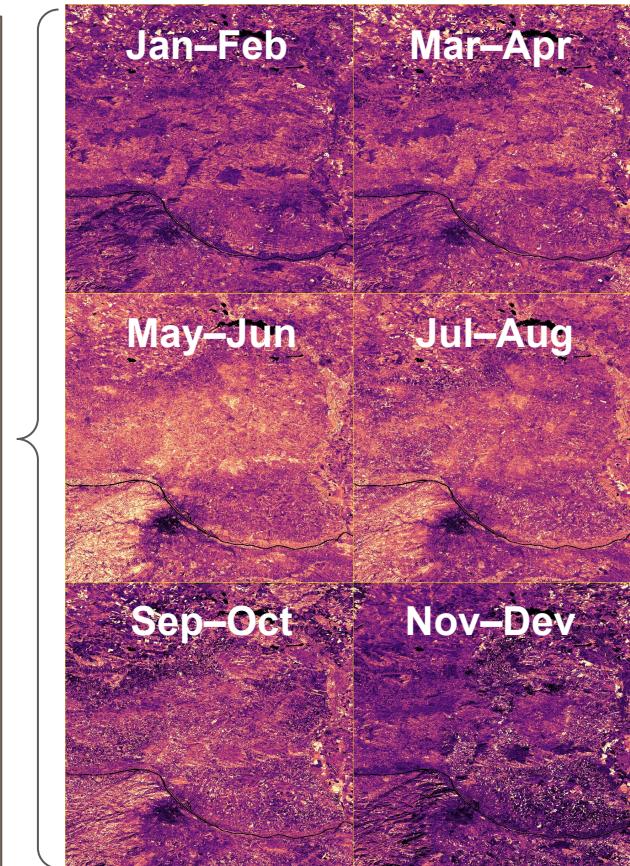
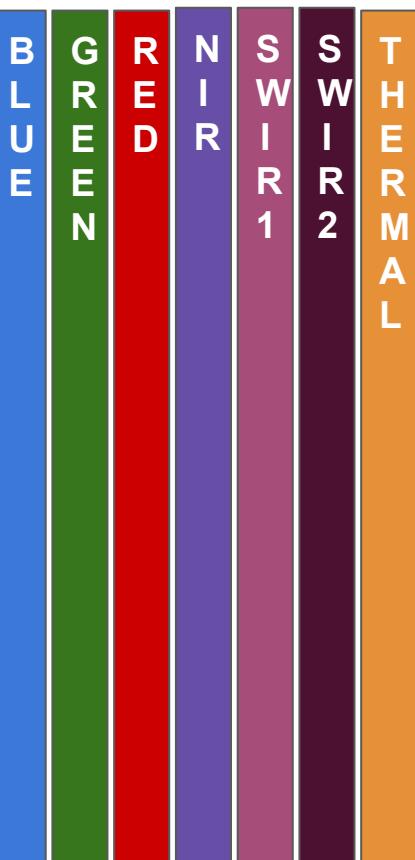
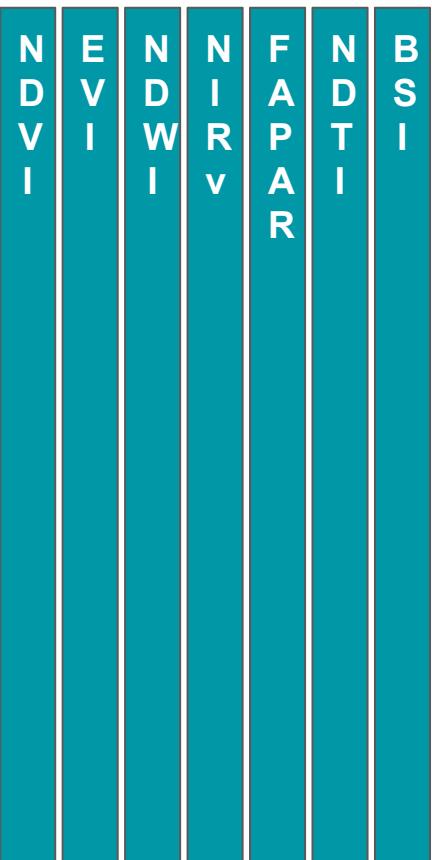
(7 indices + 7 bands) x 6 dates



4004 x 4004
pixels

Temporal layers - Landsat ARD

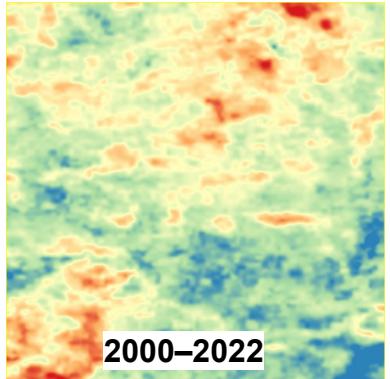
23 years x (7 indices + 7 bands) x 6 dates x 16M pixels



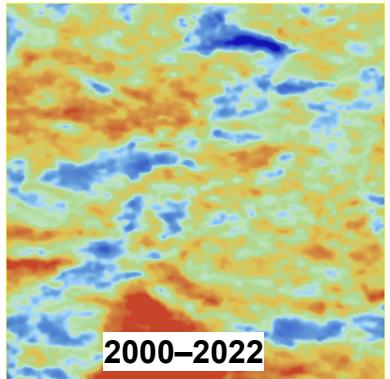
4004 x 4004
pixels

Static layers - MODIS variable, terrain, distance maps

85 static layers x 16M pixels

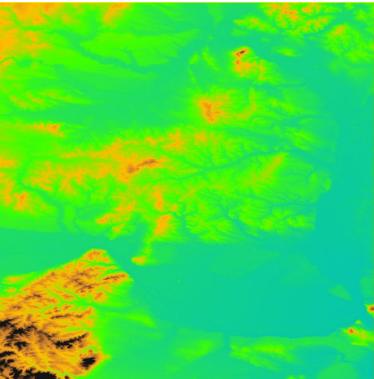


24 long-term
water vapour
($p50$ & std)

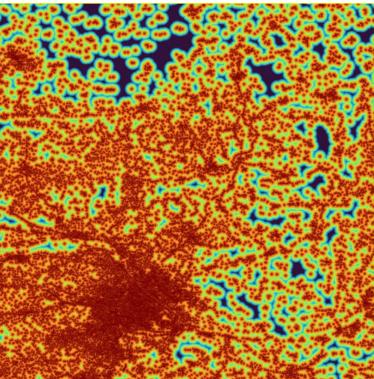


24 long-term
daytime temperature
($p50$ & std)

24 long-term nighttime
temperature ($p50$ &
 std)



2 DTM
static layers

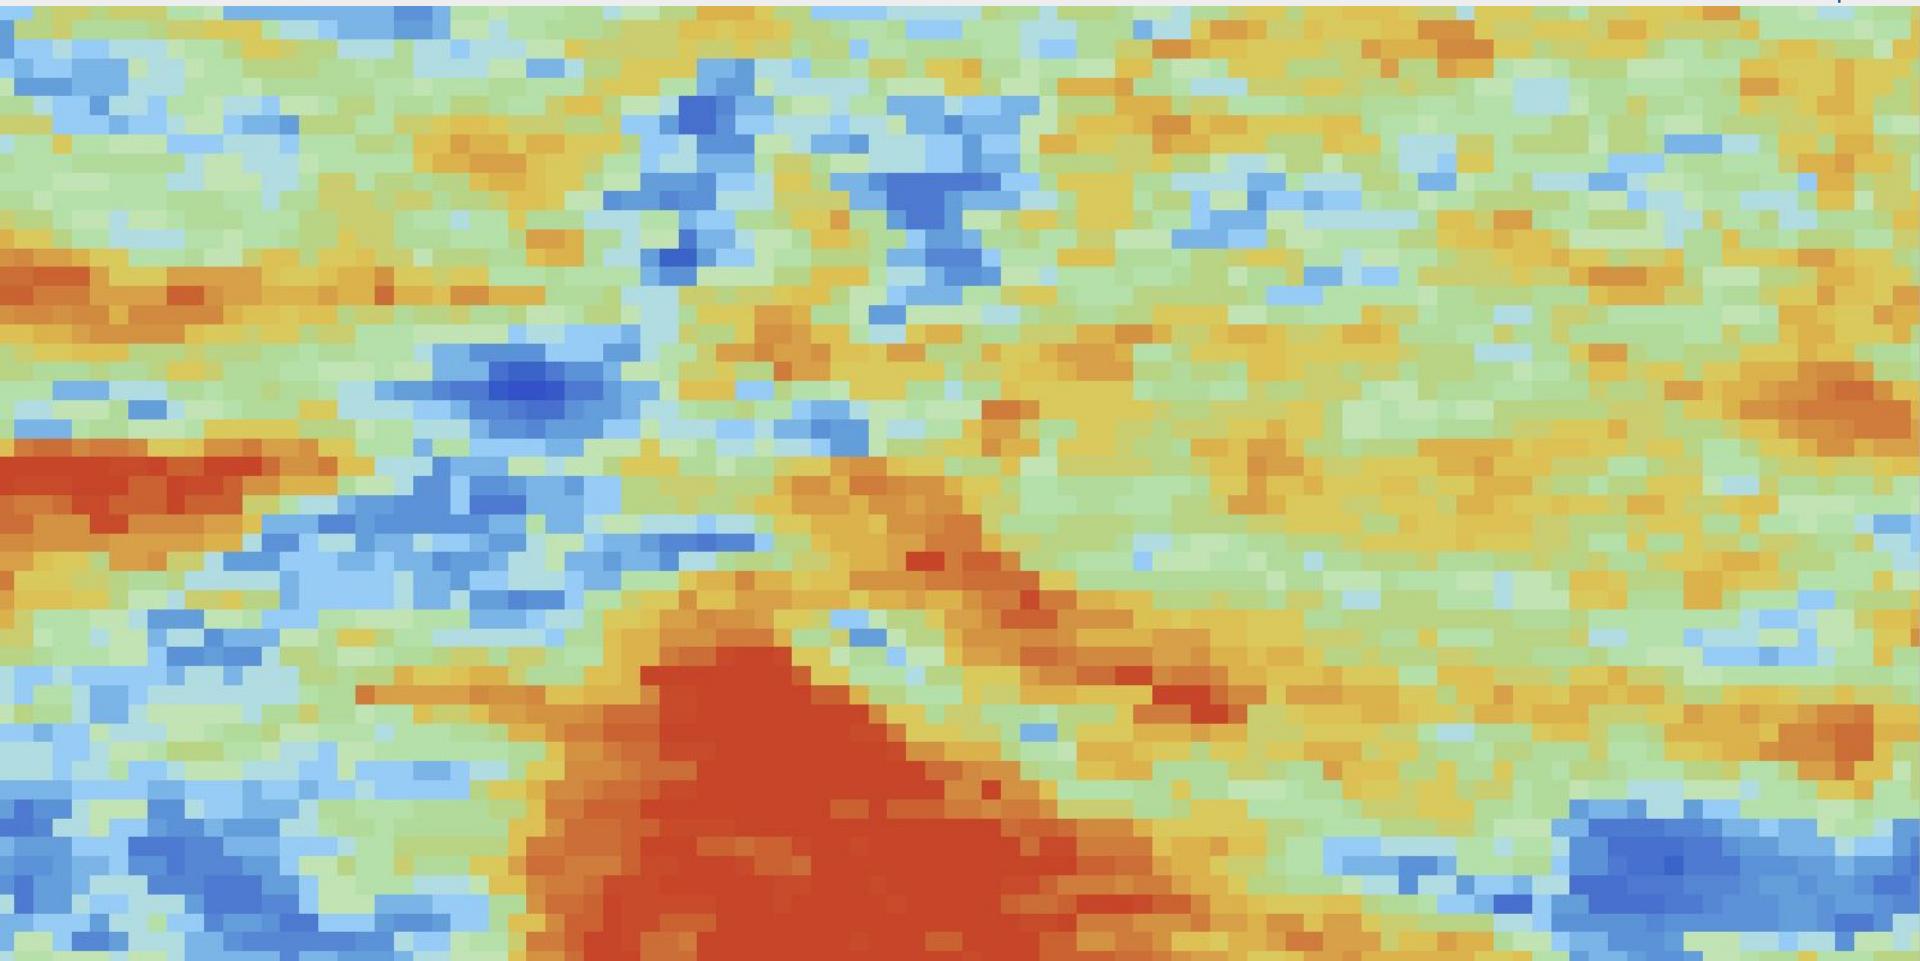


11 Static
distance maps

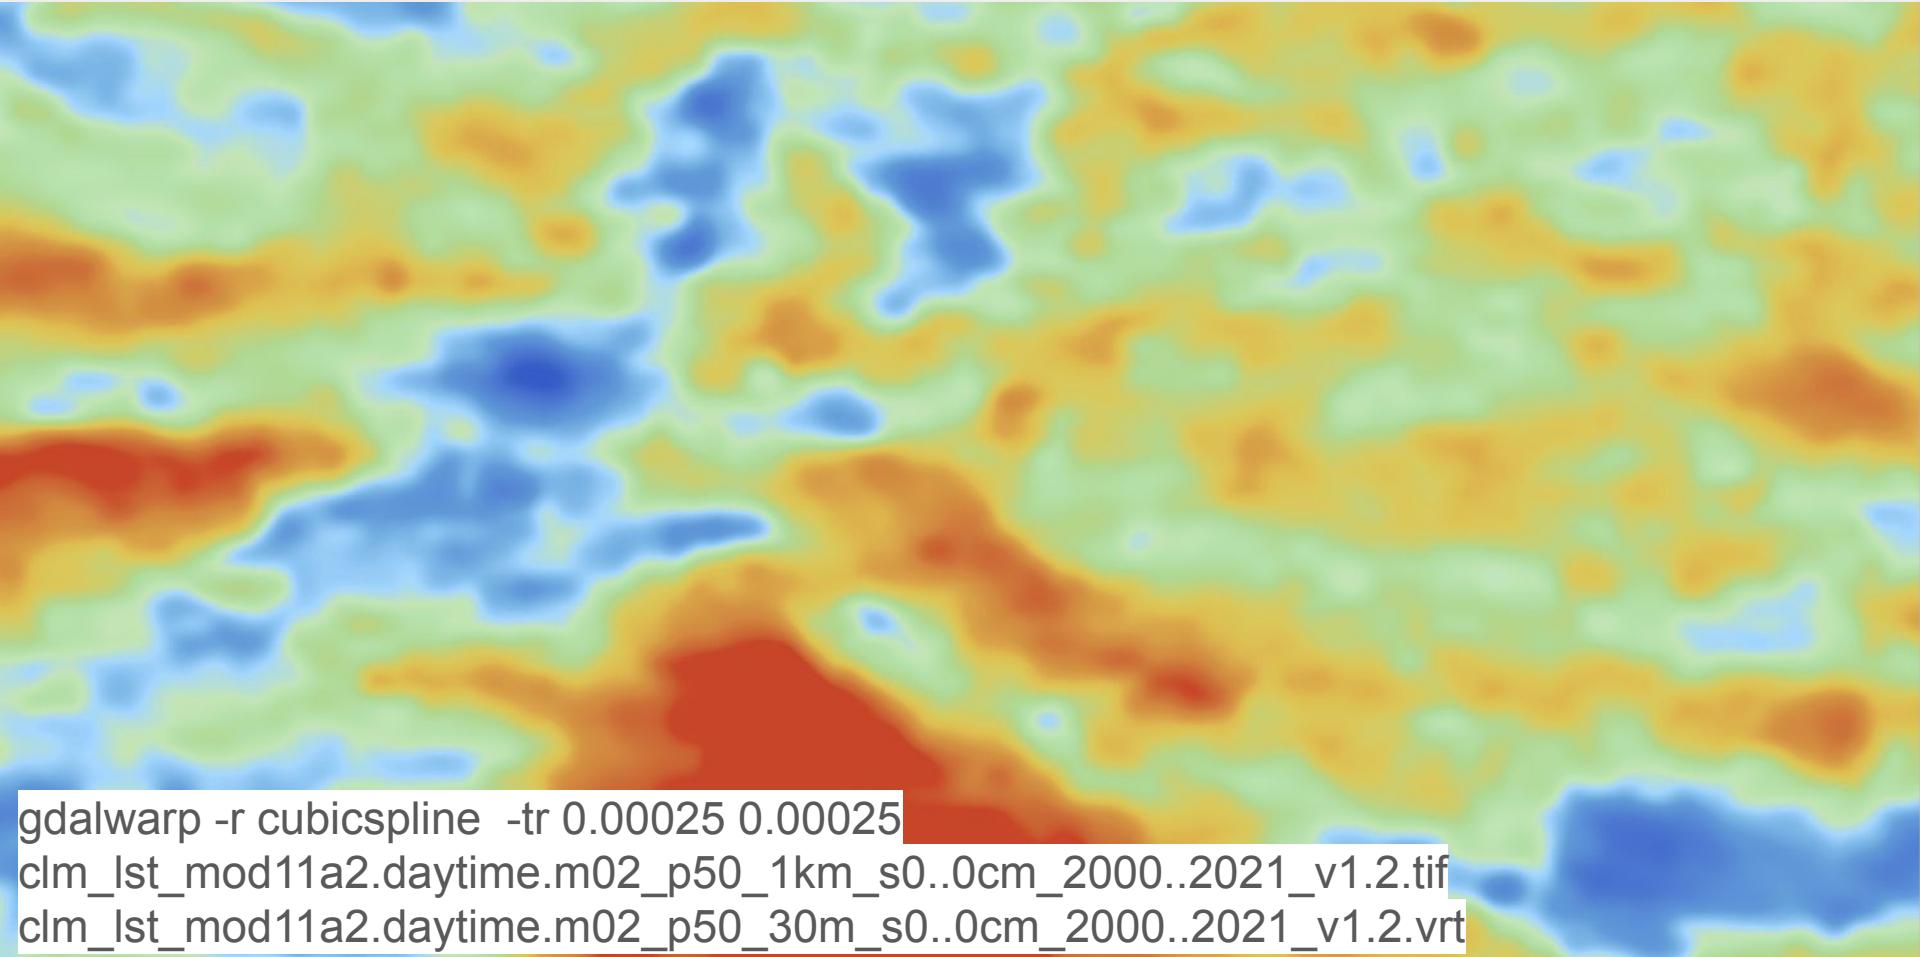


4004 x 4004
pixels

Downscale Trick (1km -> 30m)



Downscale Trick (1km -> 30m)



```
gdalwarp -r cubicspline -tr 0.00025 0.00025  
clm_lst_mod11a2.daytime.m02_p50_1km_s0..0cm_2000..2021_v1.2.tif  
clm_lst_mod11a2.daytime.m02_p50_30m_s0..0cm_2000..2021_v1.2.vrt
```

Reading all layers & years

88 static layers

6.34 secs

966 temporal layers
(Landsat bands - all years)

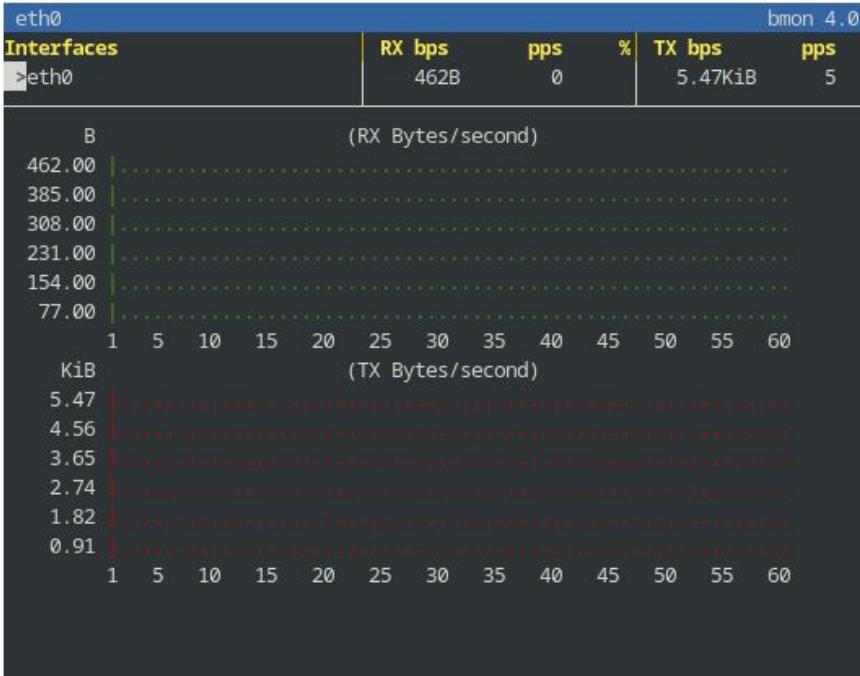
8.17 secs

1541 on-the-fly layers
(Landsat indices - all years)

25.82 secs

70 Gb RAM

(2595, 16032016)



On-the-fly layer calculations

88 static layers

6.34 secs

966 temporal layers
(Landsat bands - all years)

8.17 secs

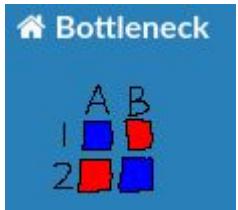
1541 on-the-fly layers
(Landsat indices - all years)

25.82 secs

(2595, 16032016)

On-the-fly layer calculations

Python libraries to optimize
Numpy Computation



<https://bottleneck.readthedocs.io>

NumExpr <https://numexpr.readthedocs.io>

 Numba <https://numba.pydata.org>



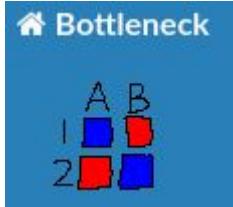
[High performance computing in Python](#)

colab

 GitLab

Computational notebooks

On-the-fly layer calculations



NumExpr

```
import numexpr as ne

expression = '( (nir - red) / (nir + red) ) * nir'
params = { 'nir': nir_data, 'red': red_data }

ne.evaluate(expression, params,
            optimization='aggressive')
```

drop-in replacement

```
import bottleneck as bn
bn.nanmedian(red_data, axis=(2))
```

```
from numba import prange, jit

last_dim = red_data.shape[-1]
out_shape = red_data.shape[:-1]

@jit(parallel=True)
def _reduce(data):
    x = np.zeros(len(data))
    for i in prange(data.shape[0]):
        x[i] = np.nanmedian(data[i])
    return x
```

15x speed up (array reduction)

| | library | average | std |
|---|------------|----------|----------|
| 2 | Numba | 0.130819 | 0.002264 |
| 1 | Bootleneck | 0.367630 | 0.004450 |
| 0 | Numpy | 2.060645 | 0.203988 |

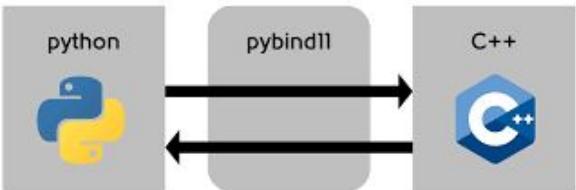
4.5x speed up (mathematical operation)

| | library | average | std |
|---|---------|----------|----------|
| 1 | Numexpr | 0.020733 | 0.000450 |
| 2 | Numba | 0.043981 | 0.000305 |
| 0 | Numpy | 0.092036 | 0.002885 |

On-the-fly layer calculations

```
band_scaling = 0.004  
result_scaling = 125.  
result_offset = 125.
```

```
skmap_bindings.computeNormalizedDifference(data, n_threads,  
    swir1_idx, swir2_idx, ndti_idx,  
    band_scaling, band_scaling,  
    result_scaling, result_offset, [0., 250.])
```

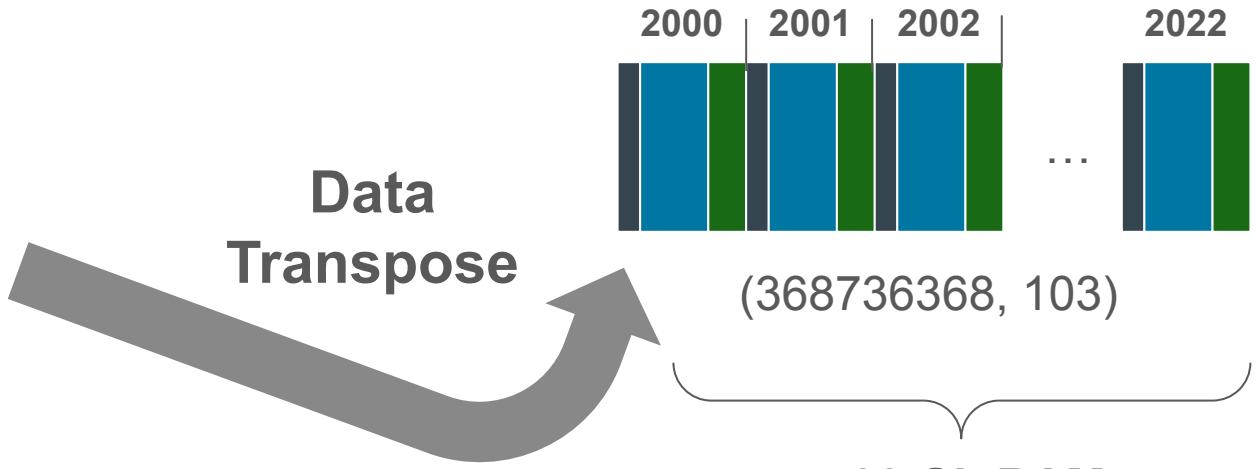


Eigen is a C++ template library for linear algebra: matrices, vectors, numerical solvers, and related algorithms.

```
void TransArray::computeNormalizedDifference(std::vector<uint_t> positive_indices,  
    std::vector<uint_t> negative_indices,  
    std::vector<uint_t> result_indices,  
    float_t positive_scaling,  
    float_t negative_scaling,  
    float_t result_scaling,  
    float_t result_offset,  
    std::vector<float_t> clip_value)  
{  
    skmapAssertIfTrue((positive_indices.size() != negative_indices.size()) || (positive_indices.size() != result_indices.size()),  
        "scikit-map ERROR 3: positive_i, negative_i, result_i must be of the same size");  
    auto computeNormalizedDifferenceRow = [&] (uint_t i, Eigen::Ref<MatFloat::RowXpr> row)  
    {  
        uint_t positive_i = positive_indices[i];  
        uint_t negative_i = negative_indices[i];  
        row = (m_data.row(positive_i) * positive_scaling - m_data.row(negative_i) * negative_scaling).array() /  
            (m_data.row(positive_i) * positive_scaling + m_data.row(negative_i) * negative_scaling).array() *  
            result_scaling + result_offset;  
        row = (row.array()).round();  
        row = (row.array() == -inf_v).select(-result_scaling + result_offset, row);  
        row = (row.array() == +inf_v).select(result_scaling + result_offset, row);  
        row = (row.array() < clip_value[0]).select(clip_value[0], row);  
        row = (row.array() > clip_value[1]).select(clip_value[1], row);  
    };  
    this->parRowPerm(computeNormalizedDifferenceRow, result_indices);  
}
```

~6x speed up

Data transposition



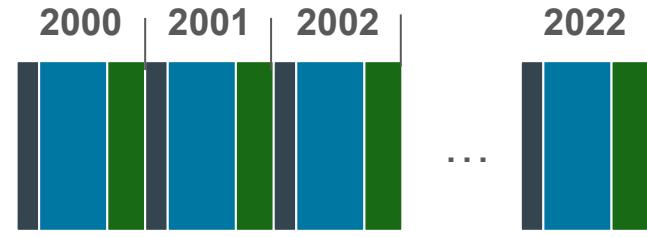
| | | | | | | | |
|----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 [0.0%] | 6 [0.0%] | 12 [0.0%] | 18 [0.0%] | 24 [0.0%] | 30 [0.0%] | 36 [0.0%] | 42 [0.0%] |
| 1 [0.0%] | 7 [0.0%] | 13 [0.0%] | 19 [0.7%] | 25 [0.0%] | 31 [0.0%] | 37 [0.0%] | 43 [0.0%] |
| 2 [0.0%] | 8 [6.0%] | 14 [0.0%] | 20 [0.0%] | 26 [0.0%] | 32 [0.0%] | 38 [0.0%] | 44 [0.0%] |
| 3 [0.0%] | 9 [0.0%] | 15 [0.0%] | 21 [0.0%] | 27 [0.0%] | 33 [0.0%] | 39 [0.0%] | 45 [0.0%] |
| 4 [0.7%] | 10 [0.0%] | 16 [0.0%] | 22 [0.0%] | 28 [0.0%] | 34 [0.0%] | 40 [0.0%] | 46 [0.0%] |
| 5 [0.0%] | 11 [0.0%] | 17 [0.0%] | 23 [0.0%] | 29 [0.0%] | 35 [0.0%] | 41 [0.0%] | 47 [0.0%] |
| Mem [| | | | | | | |
| Swp [| | | | | | | |

(2595, 16032016)

583G/1007G

0K/200G

Data transposition



(368736368, 103)

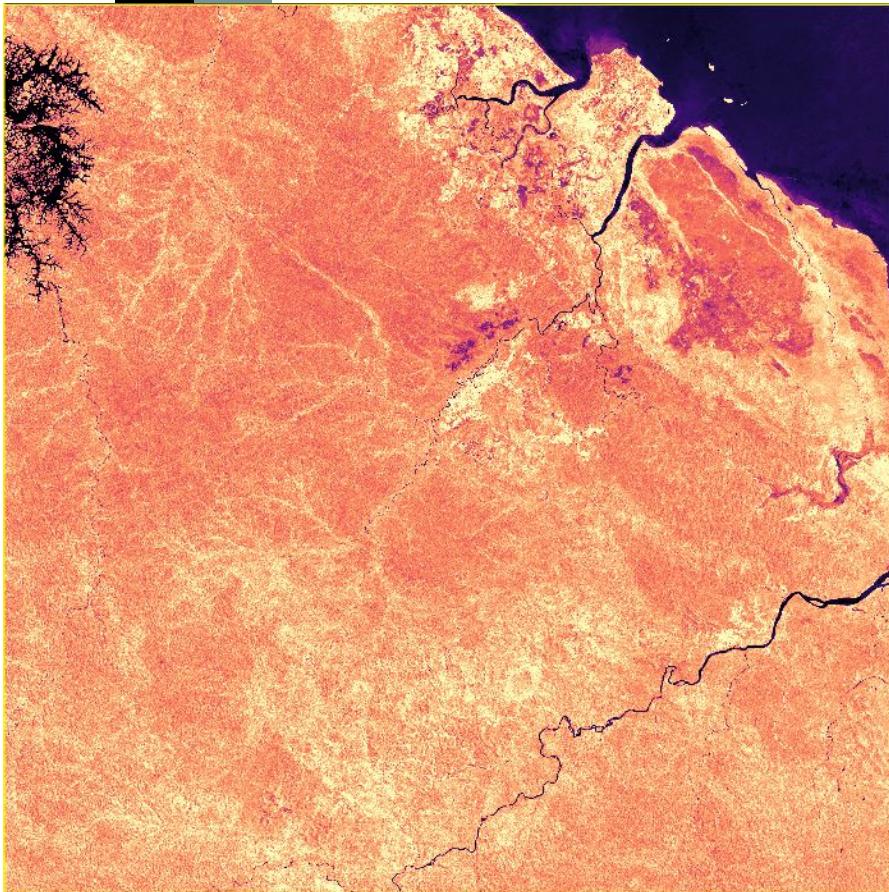
÷ 23

(16032016, 103)

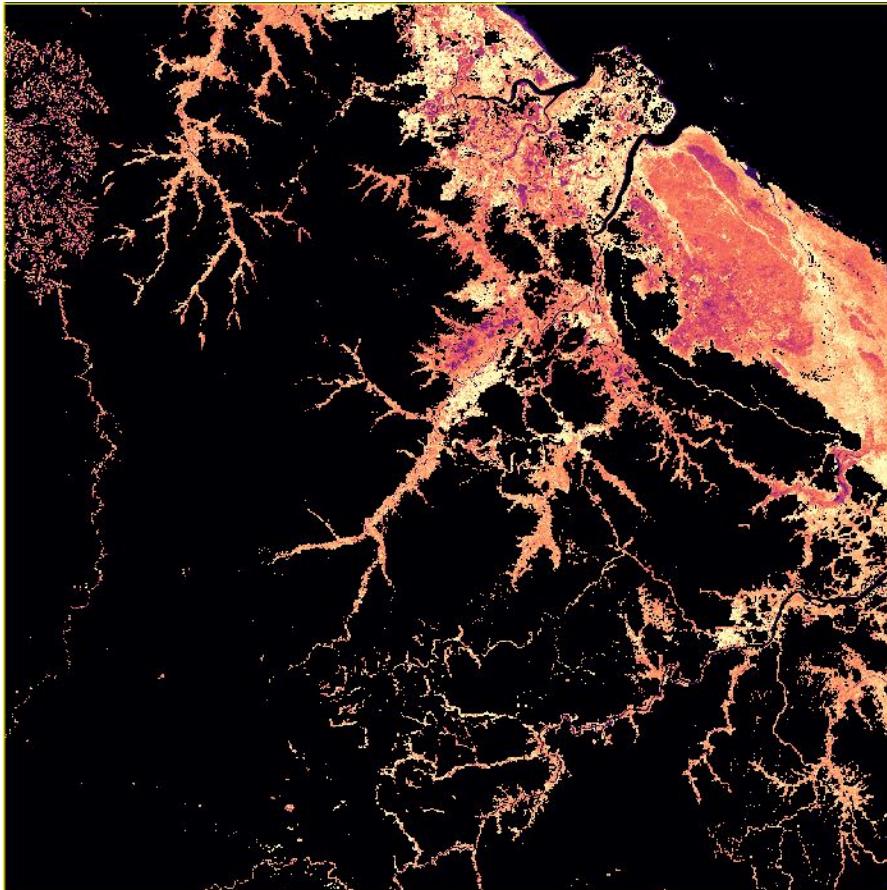
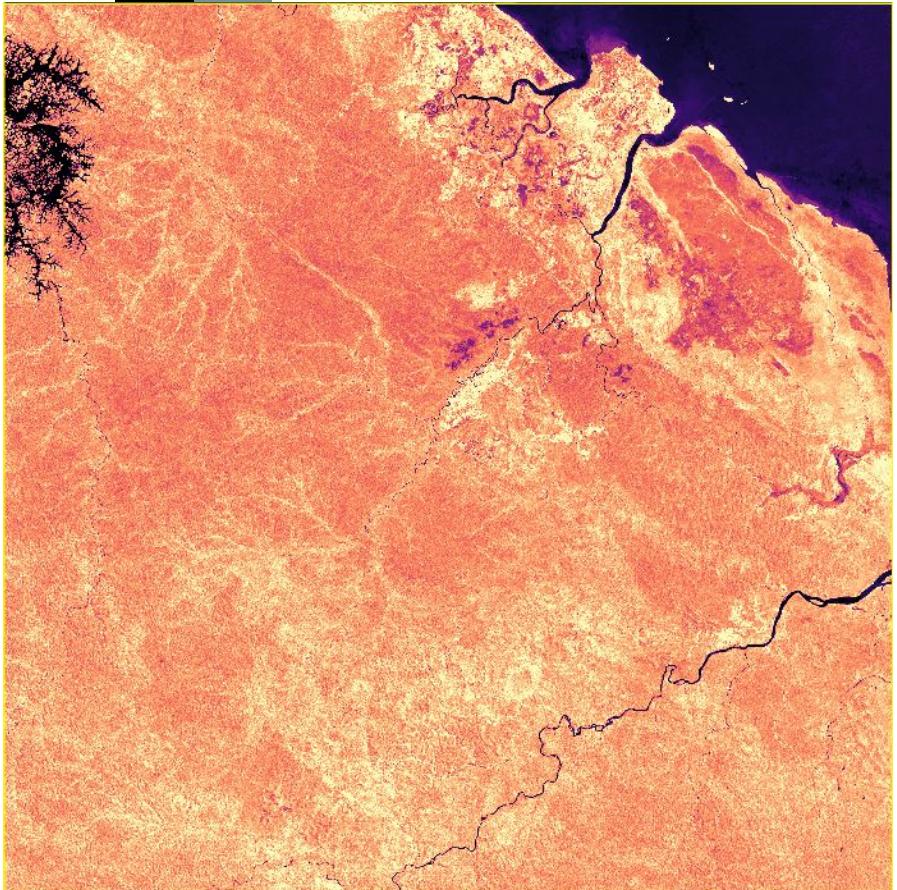
(2595, 16032016)



Global prediction mask

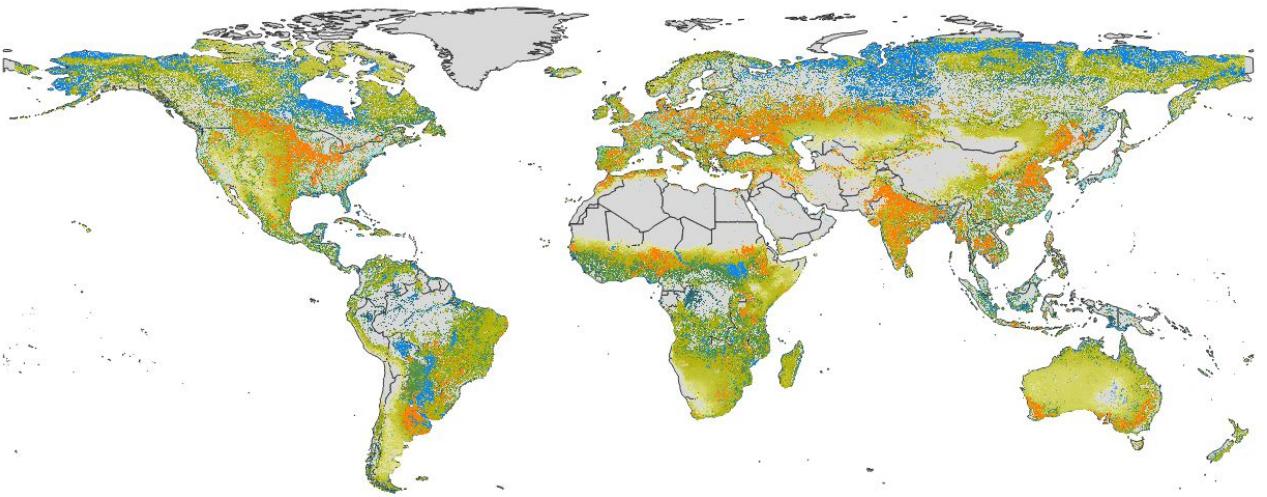


Global prediction mask



Global prediction mask

Based on [UMD GLAD annual land cover and land use change](#)

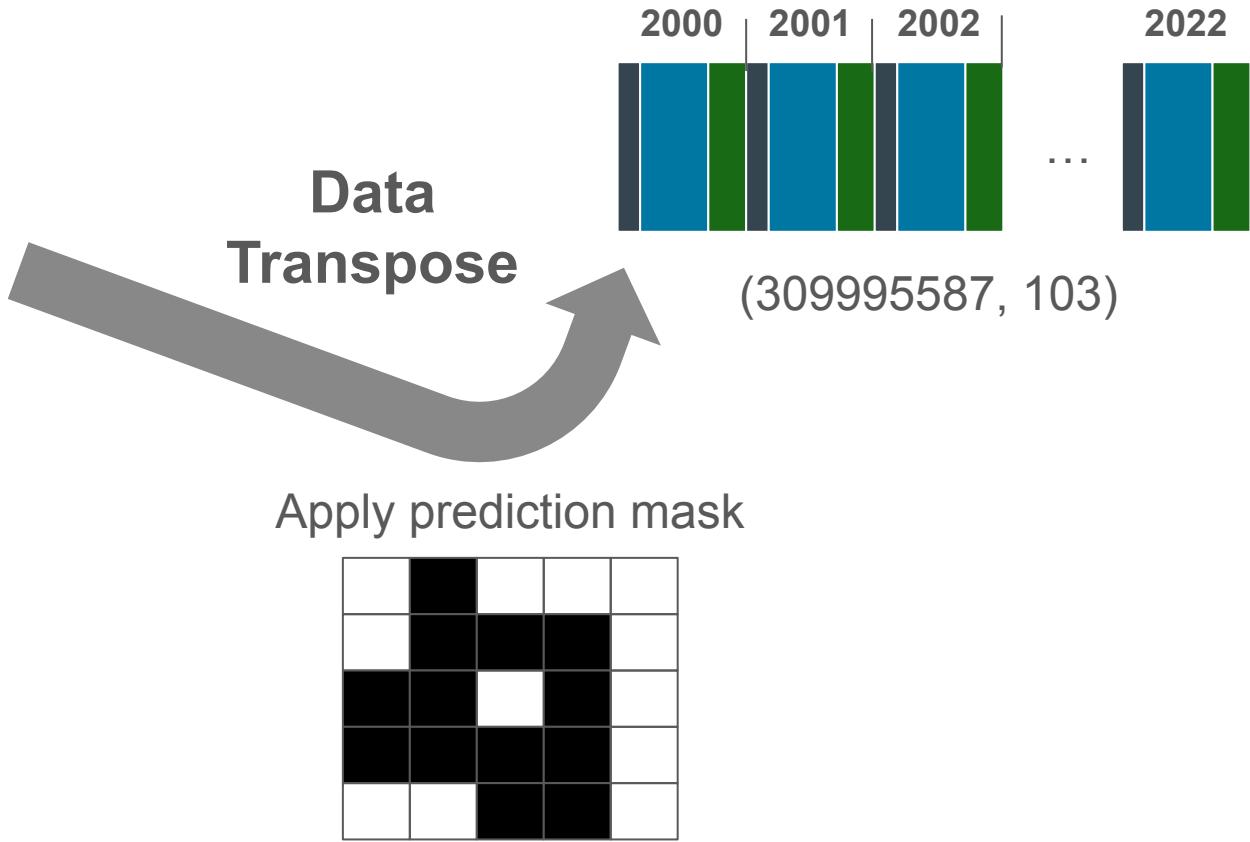


Excluded classes they appear in all maps (2000 –2020)

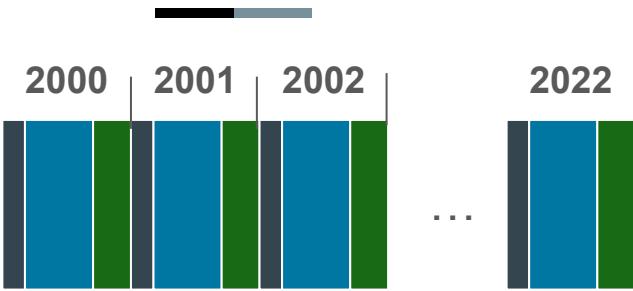
- Terra Firma, true desert - 3% short veg. cover (0)
- Terra Firma, true desert - 7% short veg. cover (1)
- Terra Firma, semi-arid - 11% short veg. cover (2)
- Terra Firma, semi-arid - 15% short veg. cover (3)
- Terra Firma, semi-arid - 19% short veg. cover (4)
- Terra Firma, semi-arid - 23% short veg. cover (5)

- Wetland, salt pan - 3% short veg. cover (100)
- Wetland, salt pan - 7% short veg. cover (101)
- Terra Firma, stable tree cover - 10m trees (32)
- Terra Firma, stable tree cover - 11m trees (33)
- Terra Firma, stable tree cover - 12m trees (34)
- Terra Firma, stable tree cover - 13m trees (35)
- Terra Firma, stable tree cover - 14m trees (36)
- Terra Firma, stable tree cover - 15m trees (37)
- Terra Firma, stable tree cover - 16m trees (38)
- Terra Firma, stable tree cover - 17m trees (39)
- Terra Firma, stable tree cover - 18m trees (40)
- Terra Firma, stable tree cover - 19m trees (41)
- Terra Firma, stable tree cover - 20m trees (42)
- Terra Firma, stable tree cover - 21m trees (43)
- Terra Firma, stable tree cover - 22m trees (44)
- Terra Firma, stable tree cover - 23m trees (45)
- Terra Firma, stable tree cover - 24m trees (46)
- Terra Firma, stable tree cover - 25m trees (47)
- Terra Firma, stable tree cover - >25m trees (48)
- Snow/ice, stable (241)
- Ocean (254)
- No data (255)

Global prediction mask

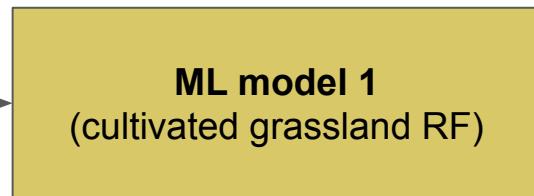


(Finally) Global ML prediction



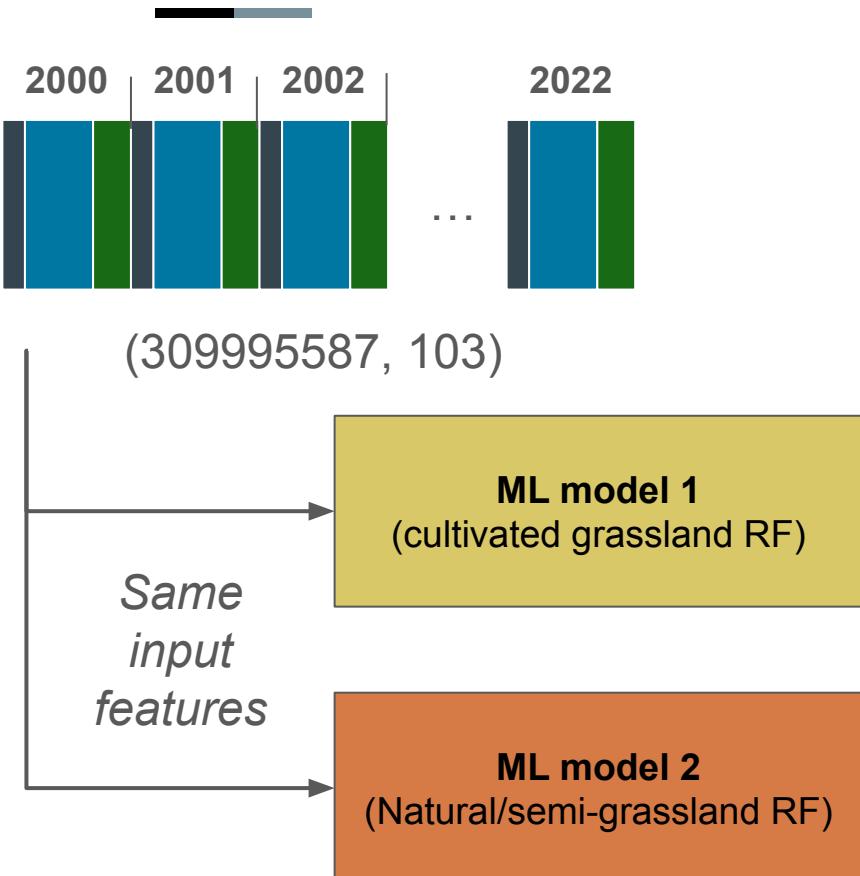
(309995587, 103)

*Same
input
features*



```
output = model.predict(data)  
# OR  
output = model.predict_proba(data)
```

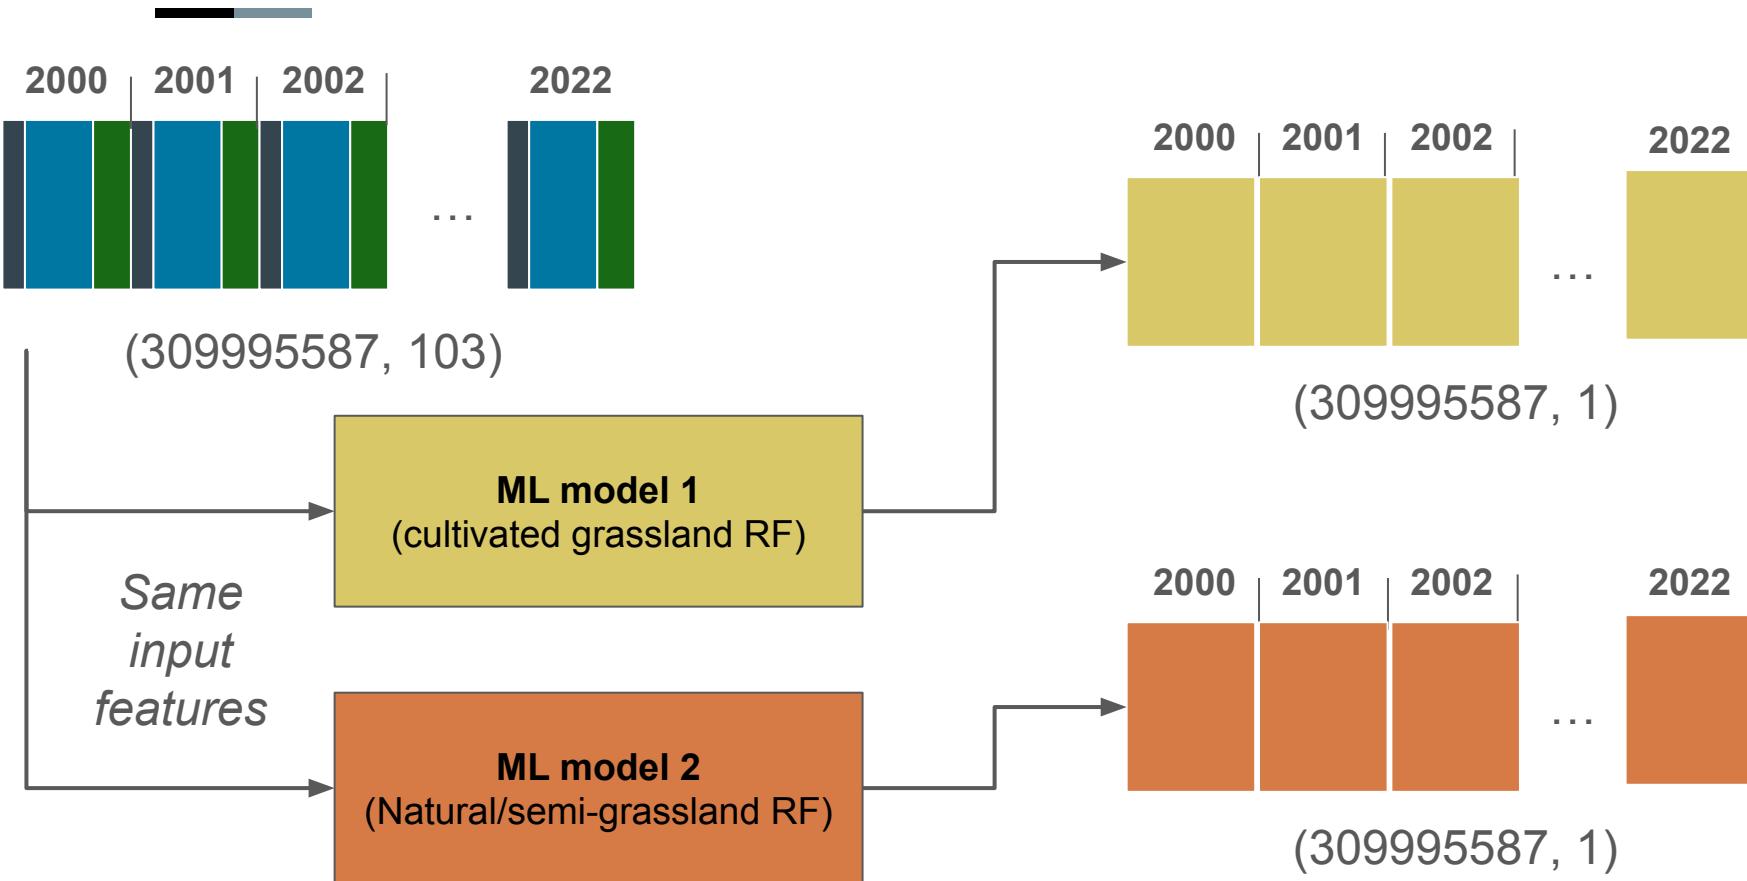
(Finally) Global ML prediction



<https://github.com/krzjoa/awesome-python-data-science>

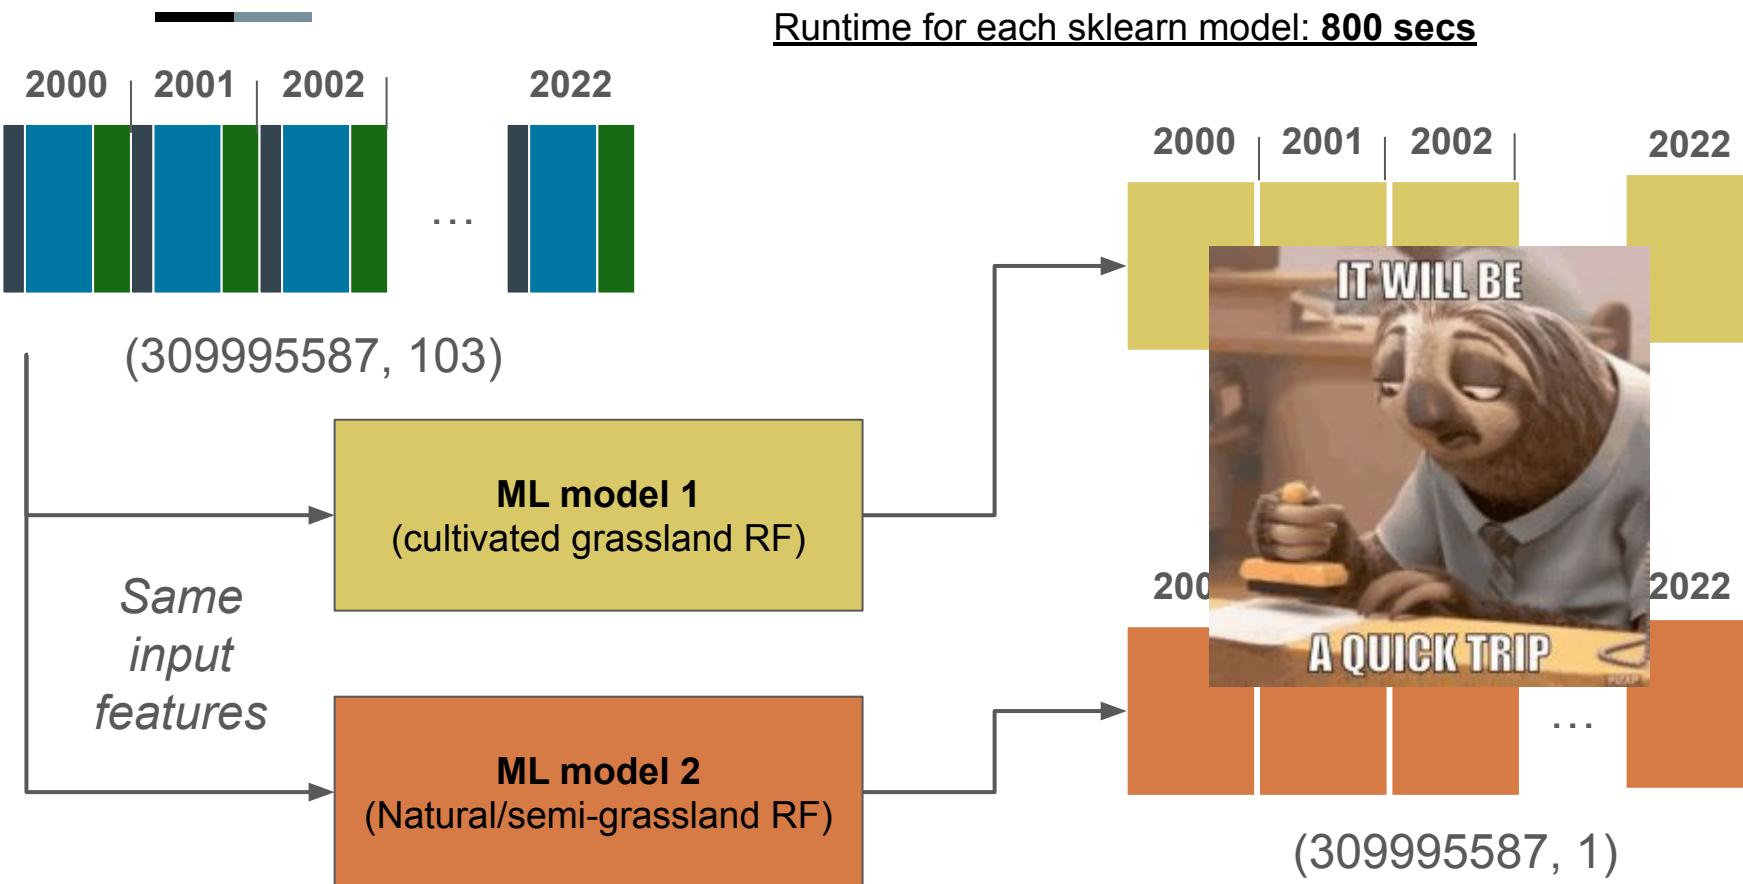
```
output = model.predict(data)  
# OR  
output = model.predict_proba(data)
```

(Finally) Global ML prediction

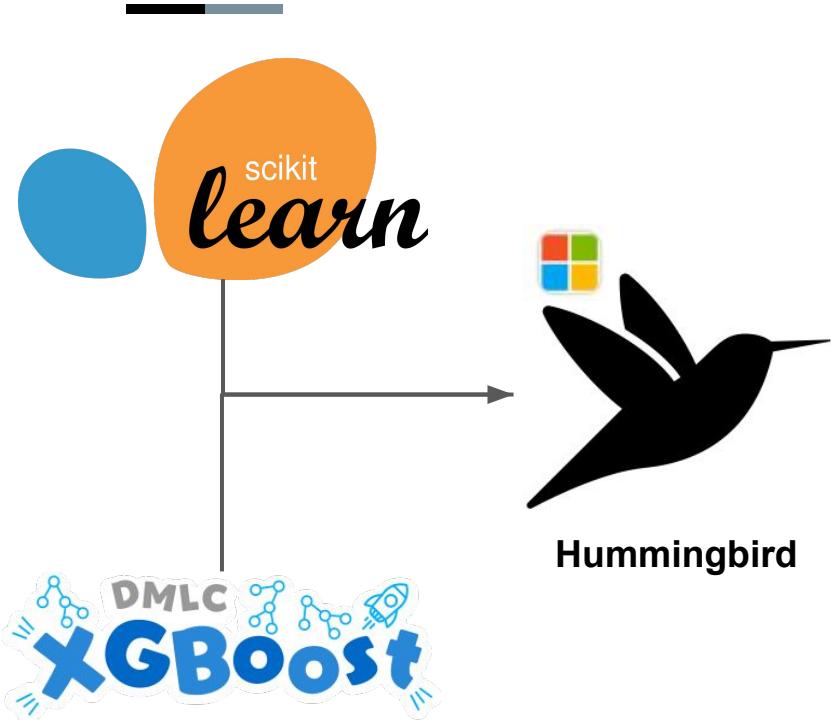


(Finally) Global ML prediction

Runtime for each sklearn model: 800 secs



Compiling ML models



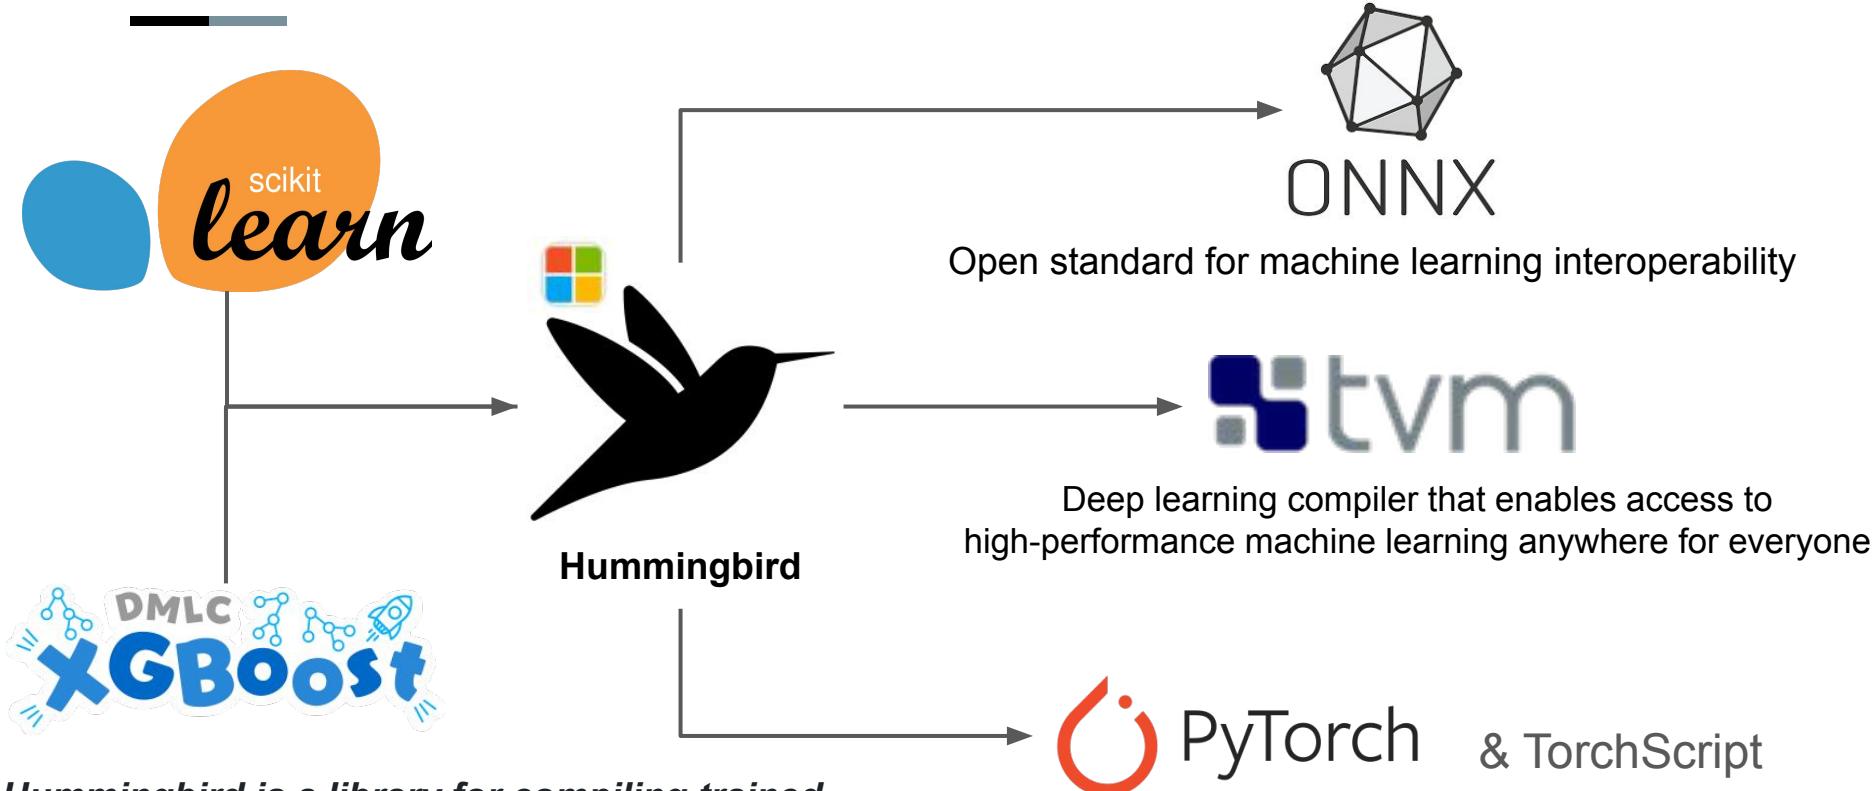
Hummingbird is a library for compiling trained traditional ML models into tensor computations.

<https://github.com/microsoft/hummingbird>

Compiling ML models



OpenGeo HUB
Connect • Create • Share • Repeat

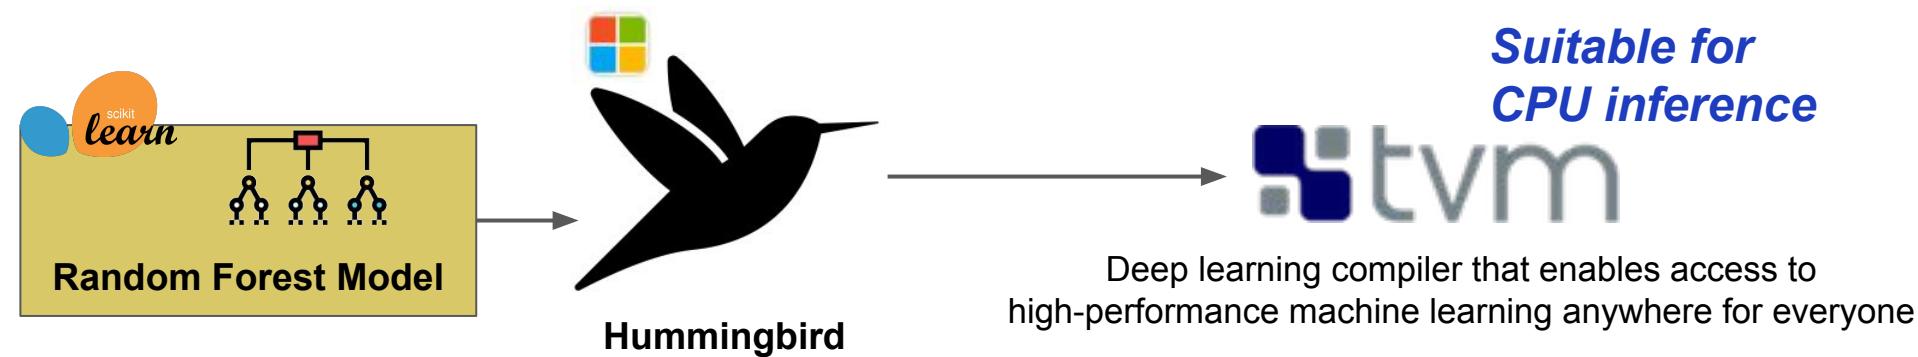


Hummingbird is a library for compiling trained traditional ML models into tensor computations.

<https://github.com/microsoft/hummingbird>

TorchScript is a way to create serializable and optimizable models from PyTorch code

Compiling ML models



Hummingbird is a library for compiling trained traditional ML models into tensor computations.

<https://github.com/microsoft/hummingbird>

Compiling ML models



OpenGeo HUB
Connect • Create • Share • Repeat

*Suitable for
CPU inference*



Hummingbird is a library for compiling trained traditional ML models into tensor computations.

<https://github.com/microsoft/hummingbird>

Compiling ML models



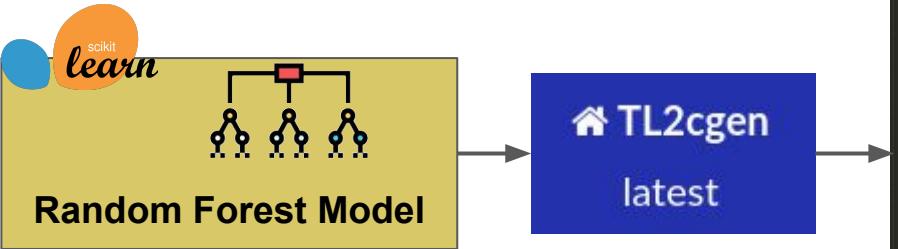
ML model 1
(cultivated grassland RF)

Number of trees: **60**
Max. depth (avg): **33.7**
Number of nodes (avg): **13,716**

ML model 2
(Natural/semi-grassland RF)

Number of trees: **60**
Max. depth (avg): **40.7**
Number of nodes (avg): **24,029**

Compiling ML models



Model compiler for decision tree models

<https://tl2cgen.readthedocs.io>

C++
compiler

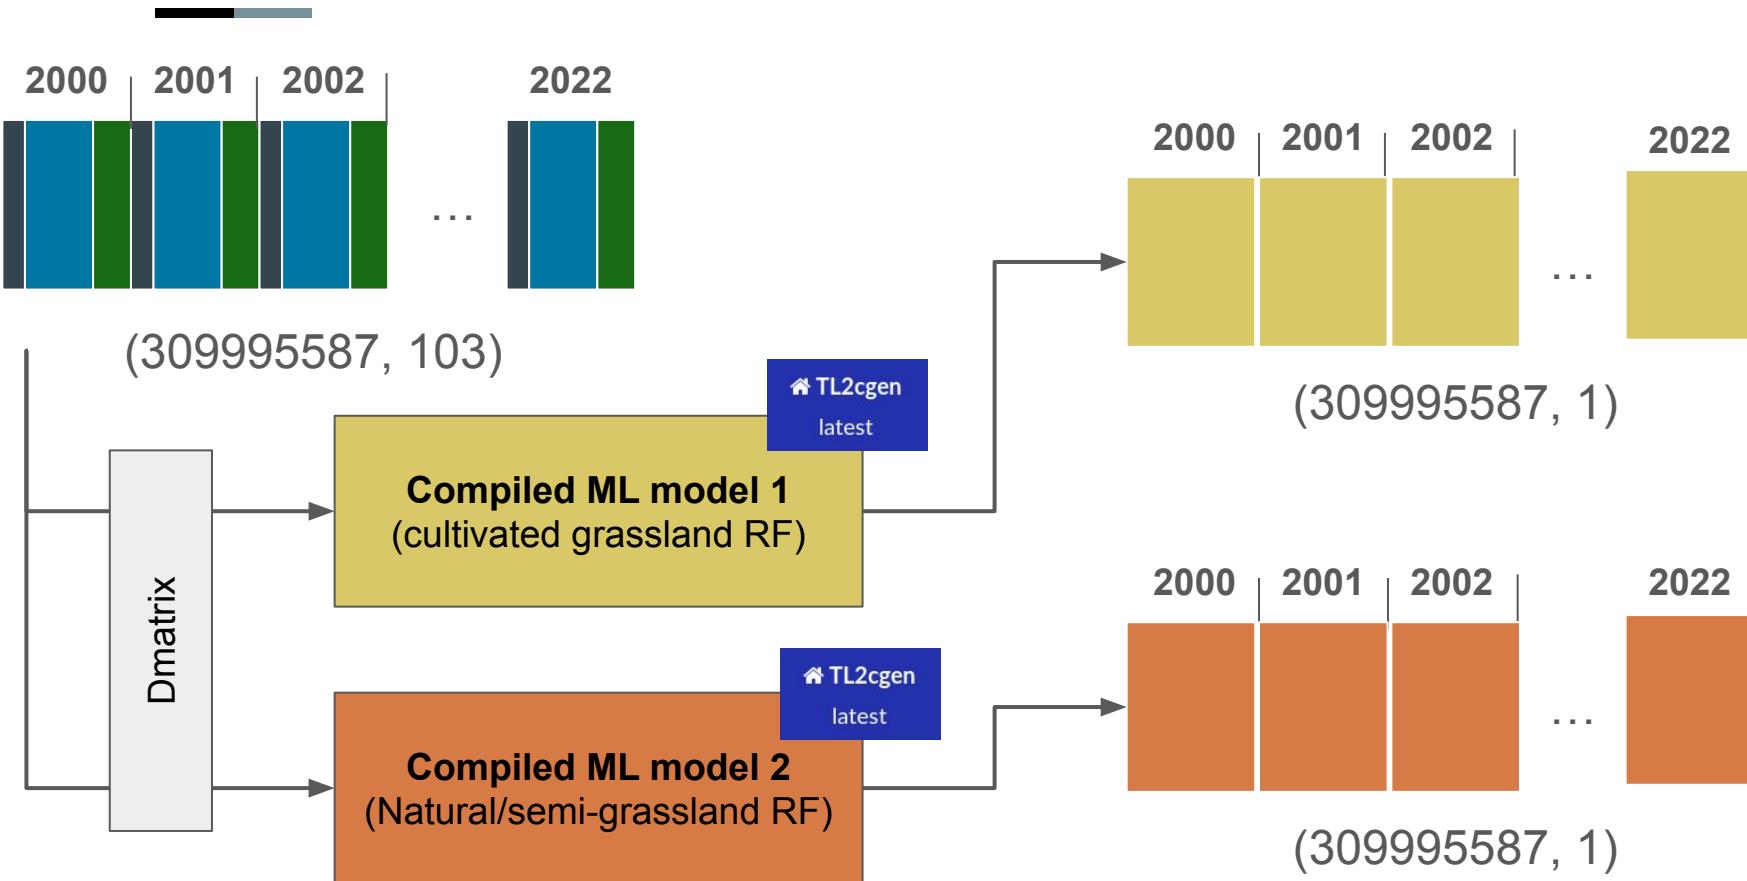


```
#include "header.h"
voi ... header.h"
d voi #include "header.h"
u d void predict_margin_multiclass_unit10(union Entry* data, double* result) {
i u double sum[3] = {0};
i i unsigned int tmp;
i int nid, cond, fid; /* used for folded subtrees */
if ( UNLIKELY( !(data[3].missing != -1) || (data[3].qvalue <= 24) ) ) {
    if ( LIKELY( !(data[5].missing != -1) || (data[5].qvalue <= 146) ) ) {
        if ( LIKELY( !(data[71].missing != -1) || (data[71].qvalue <= 652) ) ) {
            if ( UNLIKELY( !(data[84].missing != -1) || (data[84].qvalue <= 224) ) ) {
                if ( LIKELY( !(data[63].missing != -1) || (data[63].qvalue <= 678) ) ) {
                    if ( LIKELY( !(data[49].missing != -1) || (data[49].qvalue <= 8718) ) ) {
                        if ( UNLIKELY( !(data[77].missing != -1) || (data[77].qvalue <= 448) ) ) {
                            if ( UNLIKELY( !(data[93].missing != -1) || (data[93].qvalue <= 594) ) ) {
                                if ( LIKELY( !(data[25].missing != -1) || (data[25].qvalue <= 1050) ) ) {
                                    if ( UNLIKELY( !(data[29].missing != -1) || (data[29].qvalue <= 2776) ) ) {
                                        sum[0] += (double)0;
                                        sum[1] += (double)0;
                                        sum[2] += (double)1;
                                    } else {
                                        if ( LIKELY( !(data[33].missing != -1) || (data[33].qvalue <= 726) ) ) {
                                            sum[0] += (double)0;
                                            sum[1] += (double)0.5952380952380952328;
                                            sum[2] += (double)0.4047619047619047672;
                                        } else {
                                            sum[0] += (double)0;
                                            sum[1] += (double)0;
                                            sum[2] += (double)1;
                                        }
                                    } else {
                                }
                            }
                        }
                    }
                }
            }
        }
    }
}
```

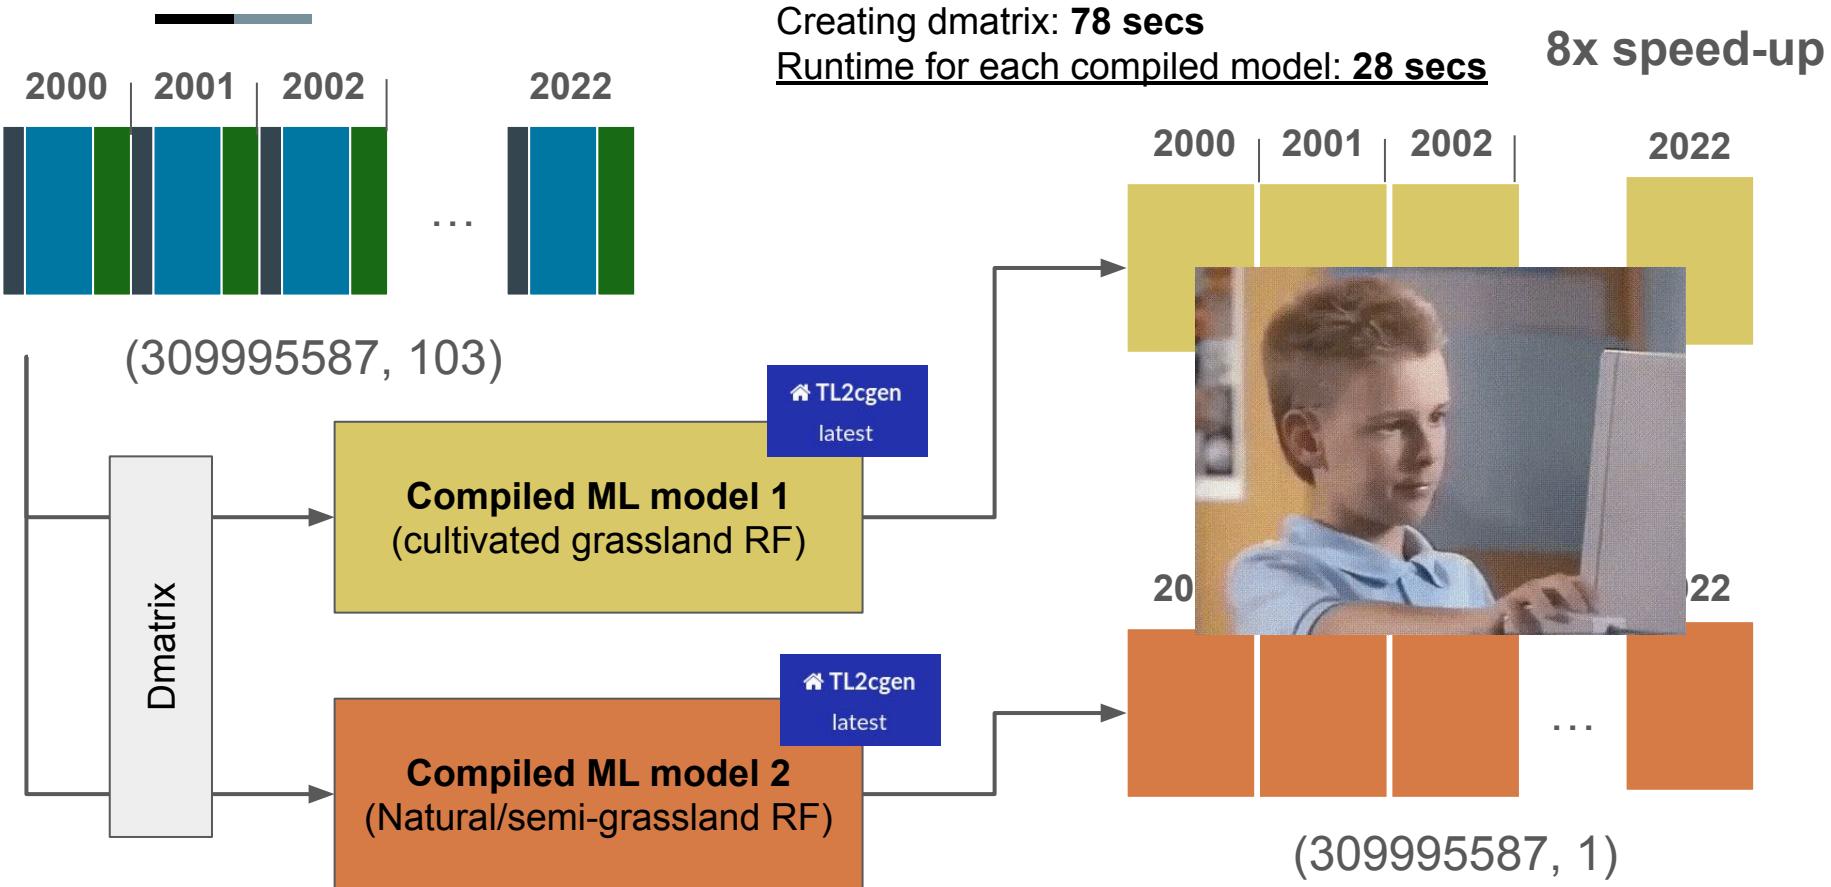
One IF-ELSE per tree node

[The Performance of Decision Tree Evaluation Strategies](#)

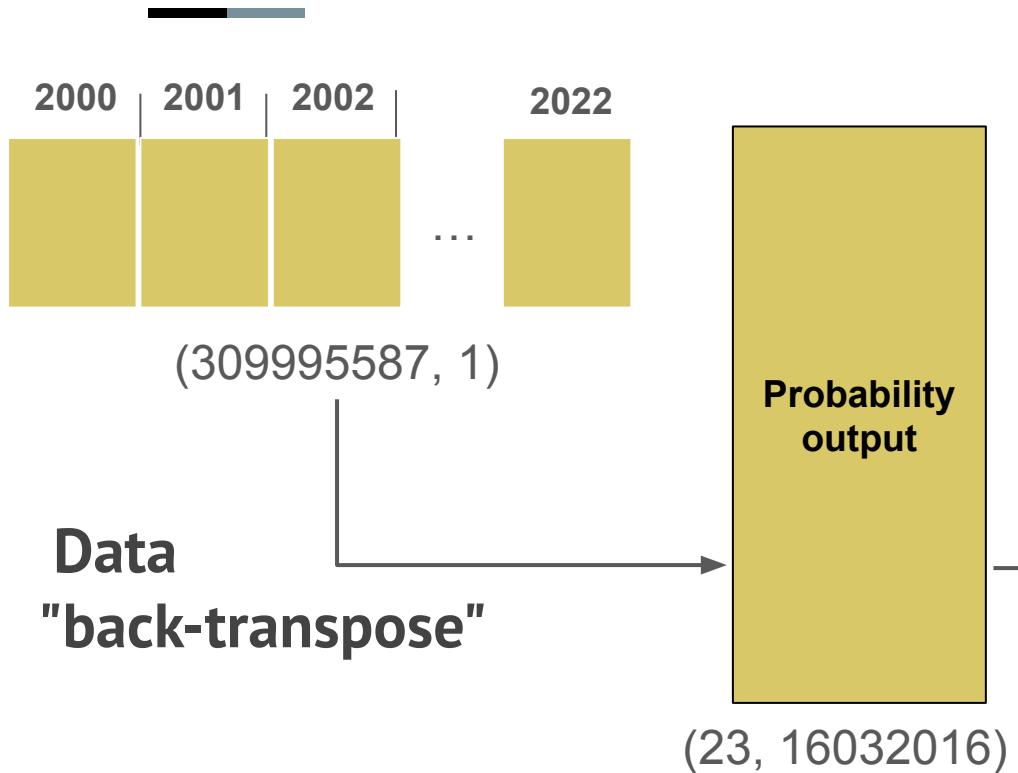
(Finally) Global ML prediction



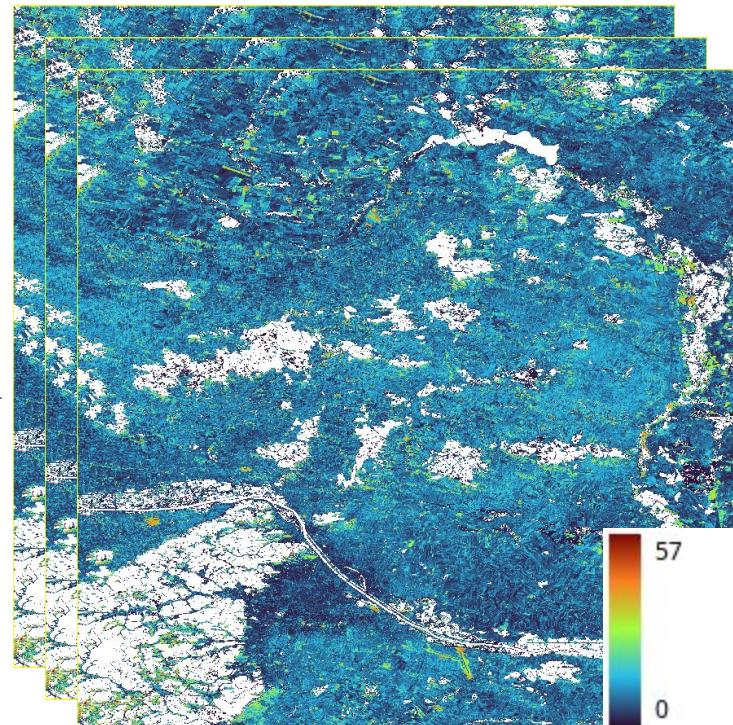
(Finally) Global ML prediction



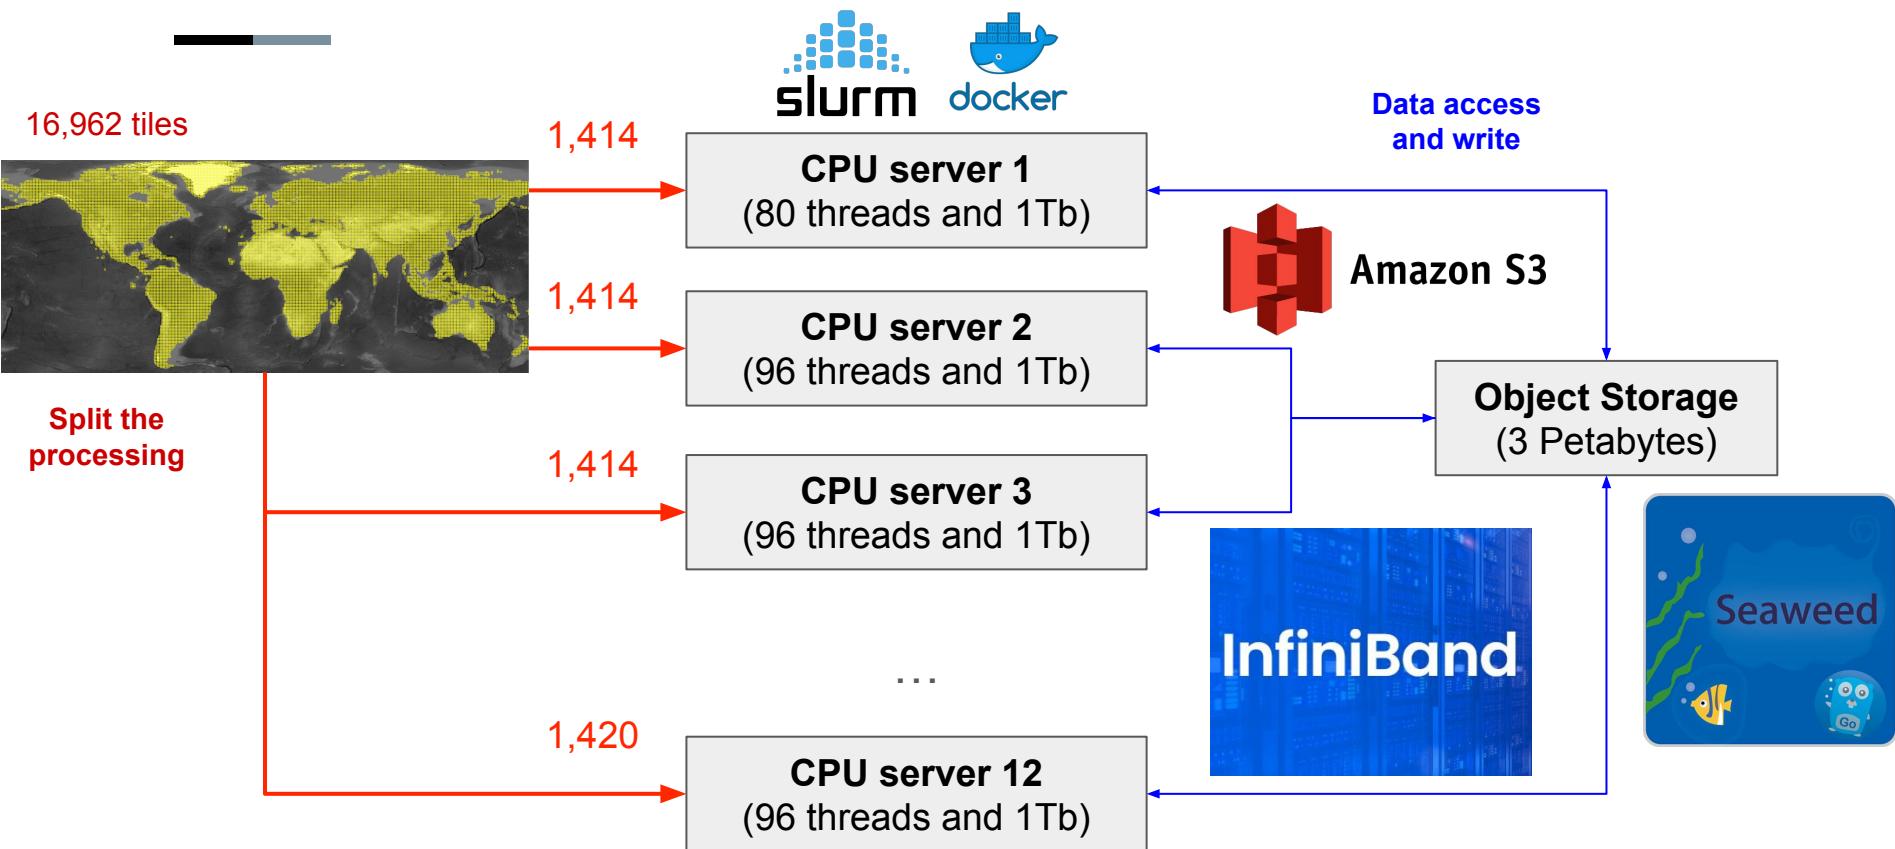
(Finally) Global ML prediction



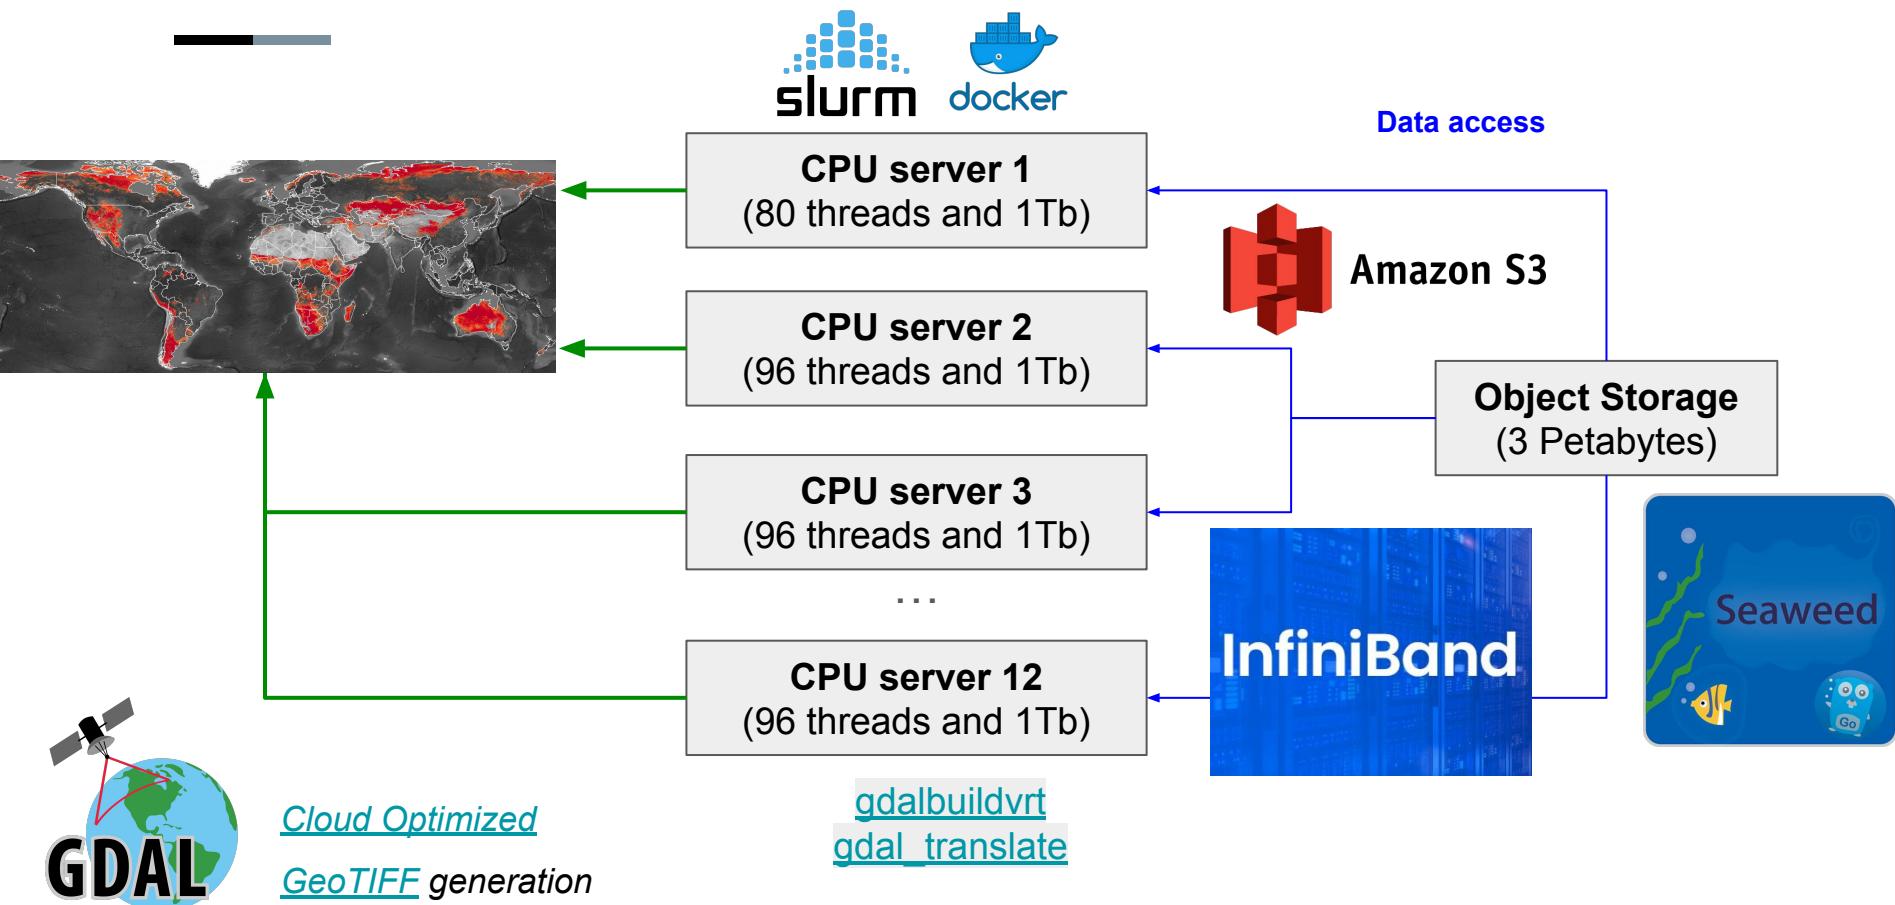
Cultivated grassland maps



(Finally) Global ML prediction

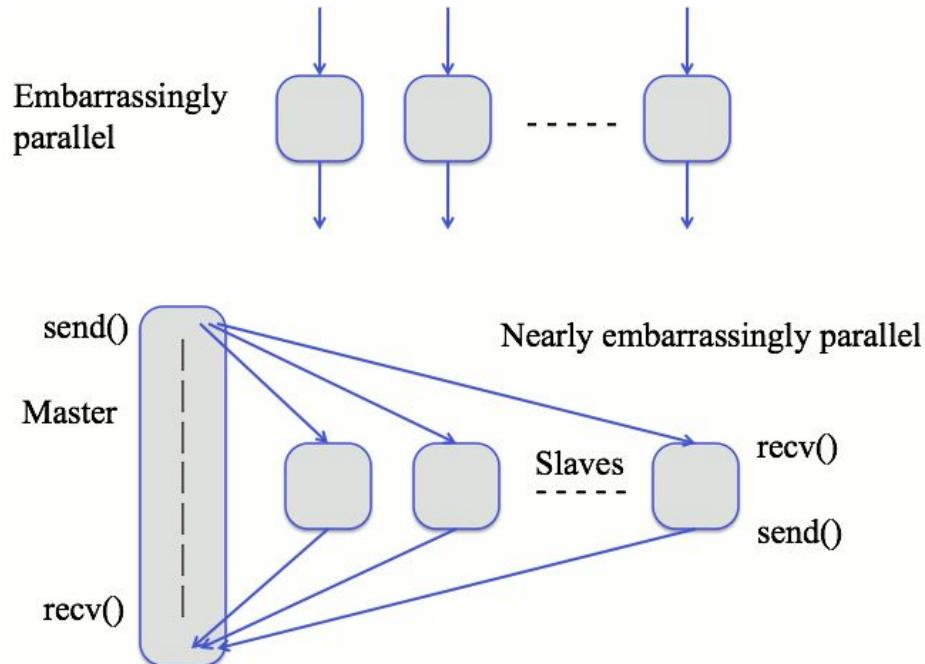


(Finally) Global ML prediction



Embarrassingly parallel problems

- Can be divided into completely **independent** parts,
- Requires none or very little communication,
- **Nearly embarrassingly parallel** is an embarrassingly parallel computation that requires initial data to be distributed and final results to be collected in some way



Implementation in slurm

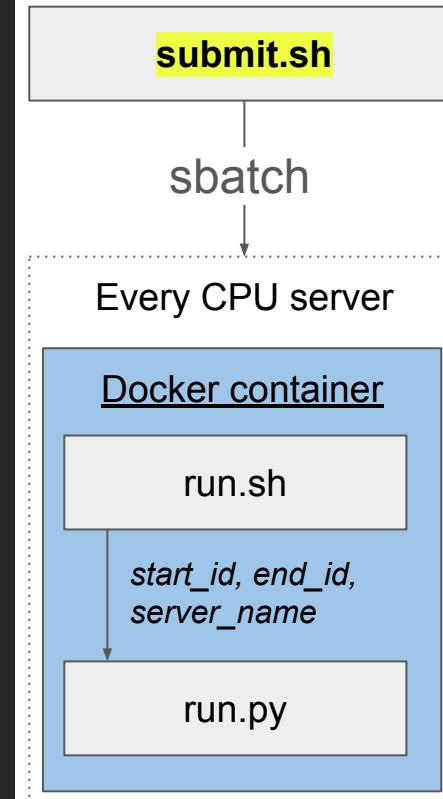
```
#!/bin/bash
#SBATCH --partition=cpu
#SBATCH --ntasks-per-node=1
#SBATCH --ntasks=1
#SBATCH --job-name=wri_pasture_class
#SBATCH --array=1-12
#SBATCH --mail-type=ALL
#SBATCH --mail-user=leandro.parente@opengeohub.org
#SBATCH --output=%x_%N.log
#SBATCH --exclude=landmark,bender # Servers to exclude (landmark,primus)

# Slurm helper functions
source /mnt/slurm/jobs/ogh_slurm.sh

# Total number of ids to split and process across the servers
N_IDS=16962 # FAPAR

# Python docker container
DOCKER_IMAGE=192.168.49.30:5000/pygeo-ide:v3.8.16-mkl-gdal362-pasture_class
LANG=Python

execute
```



Implementation in slurm

```
#!/bin/bash

job_command=( "$@" )

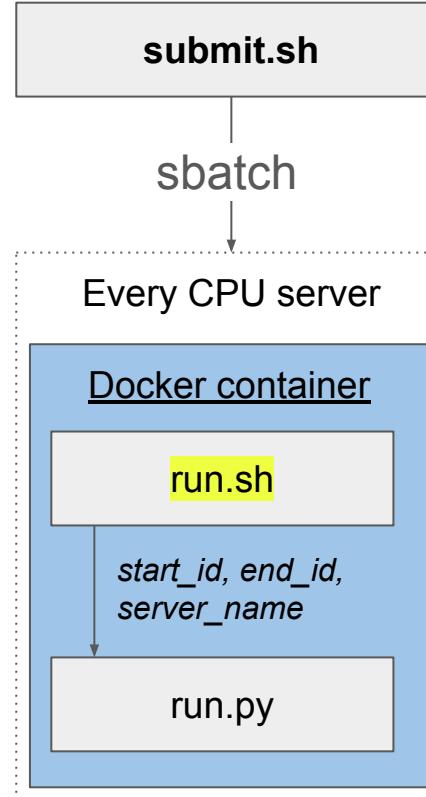
#####
# Intall all necessary libraries inside the container
#####
pip install -e 'git+https://github.com/openlandmap/scikit-map.git@pasture_class_prod#egg=scikit-map[full]'
pip install -e 'git+https://github.com/leiferlab/savitzkygolay#egg=savitzkygolay'

mc alias set g1 http://192.168.49.30:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g2 http://192.168.49.31:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g3 http://192.168.49.32:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g4 http://192.168.49.33:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g5 http://192.168.49.34:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g6 http://192.168.49.35:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g7 http://192.168.49.36:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g8 http://192.168.49.37:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g9 http://192.168.49.38:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g10 http://192.168.49.39:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g11 http://192.168.49.40:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g12 http://192.168.49.41:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4
mc alias set g13 http://192.168.49.42:8333 iwum9G1fEQ9201YV4o19 GMBME3Wsm8S7mBXw3U4CNWurkzWMqGZ0n2rXHggS0 --api S3v4

cd /mnt/slurm/jobs/wri_pasture_class/
mkdir /mnt/$(hostname)/wri_pasture_class

cp -Rv /mnt/slurm/jobs/wri_pasture_class/models/compiled /mnt/$(hostname)/wri_pasture_class/

echo #####
echo $job_command
echo #####
echo "Starting: $(date)"
$job_command
echo "Ending: $(date)"
```

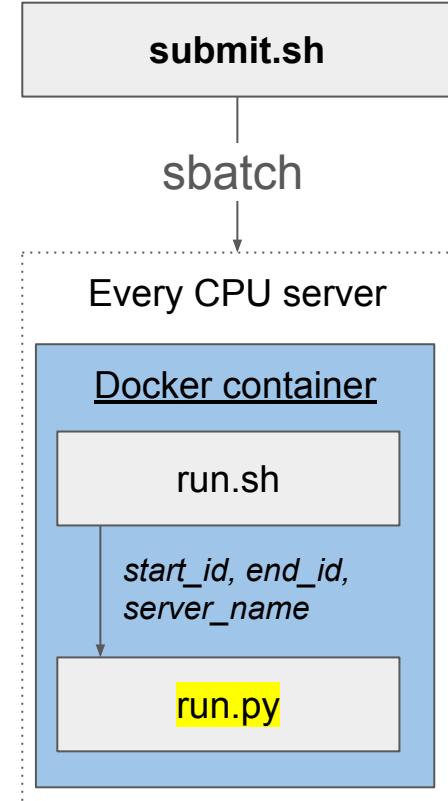


Implementation in slurm

```
import time
import sys

start_id=int(sys.argv[1])
end_id=int(sys.argv[2])
server_name=sys.argv[3]

print(f"Python: Processing the ids {start_id}-{end_id} in {server_name}")
time.sleep(10)
print(f"Python: ...actually {server_name} was just chilling out")
```



Main takeaways

- The current ML prediction pipeline is complex and difficult to maintain, however reduced the production time from **40 to 7 days**,
- We are encapsulating C++ functions into scikit-map library (<https://github.com/openlandmap/scikit-map>),
- We observed huge improvement in performance after migrate from **Python multiprocessing** to **C++ multithreading** ([more information about the Global Interpreter Lock GIL](#)),
- Most of the users **do not need** such complex pipeline, however part of the optimizations presented might be useful for other applications (ex: parallel reading, on-the-fly calculation, third-party libraries).
- The full reproducibility of this pipeline depends on have the **input EO data** publicly accessible, **container with all libraries** used and **better code abstraction** for the optimizations implemented.





Leandro Parente

leandro.parente@opengeohub.org

<https://opengeohub.org>



OpenGeoHUB
Connect • Create • Share • Repeat

GRACIAS DANKSCHEEN
ARIGATO TASHAKKUR ATU
SHUKURIA SUKSAMA EKHMET
JUSPAXAR YACHANVELAY
GOZAIMASHITA MEHRBANI
EFCHARISTO MARKE
KOMPASUNDA TINGKI
MERHBA SHÜYAN SHUKRIA
MERHBA TASHAKKUR ATU
MERHBA SUKSAMA EKHMET
MERHBA YACHANVELAY
MERHBA SHÜYAN SHUKRIA
THANK YOU OBRIGADO
BOLZİN MERCI

