

Foundation Models for Geospatial Data

Antonio Fonseca

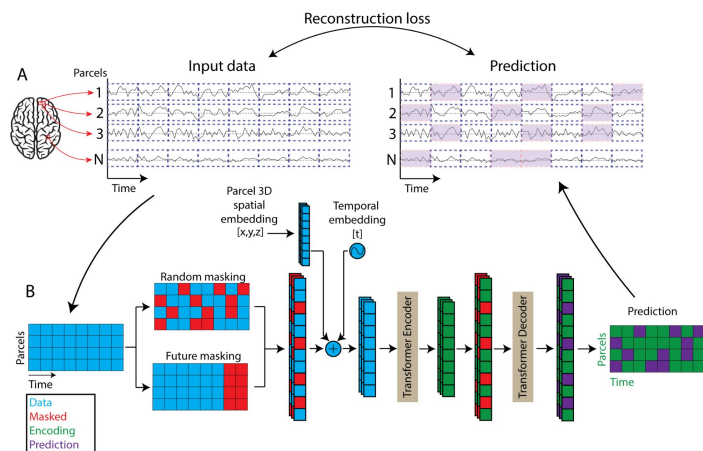


Agenda

- Intro:
 - Frameworks of learning
- The building block of modern ML
 - Transformers
 - Applications: language and image
- Foundation Models (FM)
 - FM vs conventional ML models
 - Prithvi: geospatial foundational model (IBM-Nasa)
 - Applications
- Tutorial



About me

- PhD in Neuroscience
- Foundation models for biomedical application





Brain recordings


Single-cell expression profiles



Drug discovery



From supervised learning to Foundation Models

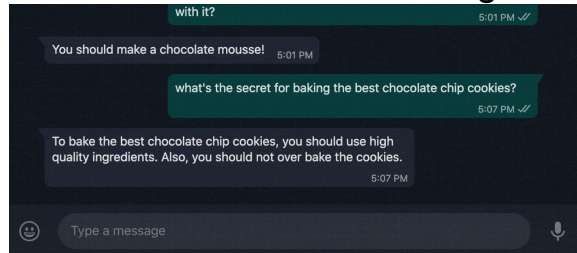


Neural Net for
Supervised Learning

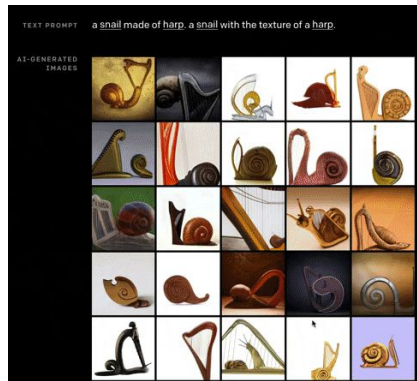
task-specific ML models

???

ChatGPT: conversational agent



Dall-E: generate image from prompt



general-purpose ML models

Frameworks for Machine Learning

Most real-world problems fall into the category of unsupervised learning.

■ "Pure" Reinforcement Learning (cherry)

- ▶ The machine predicts a scalar reward given once in a while.
- ▶ **A few bits for some samples**

■ Supervised Learning (icing)

- ▶ The machine predicts a category or a few numbers for each input
- ▶ Predicting human-supplied data
- ▶ **10→10,000 bits per sample**

■ Unsupervised/Predictive Learning (cake)

- ▶ The machine predicts any part of its input for any observed part.
- ▶ Predicts future frames in videos
- ▶ **Millions of bits per sample**

■ (Yes, I know, this picture is slightly offensive to RL folks. But I'll make it up)



Unsupervised Learning techniques

1. **Clustering:**

- Grouping similar data points together.
- Example: K-means, Hierarchical Clustering, Spectral Clustering

2. **Dimensionality Reduction:**

- Reducing the number of features.
- Example: PCA, UMAP, VAE.

3. **Sequence learning:**

- Learning patterns in series
- Example 1: RNNs, LSTMs
- Example 2: BERT, GPT, LLaMA (all based on Transformers).

Transformers vs Neural Networks

What Separates Transformers from Previous Neural Networks?

Self-Attention Mechanism:

- Allows the model to focus on different parts of the input sequence when making predictions.
- Captures long-range dependencies more effectively than traditional RNNs (Recurrent Neural Networks).

Parallel Processing:

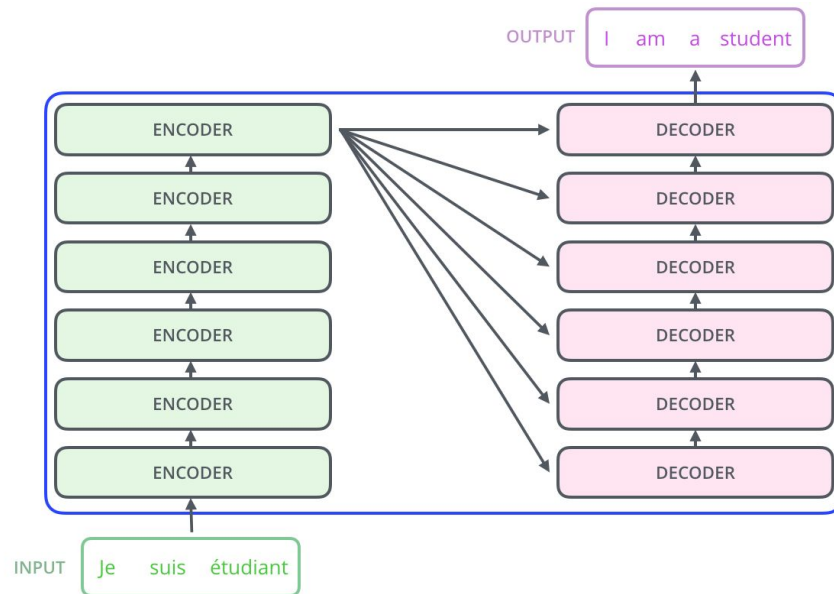
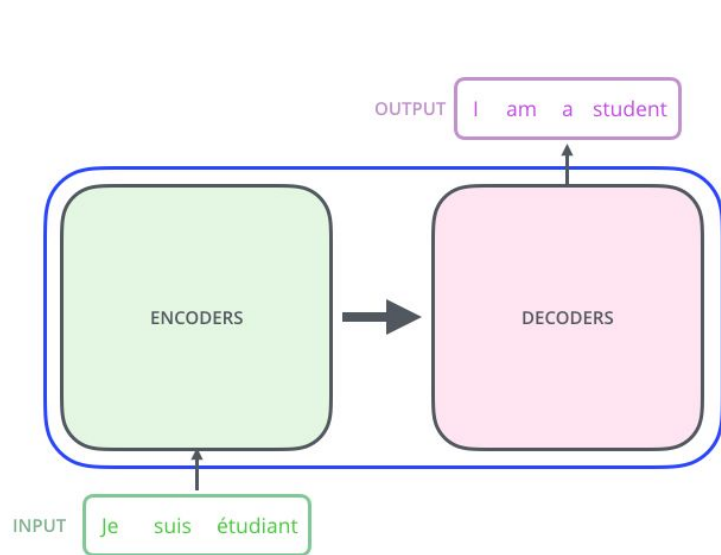
- Transformers process all tokens in the input sequence simultaneously.
- Faster training times compared to sequential processing in RNNs and LSTMs (Long Short-Term Memory networks).

Versatility:

- Applied to various tasks beyond NLP, including image processing and multimodal learning.
- Unified architecture for diverse applications.

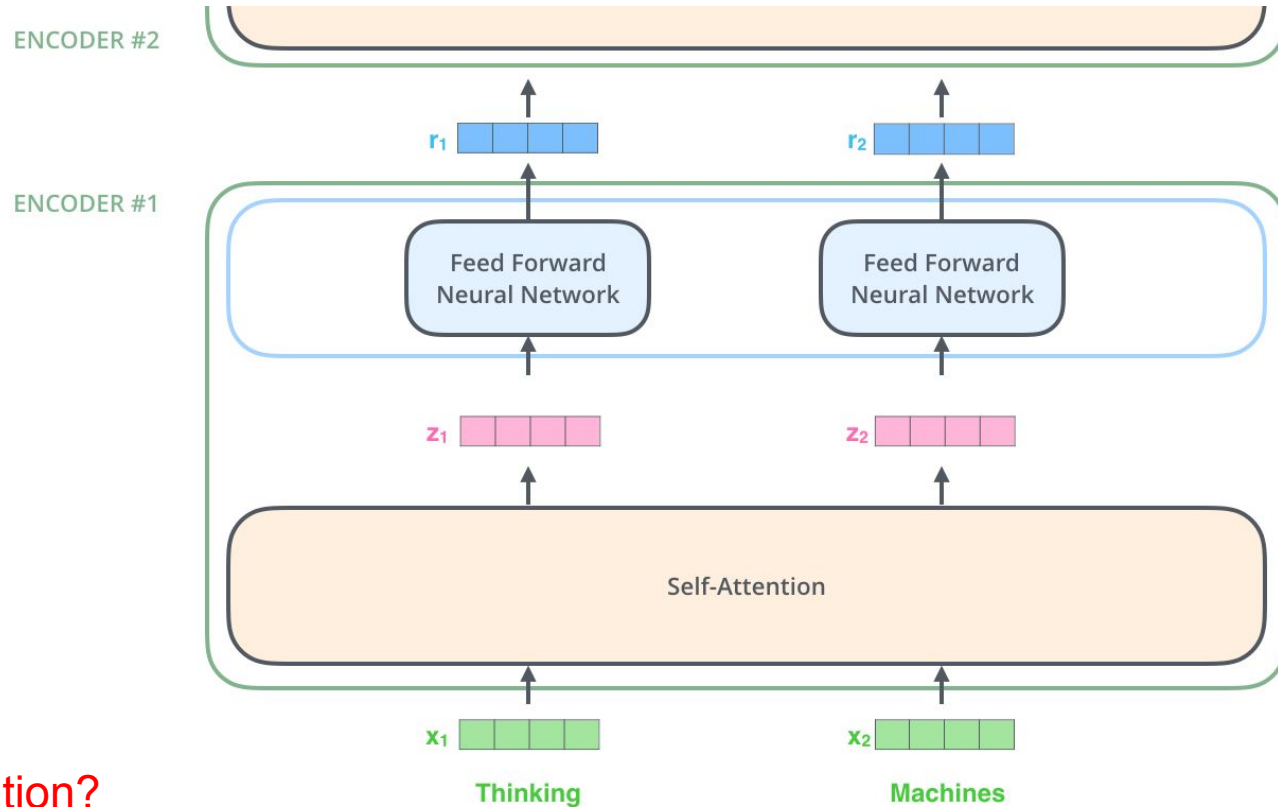
How are Transformers implemented ?

Transformer in high-level



What is inside the blocks?


Each Encoder has Self-Attention



What is attention?

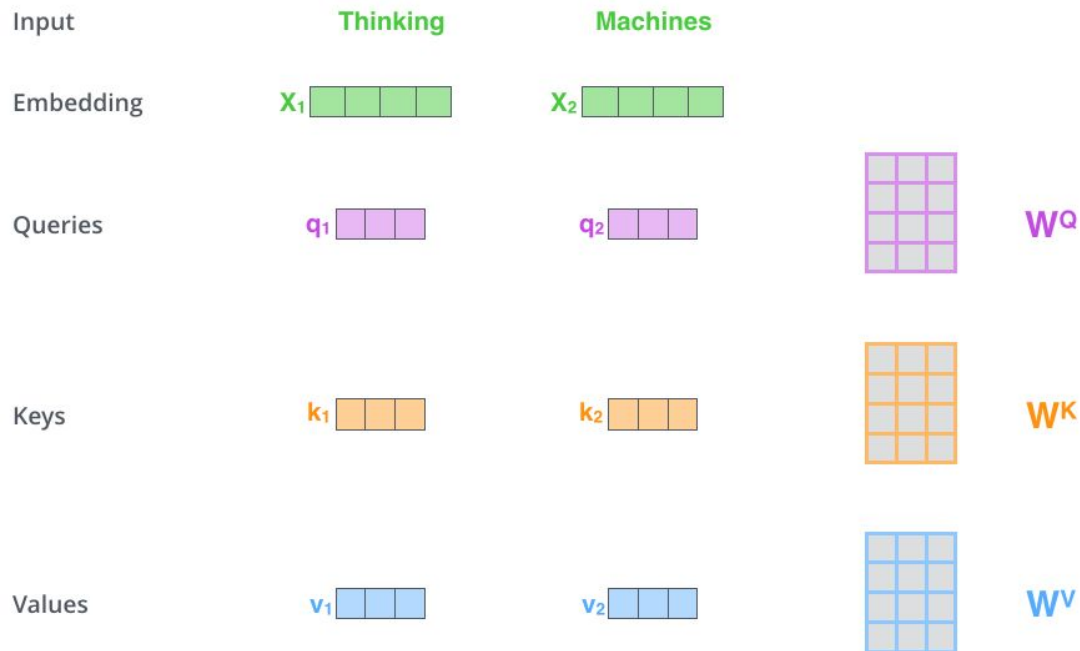
Self-Attention at a High Level

Self-Attention



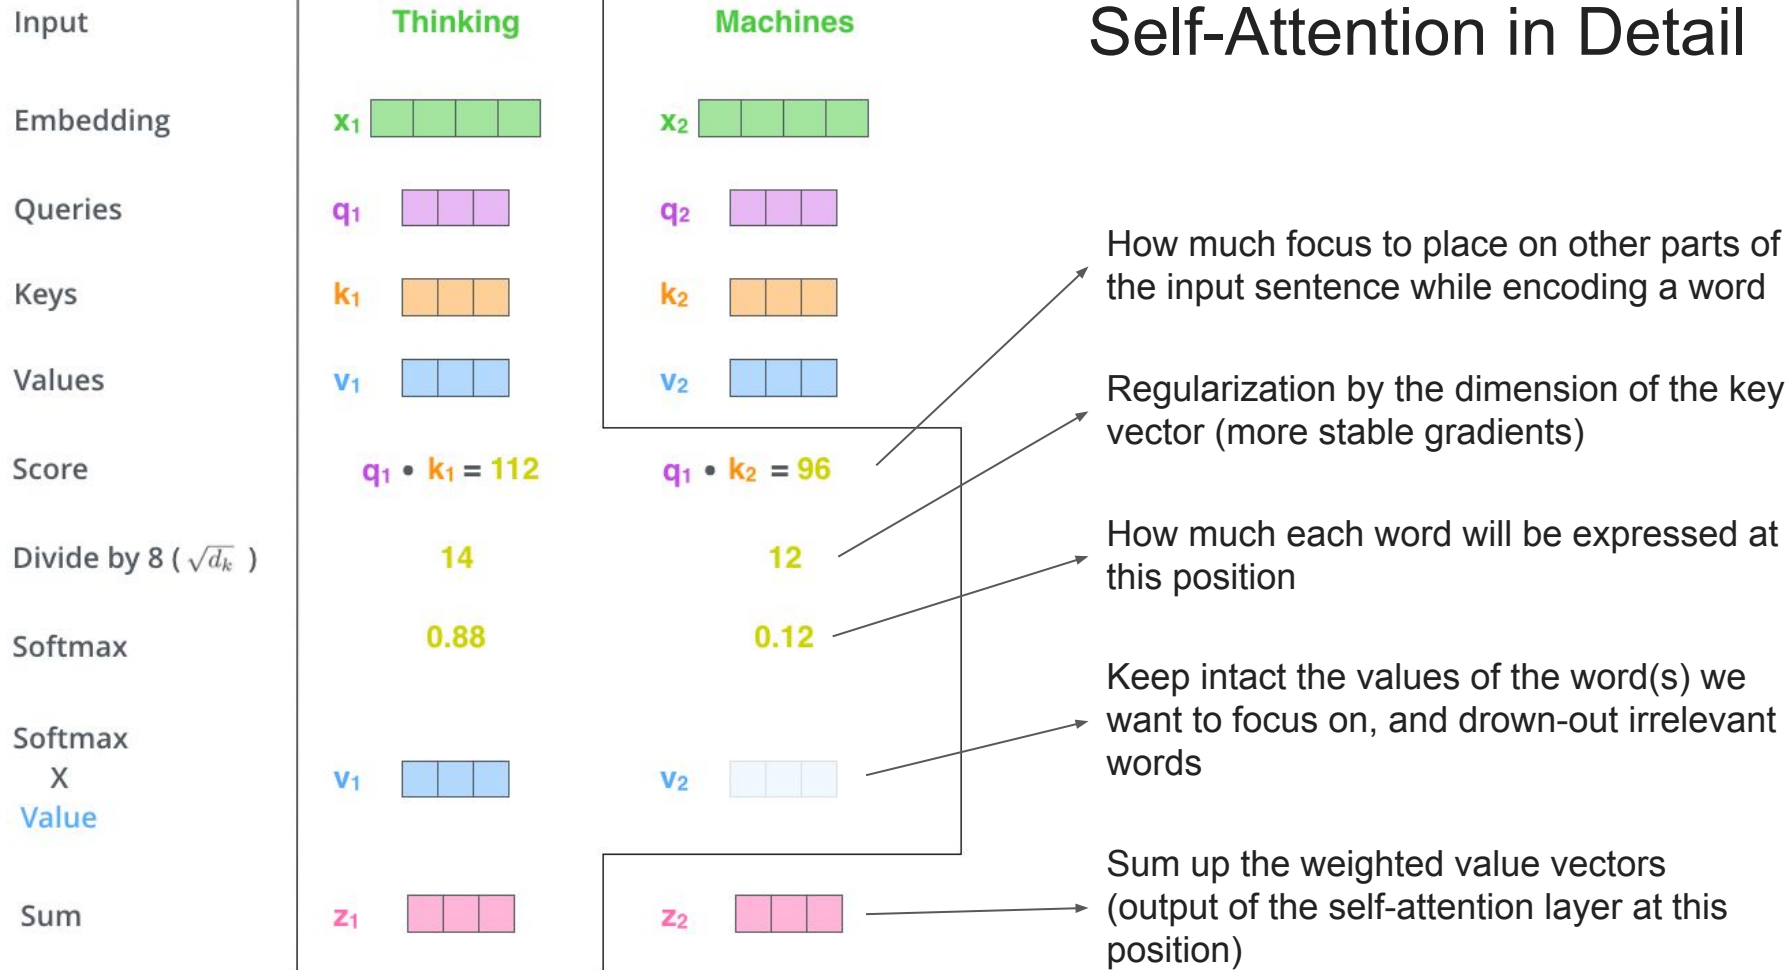
How is it implemented?

Query, Key, Value

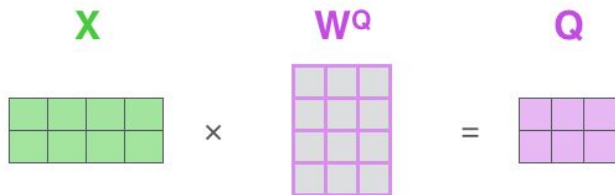


Multiplying x_1 by the W^Q weight matrix produces q_1 , the "query" vector associated with that word. We end up creating a "query", a "key", and a "value" projection of each word in the input sentence.

Self-Attention in Detail



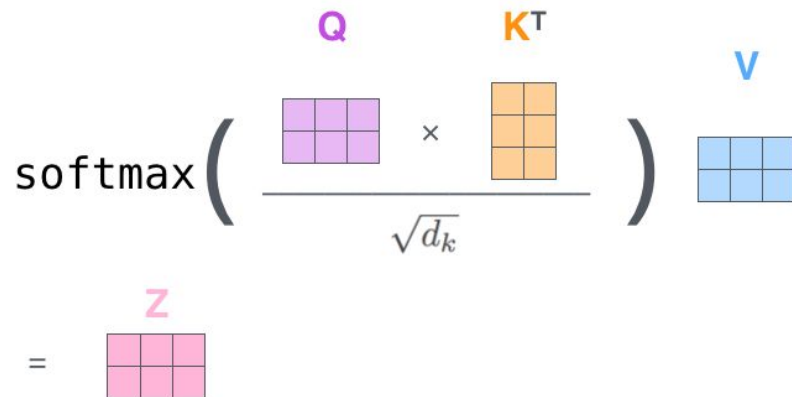
Self-Attention as Matrix Calculation

$$\mathbf{X} \times \mathbf{W}^Q = \mathbf{Q}$$



$$\mathbf{X} \times \mathbf{W}^K = \mathbf{K}$$


$$\mathbf{X} \times \mathbf{W}^V = \mathbf{V}$$


Self-attention is calculated in matrix form for faster processing

$$\text{softmax}\left(\frac{\mathbf{Q} \times \mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V} = \mathbf{Z}$$


Multi-head Attention

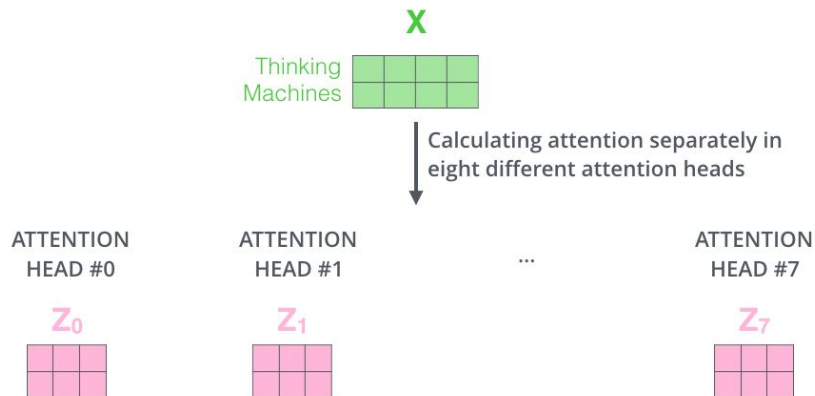


Multi-heads benefits:

- It expands the model's ability to focus on different positions.
- It gives the attention layer multiple "representation subspaces"

Self-attention with 8 heads:

- eight different times with different weight matrices, we end up with eight different Z matrices



Concatenating Multi-head Attention


1) Concatenate all the attention heads



2) Multiply with a weight matrix W^O that was trained jointly with the model

\times

3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN



Putting it all together



1) This is our input sentence*

Thinking
Machines


2) We embed each word*



3) Split into 8 heads.
We multiply X or R with weight matrices



...




4) Calculate attention using the resulting $Q/K/V$ matrices




...




5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



...



W^O

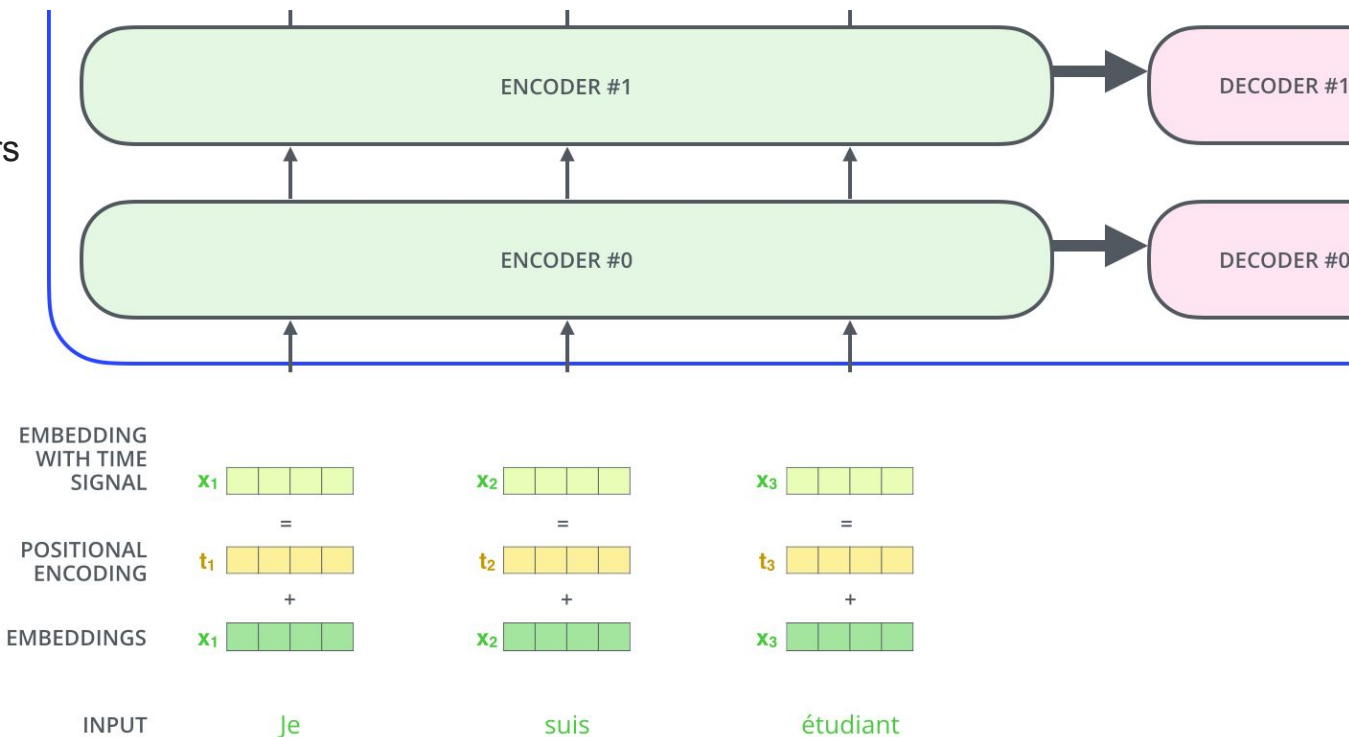


* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



Representing The Order of The Sequence Using Positional Encoding

Adding meaningful distances
between the embedding vectors



Positional Encoding

- Positions can't be encoded by numbers because sequences of different lengths can't be trained this way
- Instead use a periodic function like sin, cos to create the embedding

$$\vec{p}_t^{(i)} = f(t)^{(i)} := \begin{cases} \sin(\omega_k \cdot t), & \text{if } i = 2k \\ \cos(\omega_k \cdot t), & \text{if } i = 2k + 1 \end{cases}$$


$$\omega_k = \frac{1}{10000^{2k/d}}$$

Binary vs Positional Encoding

Why not using Binary Encoding

- Positional encoding scales better with longer sequences due to its fixed dimensionality regardless of the sequence length.
- Binary encoding requires more bits as the sequence length increases, which might not be efficient for very long sequences.

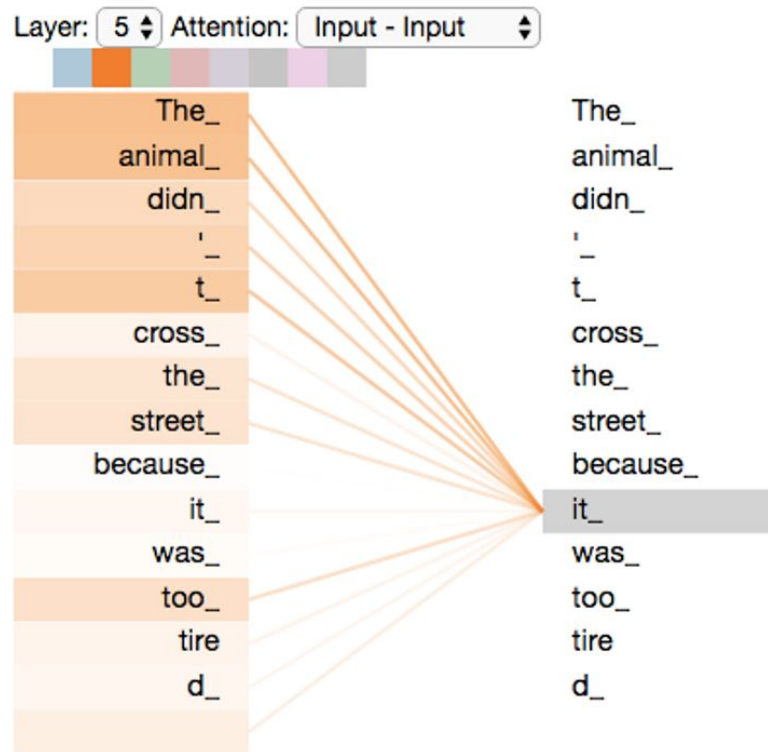
0 :	0	0	0	0	8 :	1	0	0	0
1 :	0	0	0	1	9 :	1	0	0	1
2 :	0	0	1	0	2 :	1	0	1	0
3 :	0	0	1	1	11 :	1	0	1	1
4 :	0	1	0	0	12 :	1	1	0	0
5 :	0	1	0	1	13 :	1	1	0	1
6 :	0	1	1	0	14 :	1	1	1	0
7 :	0	1	1	1	15 :	1	1	1	1



Attention maps over tokens

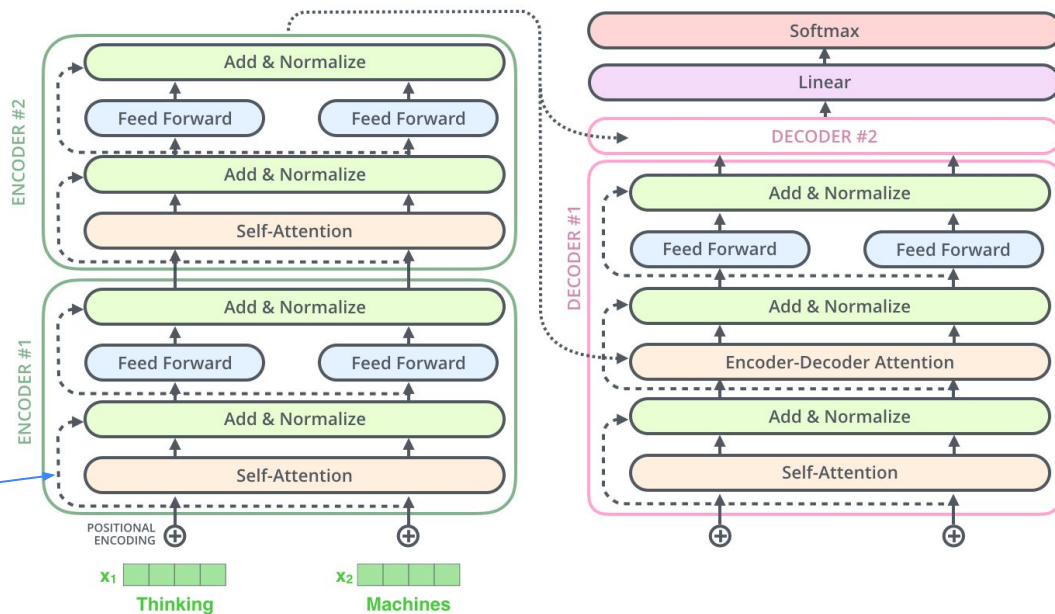
“The animal didn’t cross the street because it was too tired”

Self-attention helps a word gather information from other positions in the input sequence to enhance its representation.



Transformer architecture

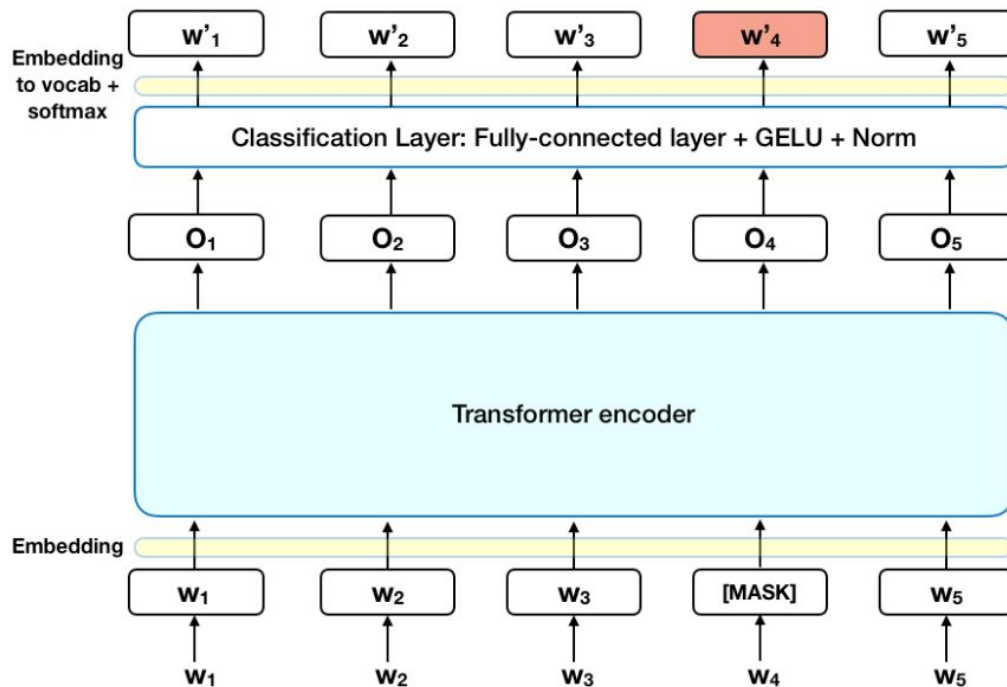
Residual
connection



Learning language with Transformer-based models

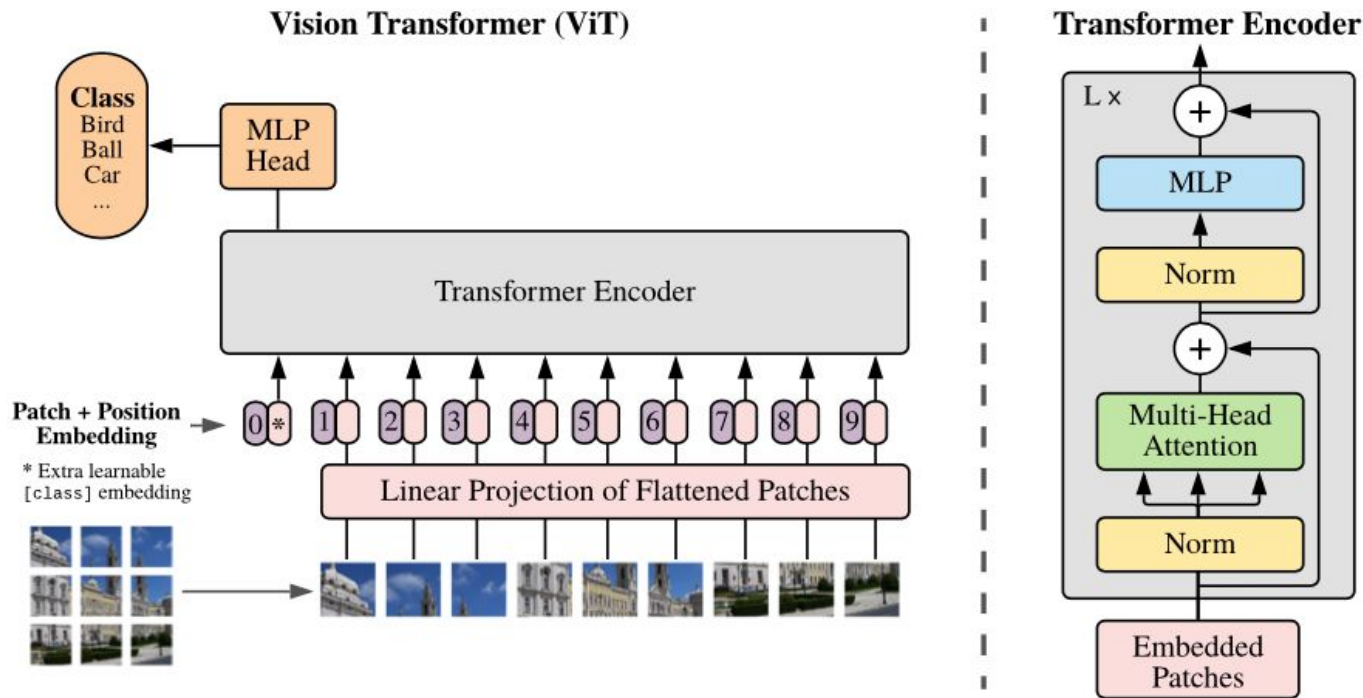
BERT: Bidirectional Encoder Representations from Transformers

- The meaning is heavily dependent on the context
- Learning a language via predicting masked tokens (words/characters)
- Contextualized prediction



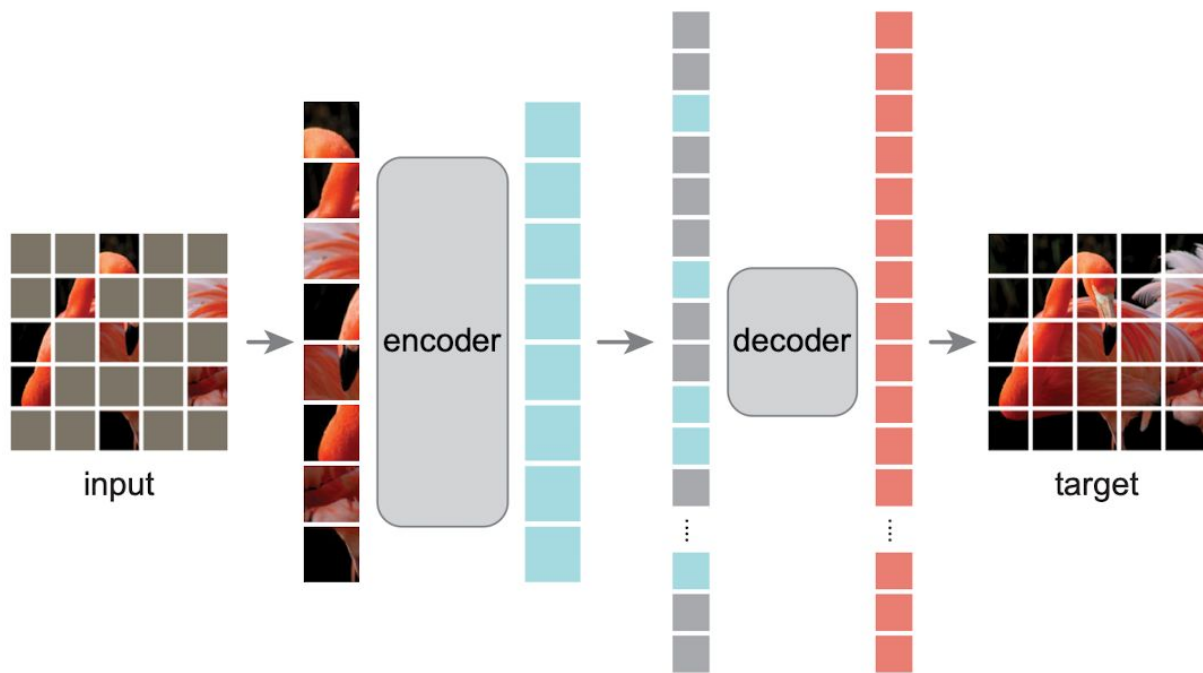
Can we apply this to other domains?

Vision Transformer (ViT)



How can we turn this into an unsupervised task?

Masked Autoencoder (MAE)



Foundation Models

Large Models + Big data = Foundation Models

Key Characteristics:

- **Scale:** Large models that are trained on extensive datasets.
- **Unsupervised:** Trained to predict the data itself
- **Transfer Learning:** Use pre-trained knowledge for task-specific applications.

Examples:

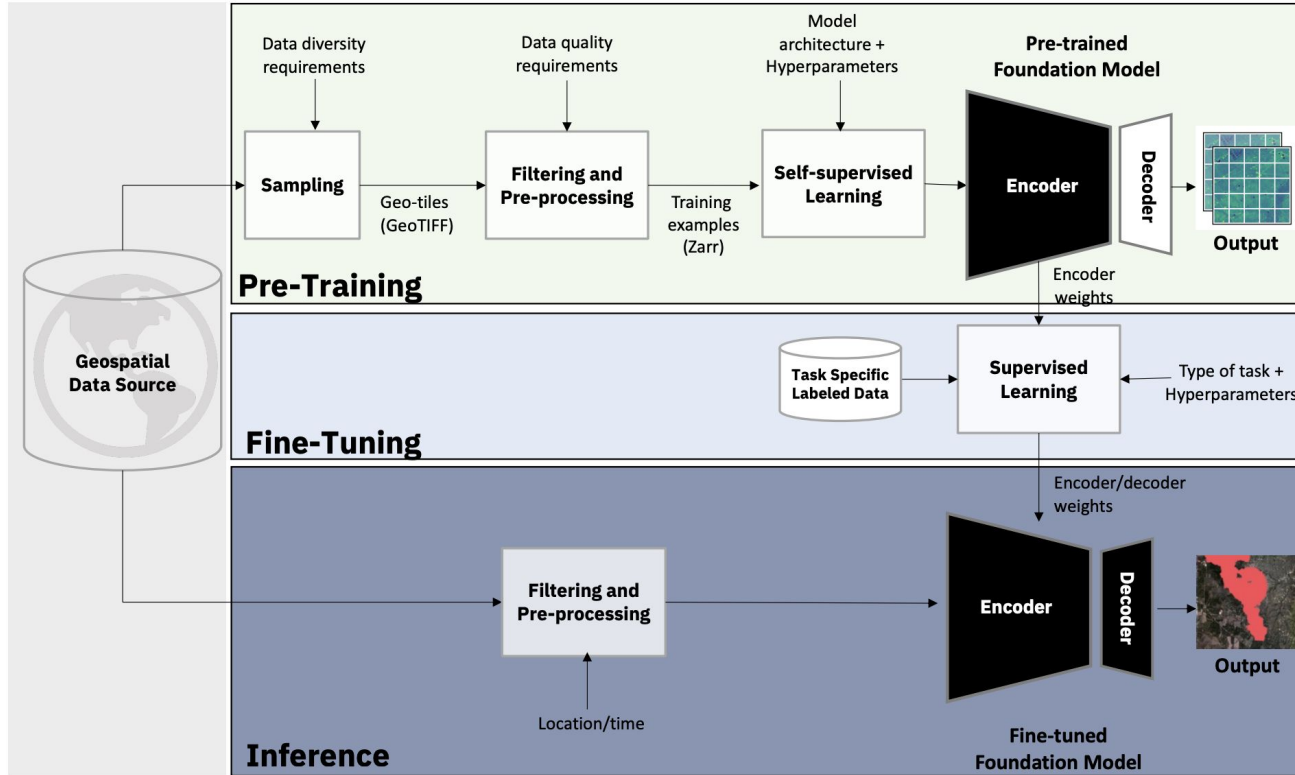
1. **GPT-3 (OpenAI):**
 - **Parameters:** 175 billion
 - **Dataset Size:** 410 billion tokens (source: Common Crawl, WebText2, Books1, Books2, Wikipedia)
2. **BERT (Google):**
 - **Parameters:** 340 million (BERT Large)
 - **Dataset Size:** 3.3 billion tokens (source: BooksCorpus and English Wikipedia)
3. **CLIP (OpenAI):**
 - **Parameters:** 428 million (Transformer + ViT-Large)
 - **Dataset Size:** 400 million image-text (source: Common Crawl)

Prithvi

A transformer-based geospatial foundational model

- Model size: 100 million parameters
- Dataset: 1TB of multispectral satellite imagery from the Harmonized Landsat-Sentinel 2 (HLS) dataset
- HLS dataset:
 - Spatial resolution: 30 m
 - Operational Land Imager (OLI) and Multi-Spectral Instrument (MSI) aboard the Landsat 8/9 and Sentinel-2 remote sensing satellites
 - Global observations of the land every 2–3 days.

Prithvi's framework



Prithvi's framework

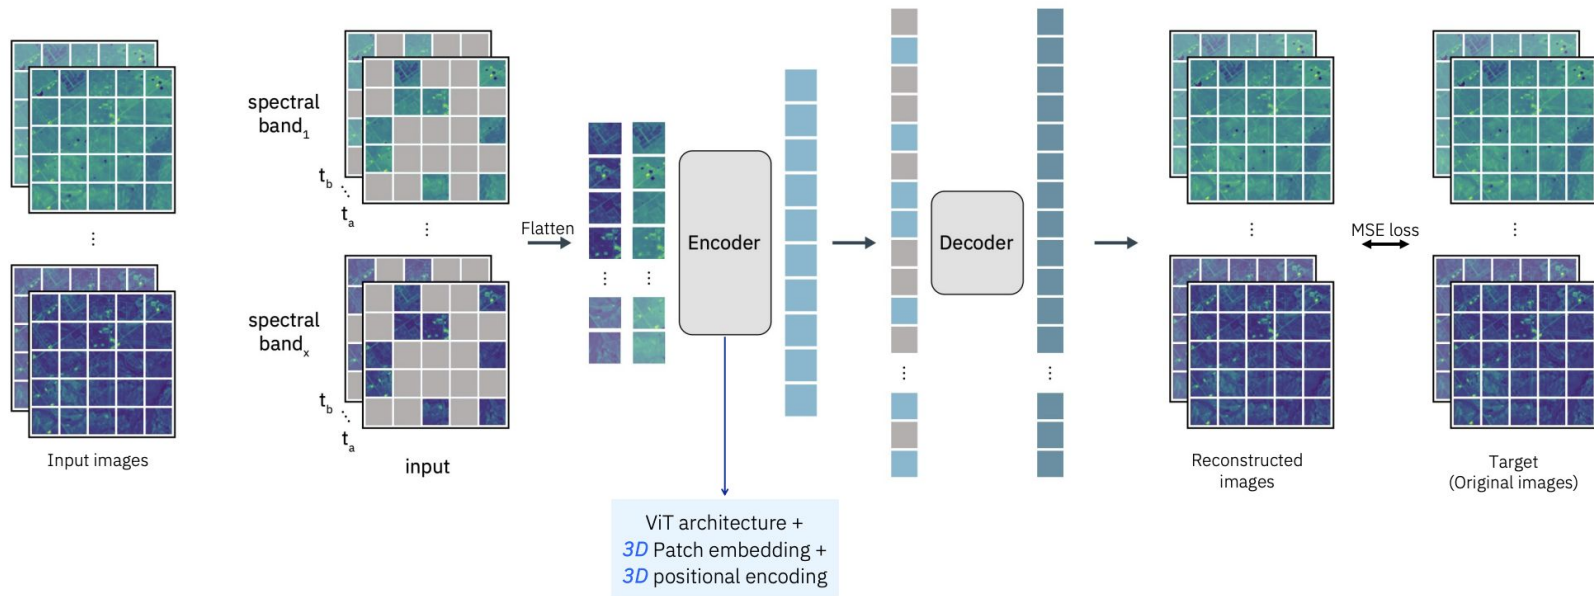


Fig. 3: The masked autoencoder (MAE) structure for pre-training Prithvi on large-scale multi-temporal and multi-spectral satellite images.

Fine-tuning

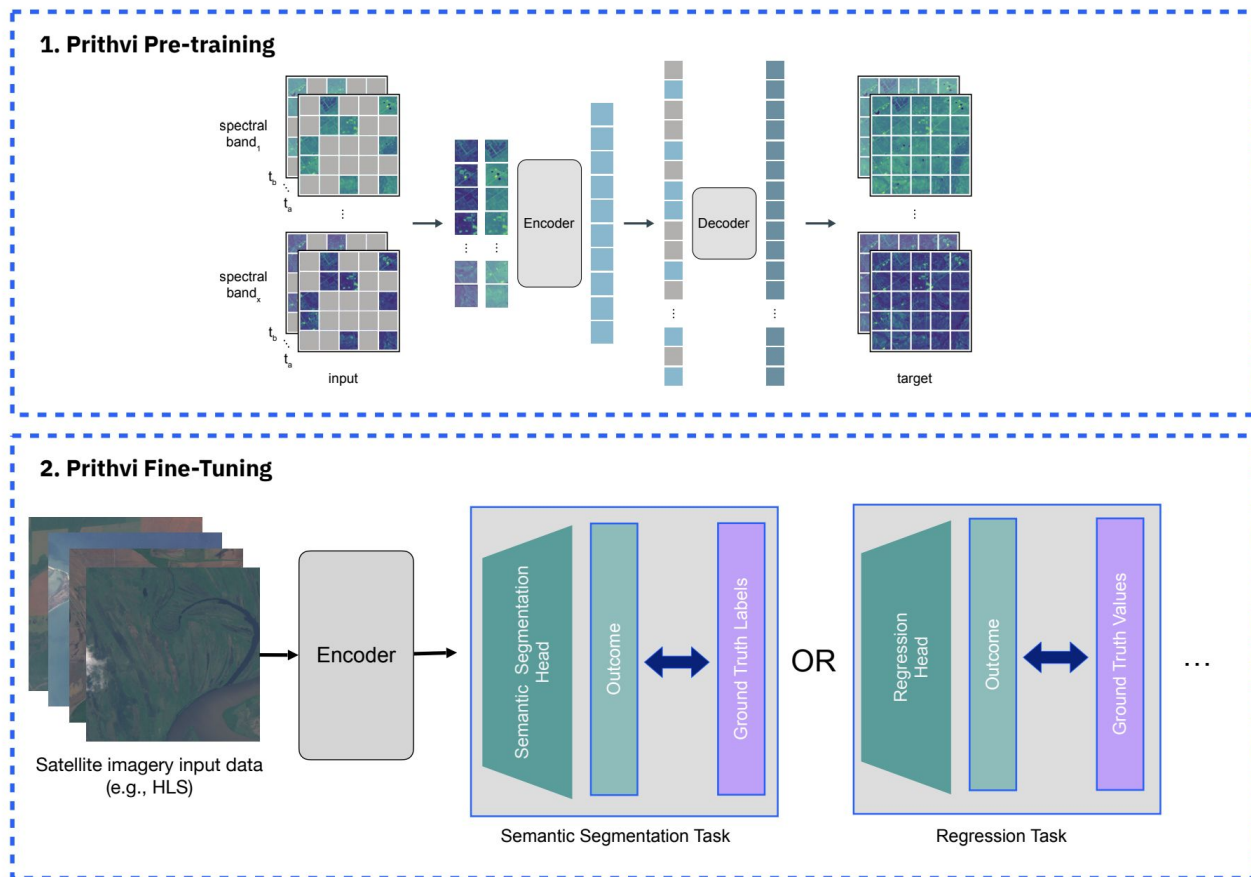
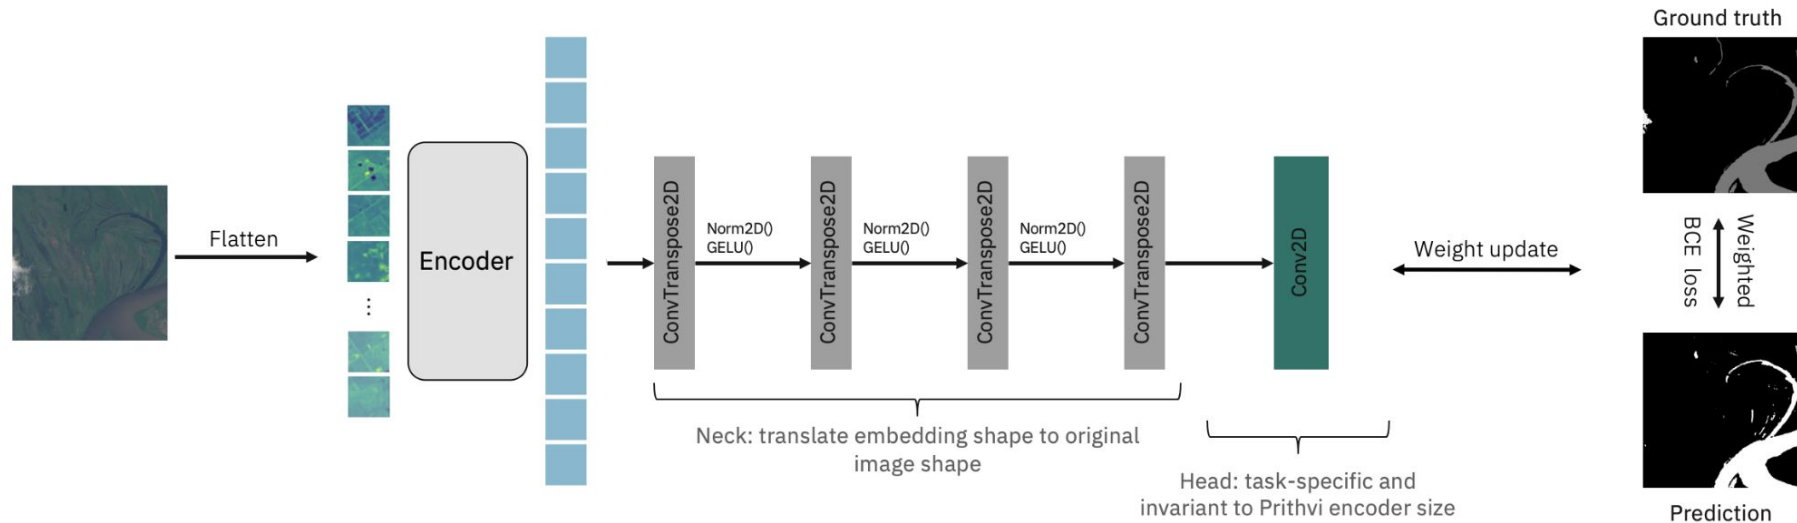
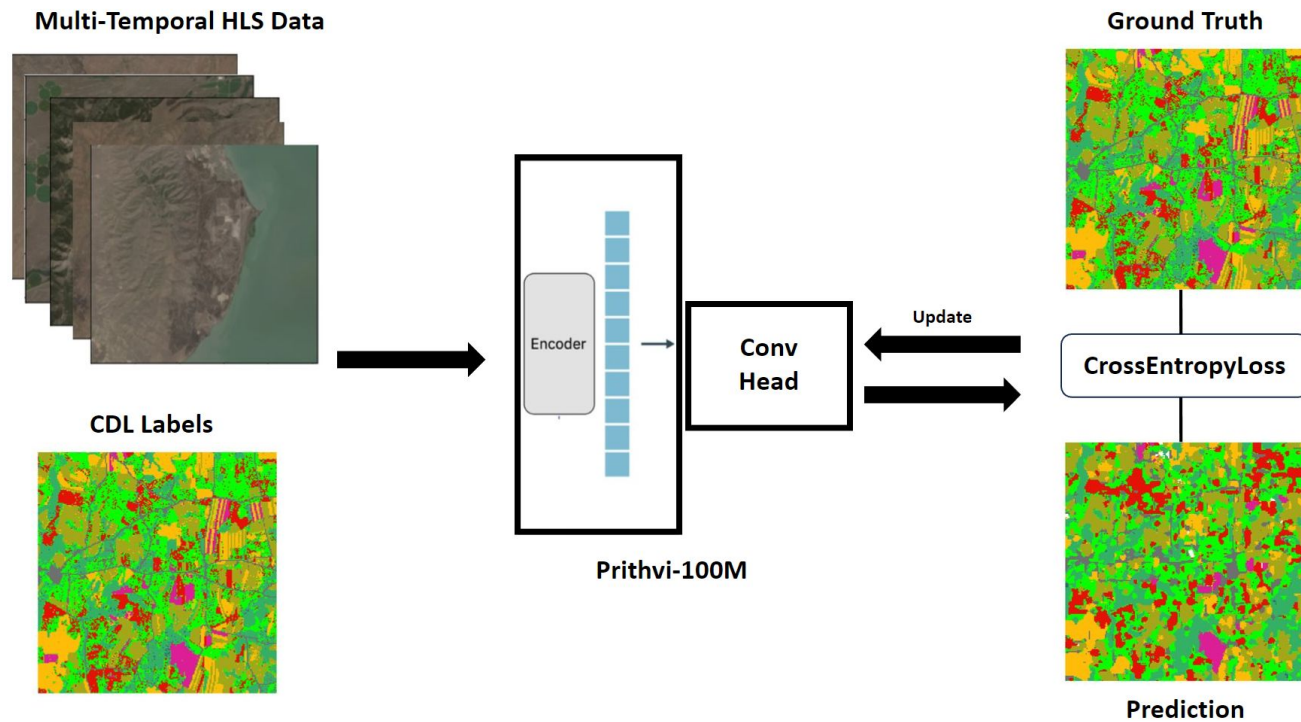


Fig. 4: Pre-training and fine-tuning in Prithvi for various types of downstream tasks.

Flood detection



Crop classification



Conclusion

- Foundation Models: large models that are trained on extensive datasets.
- Advantage: A model that already understands the “basics” about the data
- Few-shot learning: Much smaller number of samples to tune it for your application
- Fine-tuning: Use the model’s knowledge for a downstream application

Tutorial

Let's explore Prithvi and its fine-tuned models:

- 1) Predict masked regions of an image (pretrained)
- 2) Predict regions with likelihood of flooding (fine-tuned)
- 3) Predict different types of crop (fine-tuned)
- 4) **BONUS: download an image for your hometown and do the prediction tasks above [did the model get it right?]**